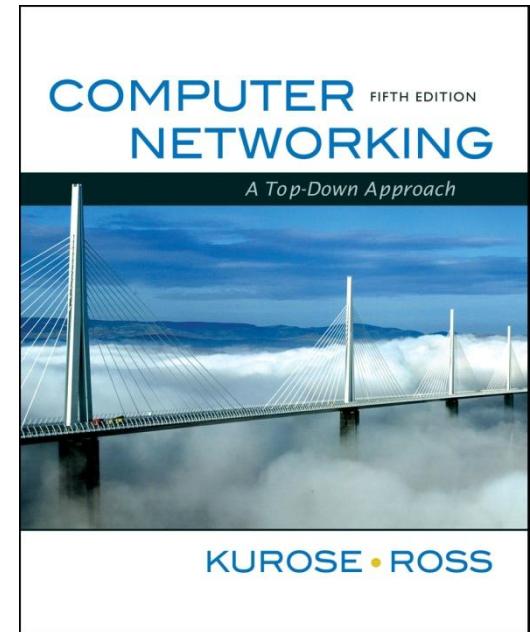


Chapter 5

Link Layer and LANs



A note on the use of these ppt slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:

- ❖ If you use these slides (e.g., in a class) in substantially unaltered form, that you mention their source (after all, we'd like people to use our book!)
- ❖ If you post any slides in substantially unaltered form on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR

All material copyright 1996-2010
J.F Kurose and K.W. Ross, All Rights Reserved

*Computer Networking:
A Top Down Approach
5th edition.
Jim Kurose, Keith Ross
Addison-Wesley, April
2009.*

Link Layer

5.1 Introduction and services

5.2 Error detection and correction

5.3 Multiple access protocols

New: WiFi

5.4 Link-layer Addressing

5.5 Ethernet

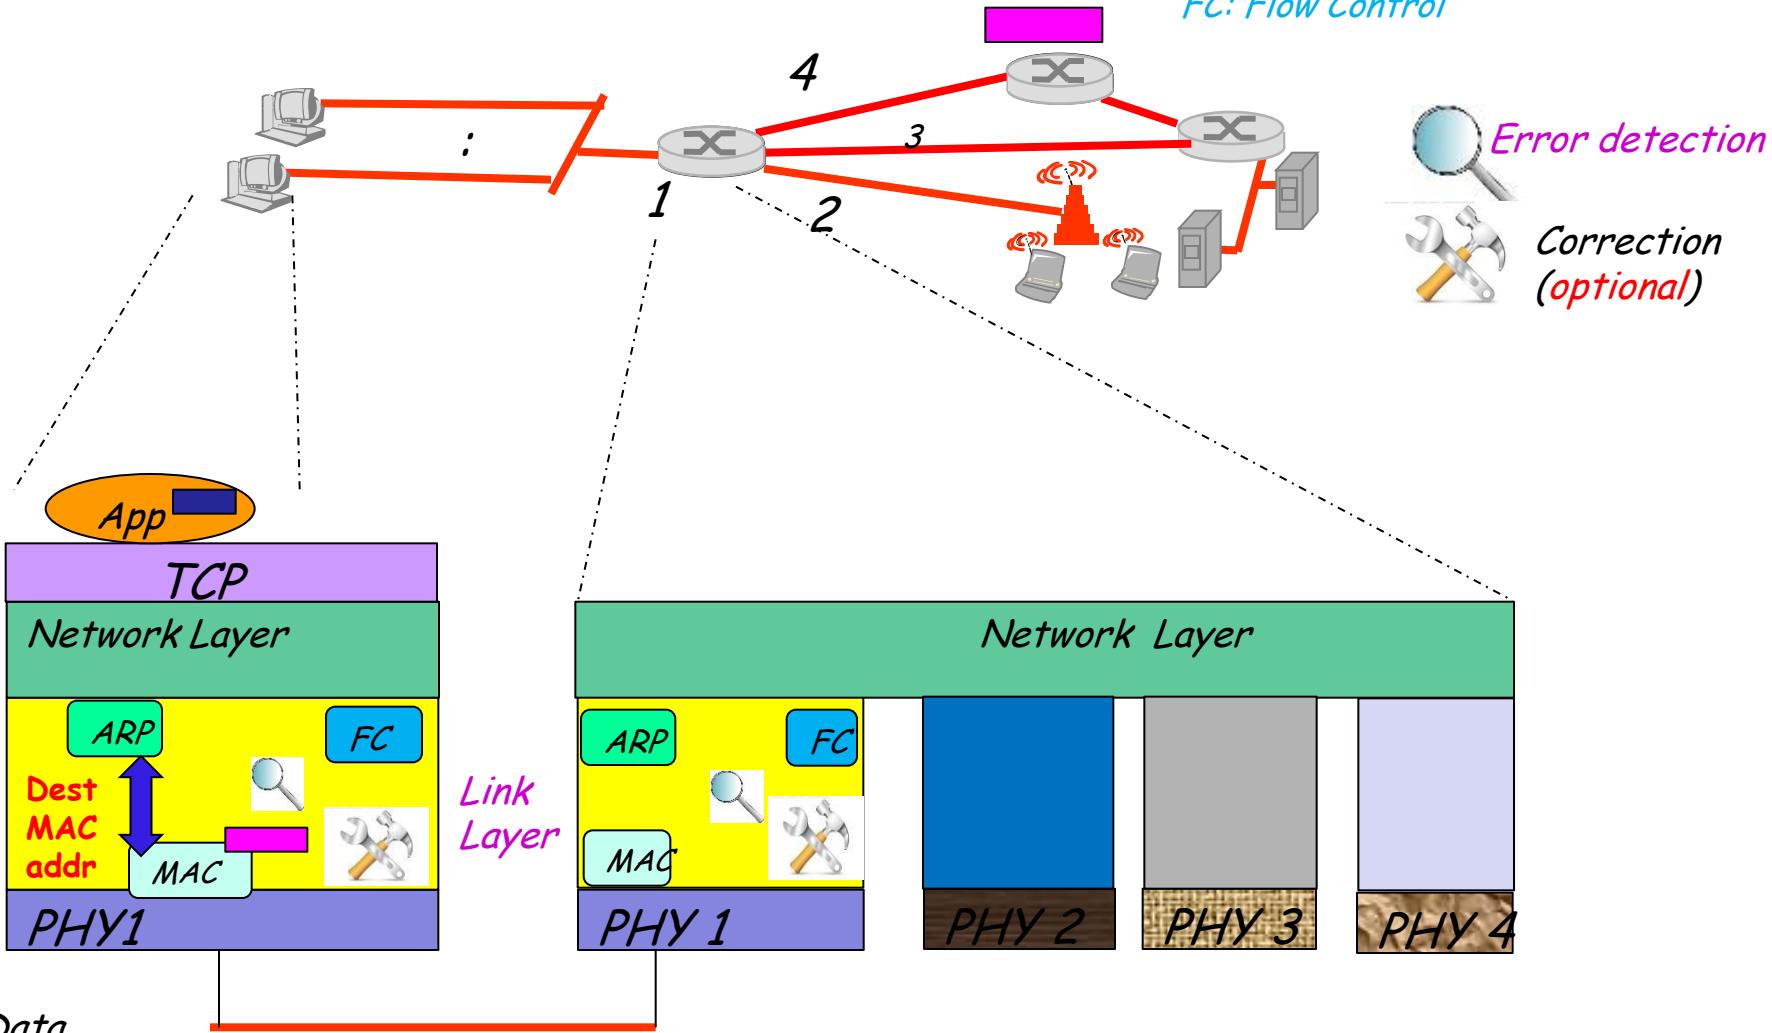
5.6 Link-layer switches

3.4 New: Reliable data transfer + flow control

Link layer: context

MAC: Medium Access Control

FC: Flow Control



ARP (Address Resolution Protocol): IP address \rightarrow MAC address
(%ipconfig /all \leftarrow for MAC addresses.)

PHY: Physical layer \leftarrow hardware

Link Layer: Introduction

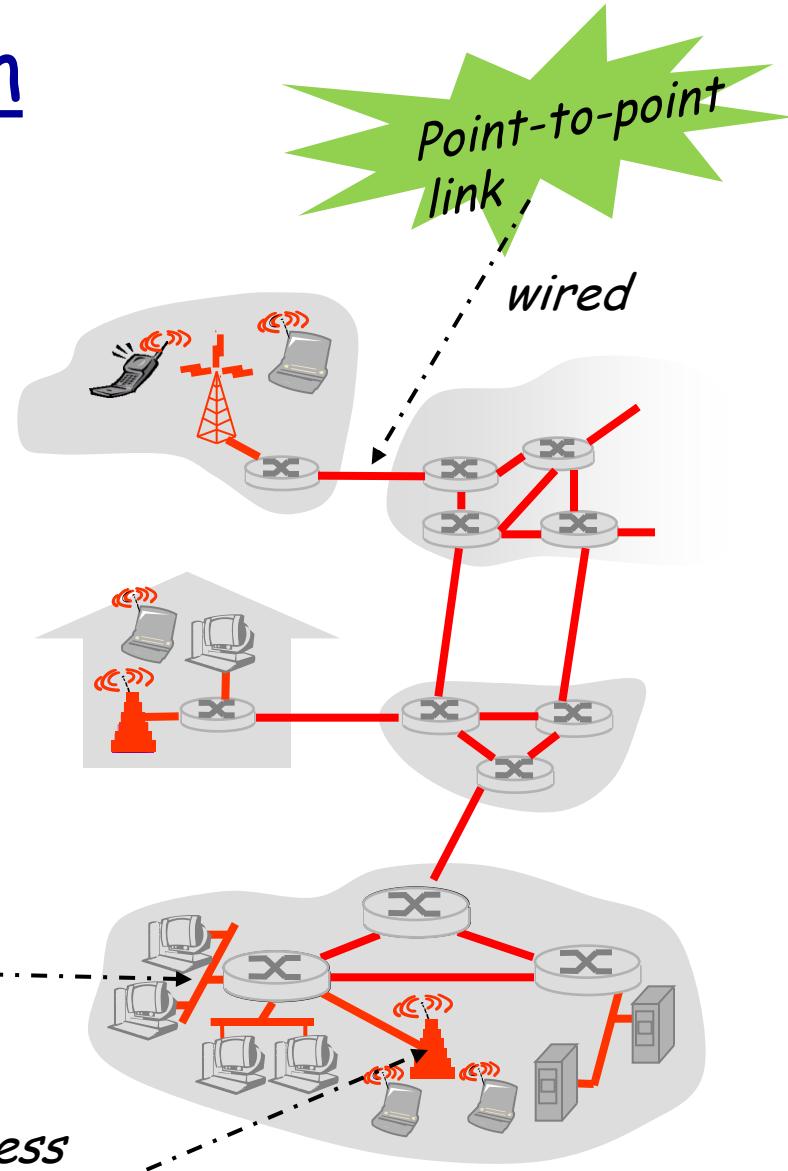
Terminology:

- ❖ hosts and routers are **nodes**
- ❖ **communication channels** that connect adjacent nodes are **links**
 - wired links
 - wireless links

We will focus on
broadcast links.

wired

wireless



Chapter 5: The Data Link Layer

Our goals:

- ❖ understand principles behind data link layer services:

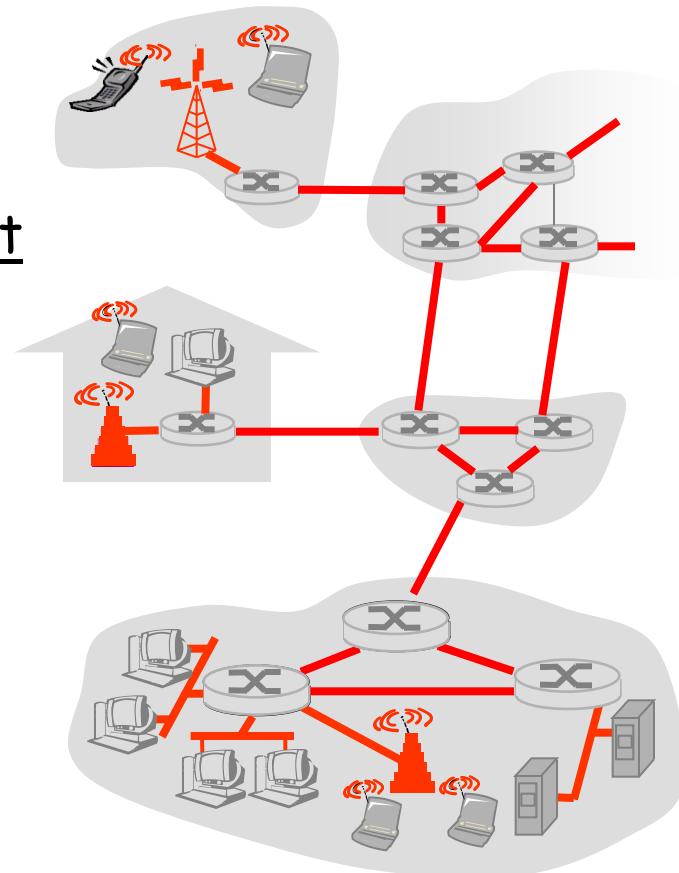
1

link access by sharing a broadcast channel: multiple access

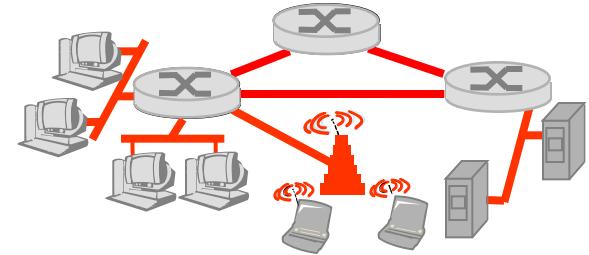
- Instruct the hardware (PHY layer) **when** to transmit (... MAC protocols)

2

link layer addressing



Link Layer Services (more)



❖ *error detection:*

3

- errors caused by signal attenuation and noise
- receiver detects presence of errors

❖ *error correction:*

- 4
- receiver identifies and *corrects* bit error(s) without resorting to *retransmission*
 - receiver signals sender for *retransmission*

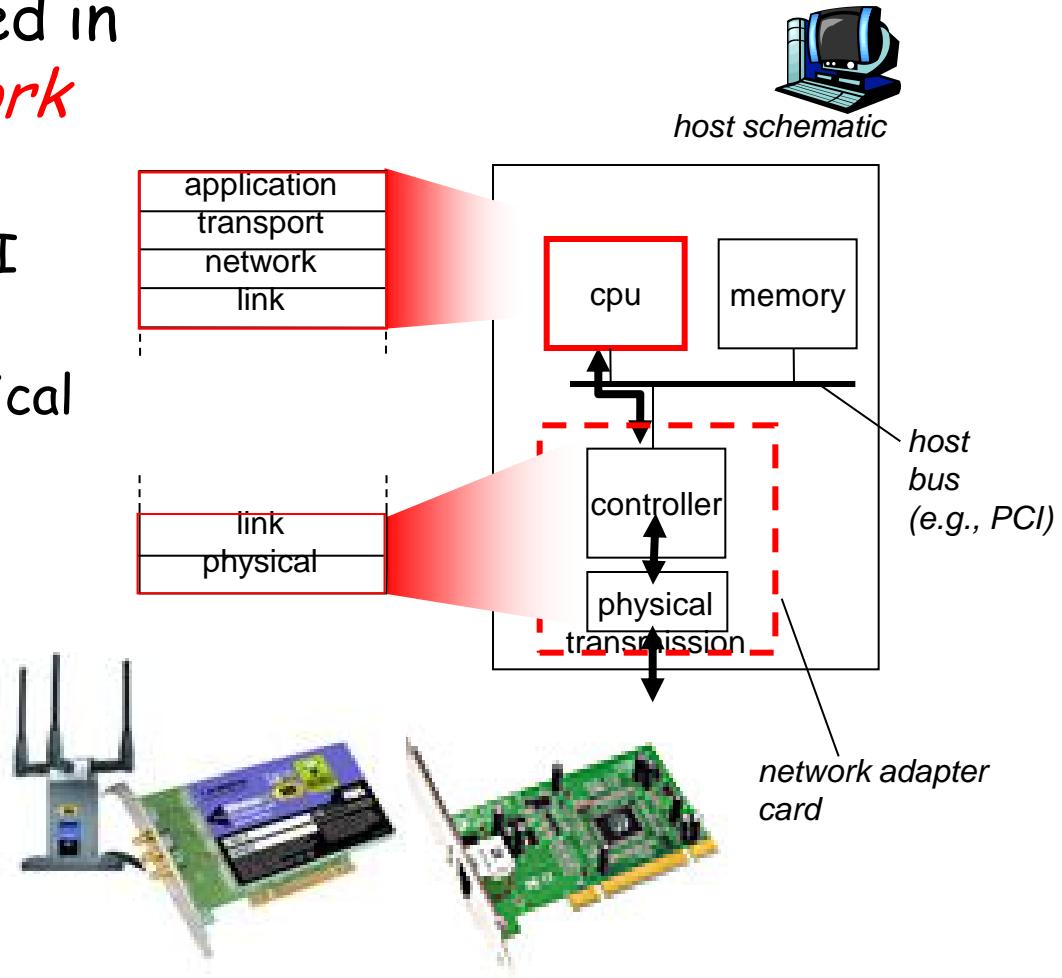
❖ *flow control:*

- pacing between (adjacent) sending and receiving nodes

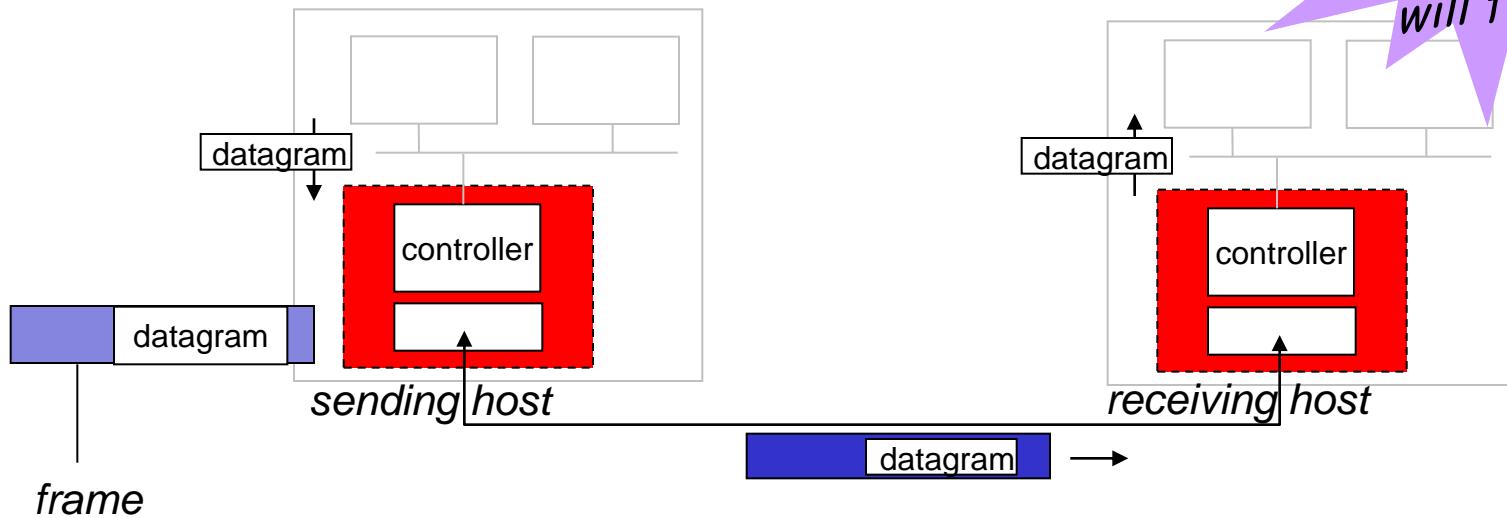


Where is the link layer implemented?

- ❖ in each and every host
- ❖ link layer implemented in "adaptor" (aka *network interface card NIC*)
 - Ethernet card, PCMCIA card, 802.11 card
 - implements link, physical layer
- ❖ attaches into host's system buses
- ❖ combination of **hardware, software, firmware**



Adaptors Communicating



Wait...
details
will follow

❖ sending side:

- encapsulates datagram in frame
- adds **error checking bits**, and sequence # for flow control

❖ receiving side

- looks for errors
- extracts datagram, passes to upper layer at receiving side

Link Layer

5.1 Introduction and services

5.2 Error detection and correction

5.3 Multiple access protocols

5.4 Link-layer Addressing

5.5 Ethernet

5.6 Link-layer switches

5.7 PPP

5.8 Link virtualization:
MPLS

5.9 A day in the life of a web request

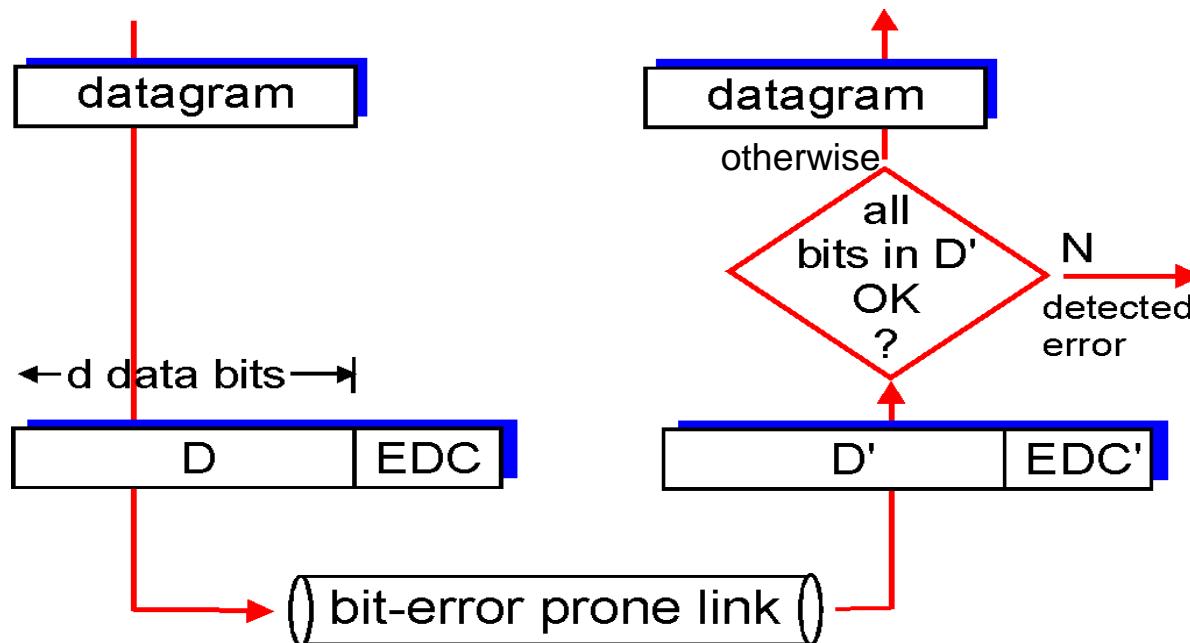
3.4 Reliable data transfer

Error Detection

EDC= Error Detection and **C**orrection bits (redundancy)

D = Data protected by error checking, may include header fields

- Error detection **not** 100% reliable!
 - protocol may miss some errors, but rarely
 - **larger** EDC field yields **better** detection and correction



What are all the *error detection mechanisms* that I will learn



Parity bits (today ...)

You will learn ...



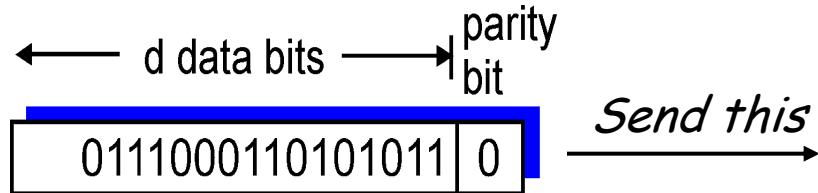
CRC:
Cyclic Redundancy Check (Link layer)

Internet Checksum
(Network Layer, Transport Layer)

Parity Checking

Single Bit Parity:

Detect single bit errors



(Example of **odd** parity)

Two kinds of parity:

* **Odd parity:** Set the parity bit to 1 or 0 to make the total # of 1's an odd number.

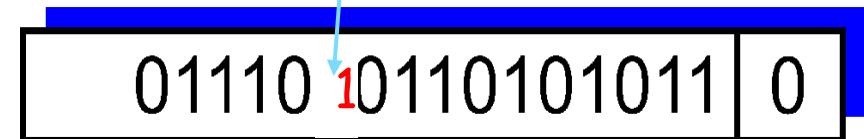
* **Even parity:** Set the parity bit to 1 or 0 to make the total # of 1's an even number



You find the frame to be erroneous, but you cannot locate the bit error.



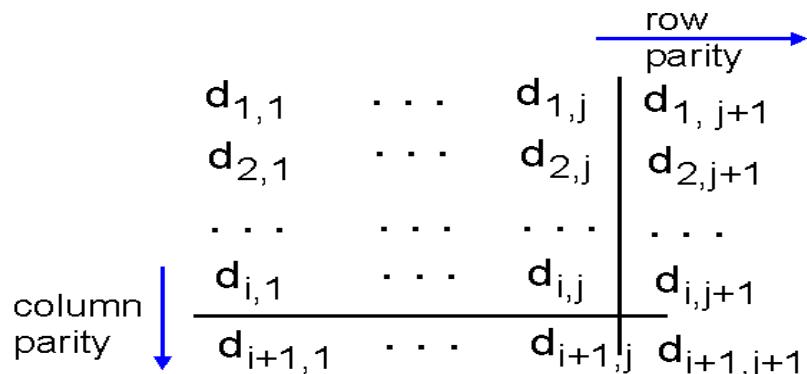
..... But, this is received



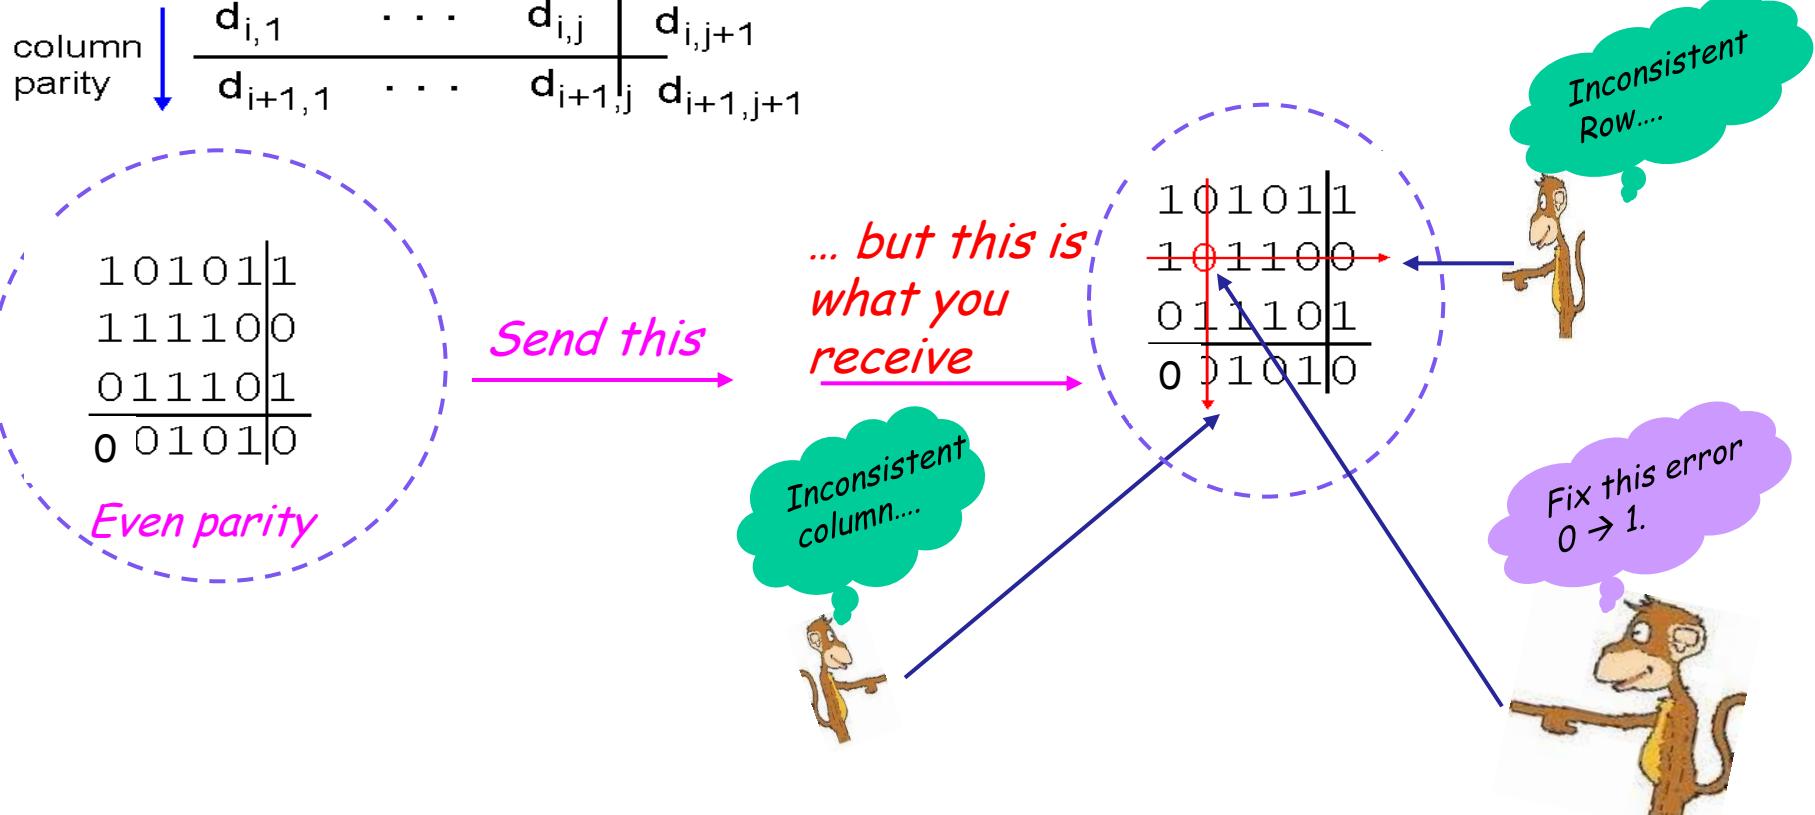
When you compute the parity bit, you know that it should be a '1'.

Use two-dimensional bit parity to detect and correct single bit errors

Parity Checking (Two dimensional)



Detect **and** correct single bit errors



Applications of Parity Checking

Where is parity checking used?

Not generally in network protocols
these days...



RAID: Redundant Array of Inexpensive Drives

*DES: Data Encryption Standard for
symmetric-key cryptography*



But, it is used
here and here
and here



Media access is a **noisy process.....**
because of **imperfect fabrication**
and **hardware aging**.



What kinds of questions
can be asked on the exams?



For the following **four bytes**
of data, compute the
two dimensional **odd** parity bits.

1	1	0	0	1	1	0	1	
1	0	1	1	1	0	1	1	
0	1	1	0	1	0	1	1	
<u>1</u> 0000101								



For the data and parity bits
generated above, **introduce four bit
errors** that the **receiver** will not be
able to detect...



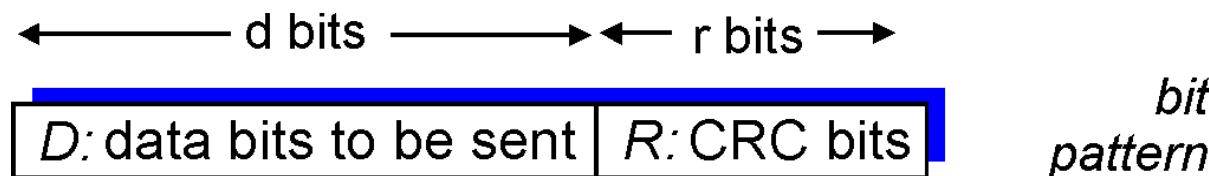
ECE358
Mid-term S'12



ECE358
Mid-term F'12

Checksumming: Cyclic Redundancy Check

- ❖ view data bits, D , as a binary number
- ❖ choose $r+1$ bit pattern (generator), G
- ❖ **goal:** choose r CRC bits, R , such that
 - $\langle D, R \rangle$ exactly divisible by G (modulo 2) ← zero remainder
- ❖ receiver knows G :
 - Divides $\langle D, R \rangle$ by G . If non-zero remainder: error detected!
- ❖ can detect all burst errors less than $r+1$ bits
- ❖ widely used in practice (Ethernet, 802.11 WiFi, ATM)



$D * 2^r \text{ XOR } R$ *mathematical formula*

CRC Example:

Want:

$$D \cdot 2^r \text{ XOR } R = nG$$

equivalently:

$$D \cdot 2^r = nG \text{ XOR } R$$

equivalently:

if we **divide** $D \cdot 2^r$ by G ,
we get remainder R

Properties of XOR:

$$A \text{ XOR } A = 0$$

$$A \text{ XOR } 0 = A$$

$$(A \text{ XOR } B) \text{ XOR } C = A \text{ XOR } (B \text{ XOR } C)$$

$$R = \text{remainder} \left[\frac{D \cdot 2^r}{G} \right]$$

Process: Calculation of CRC

If the **input** bit above the **leftmost divisor** bit is **1**,
the divisor is **XORed** into the input.

Else (the **input** bit above the **leftmost divisor** bit is **0**) do nothing.

The **divisor** is then **shifted one bit to the right** (\Rightarrow)

The process is **repeated until**
the **divisor reaches the right-hand end of the input row**.

Input: $D.2^r$ \longrightarrow $1\ 0\ 1\ 1\ 1\ 0\ \underline{0\ 0\ 0}$

Divisor: G \longrightarrow $1\ 0\ 0\ 1$

Align input and divisor on MSB

$0\ 0\ 1\ 0\ 1\ 0\ \underline{0\ 0\ 0}$

$D = 101110$

$1\ 00\ 1$

$0\ 0\ 1\ 0\ 1\ 0\ \underline{0\ 0\ 0}$

$G = 1001$

$1\ 00\ 1$

$0\ 0\ 0\ 0\ 1\ 1\ \underline{0\ 0\ 0}$

$1\ 00\ 1$

$0\ 0\ 0\ 0\ 1\ 1\ \underline{0\ 0\ 0}$

$1\ 0\ 0\ 1$

$0\ 0\ 0\ 0\ 0\ 1\ \underline{0\ 1\ 0}$

$1\ 00\ 1$

$0\ 0\ 0\ 0\ 0\ 0\ \underline{0\ 1\ 1}$

R

Detailed
Calculation of CRC



What questions
can be asked on the exams?



For the data bits $D = 11001101$ and
generator bits $G = 1011$,
compute the CRC bits.



What is the *disadvantage* of making
the CRC field as long as the data portion of a frame?

D.2^r: 11001101 000?
G: 1011

What is the length
of the CRC field?



Link Layer

5.1 Introduction and services

5.2 Error detection and correction

5.3 Multiple access protocols

5.4 Link-layer Addressing

5.5 Ethernet

5.6 Link-layer switches

5.7 PPP

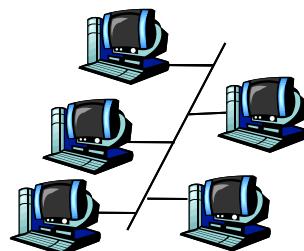
5.8 Link virtualization:
MPLS

5.9 A day in the life of a web request

Multiple Access Links and Protocols

Two types of "links":

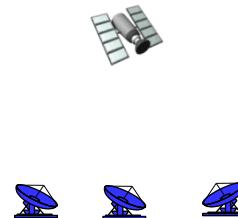
- ❖ **point-to-point**
 - PPP for dial-up access
- ❖ **broadcast** (shared wire or medium)
 - old-fashioned Ethernet
 - 802.11 wireless LAN



shared wire (e.g.,
cabled Ethernet)



shared RF
(e.g., 802.11 WiFi)



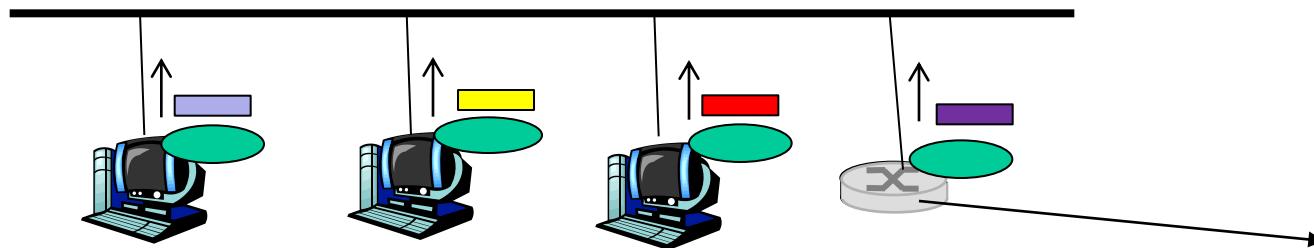
shared RF
(satellite)



humans at a
cocktail party
(shared air, acoustical)

Multiple Access protocols

- ❖ single shared broadcast channel
- ❖ two or more simultaneous transmissions cause collision
- ❖ multiple access protocol
 - distributed algorithm that determines when a node can transmit

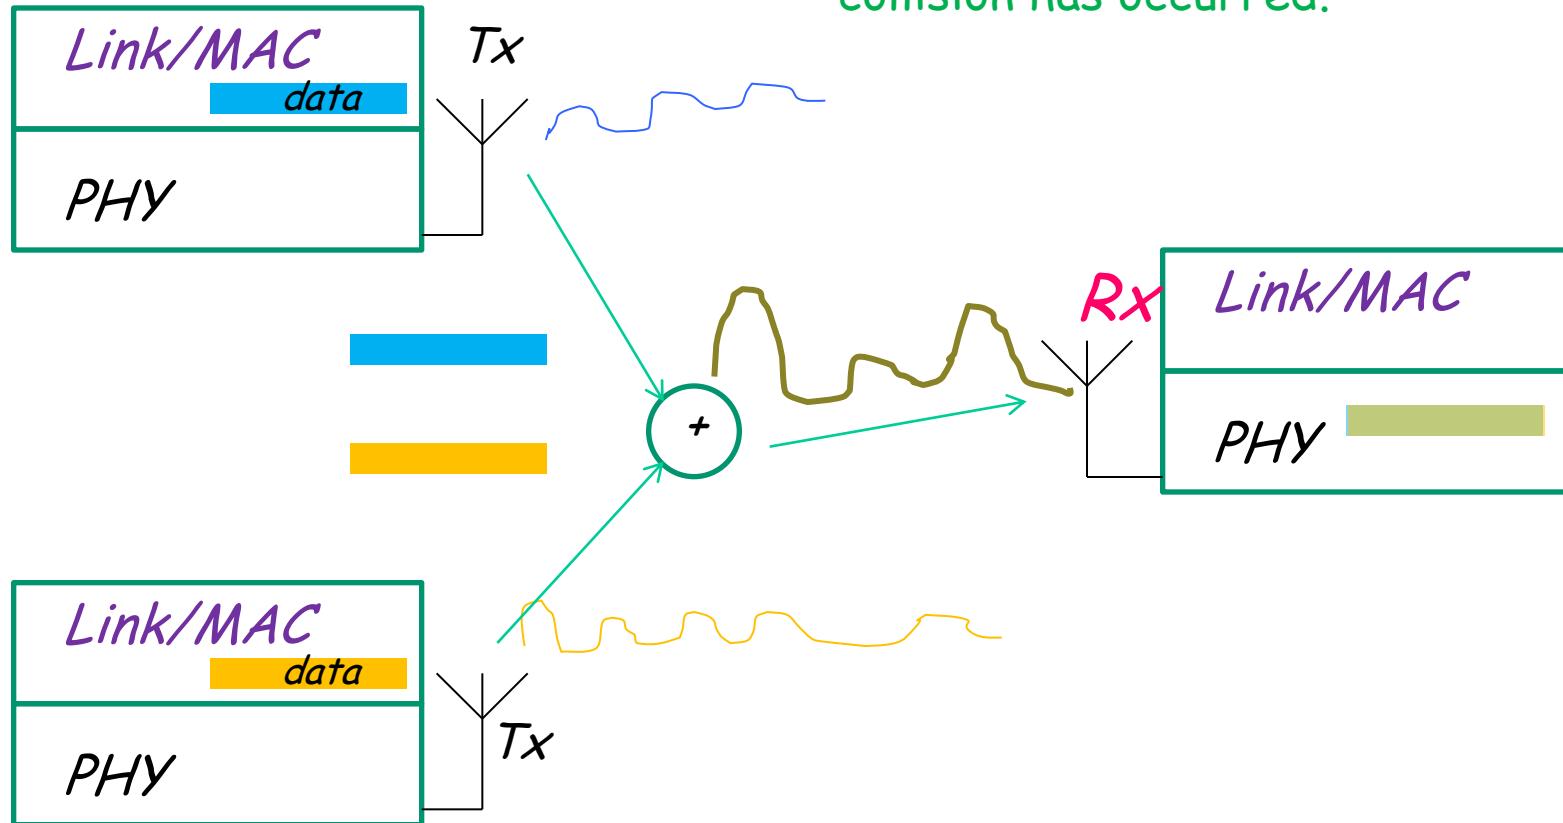


Packet Collision



Note

Packet collision occurs at the receiver, but transmitters must know that collision has occurred.



Ideal Multiple Access Protocol

Broadcast channel of rate R bps

1. when one node wants to transmit, it can send at rate R.
2. when M nodes want to transmit, each can send at average rate R/M
3. fully decentralized:
 - no special node to coordinate transmissions
 - no synchronization of clocks, slots
4. simple (plug-and-play, no complex hardware, ...)

MAC Protocols: a taxonomy

Three broad classes:

- ❖ **Channel Partitioning** (Commonly done in cellular networks)

- divide channel into smaller “pieces” (time slots, frequency)
 - allocate piece to node for **exclusive use**

- ❖ **Random Access**

- channel not divided, allow collisions
 - “recover” from collisions

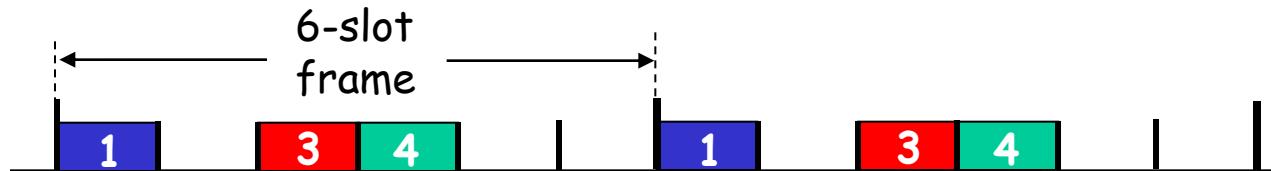
- ❖ **“Taking turns”**

- nodes take turns

Channel Partitioning MAC protocols: TDMA

TDMA: time division multiple access

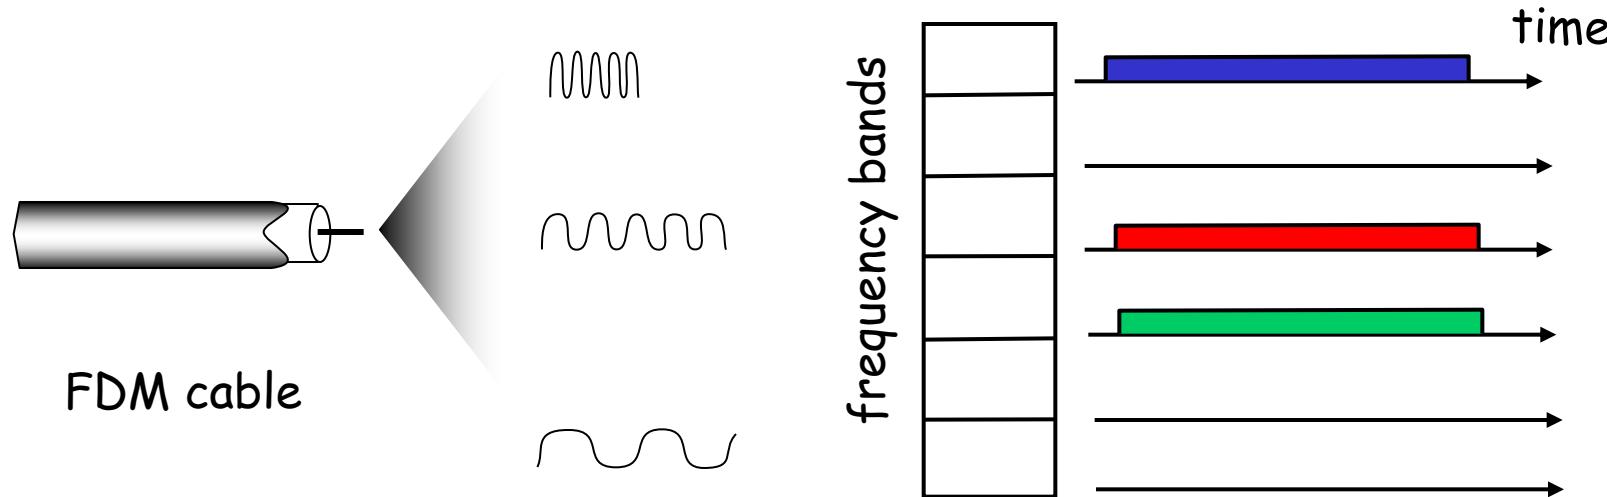
- ❖ access to channel in "rounds"
- ❖ each station gets fixed length slot (length = pkt trans time) in each round
- ❖ unused slots go idle
- ❖ example: 6-station LAN, 1,3,4 have pkt, slots 2,5,6 idle



Channel Partitioning MAC protocols: FDMA

FDMA: frequency division multiple access

- ❖ channel spectrum divided into frequency bands
- ❖ each station assigned fixed frequency band
- ❖ unused transmission time in frequency bands go idle
- ❖ example: 6-station LAN, 1,3,4 have pkt, frequency bands 2,5,6 idle



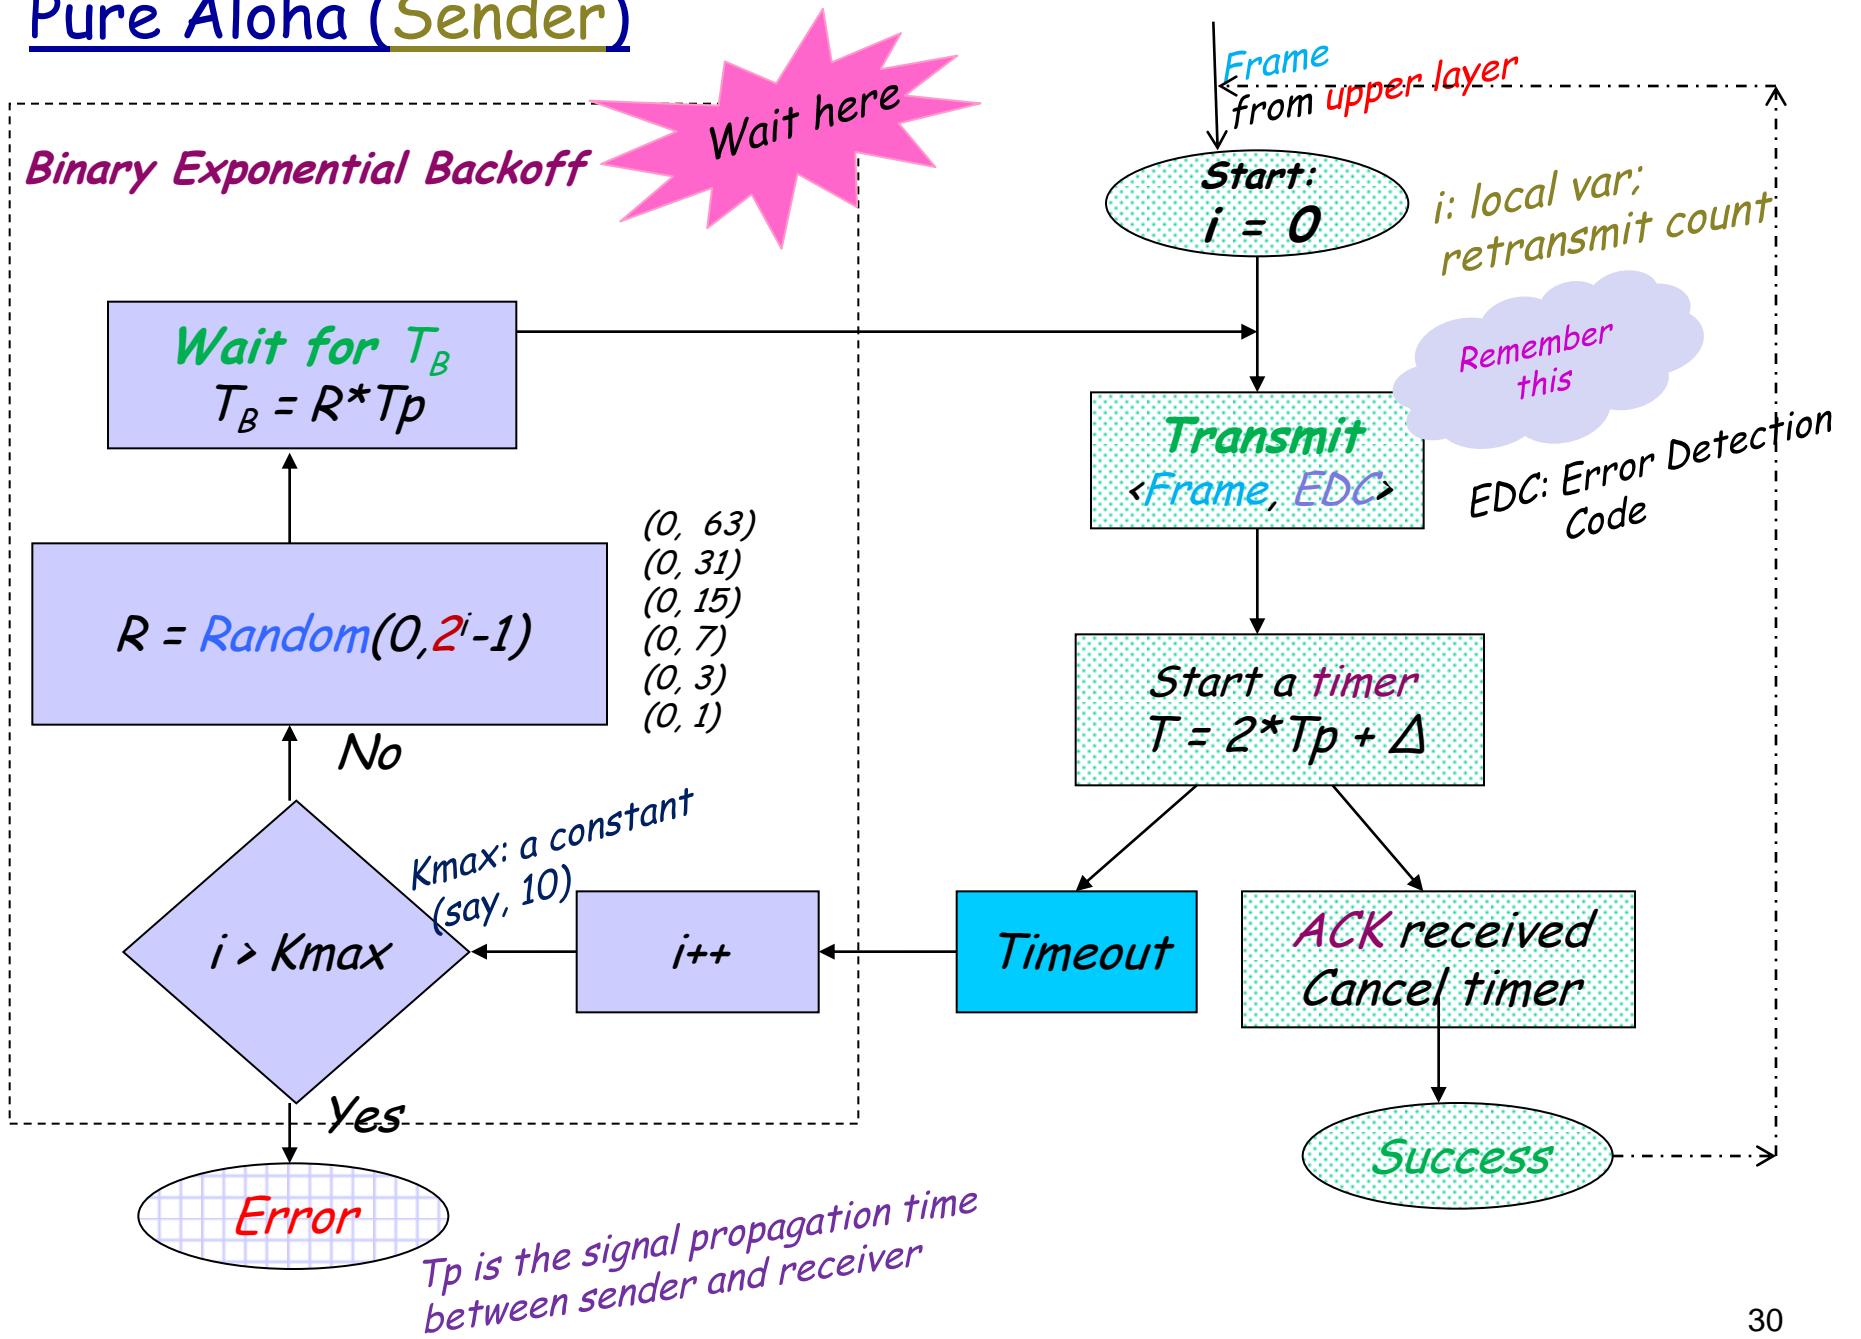
Random Access Protocols

- ❖ When node has packet to send
 - transmit at full channel data rate R.
 - no a priori coordination among nodes (...partial exception in WiFi)
- ❖ two or more transmitting nodes → "collision",
- ❖ random access MAC protocol specifies:
 - how to **detect** collisions
 - how to **recover** from collisions
- ❖ Examples of random access MAC protocols:
 - ALOHA and slotted ALOHA
 - CSMA/**CD**, CSMA/**CA** CSMA: Carrier Sense Multiple Access
 - **CD**: Collision Detection ← Reactive scheme
 - **CA**: Collision Avoidance ← Preventive scheme

Aloha Protocol

- ❖ Developed in the 1970s at U of Hawaii
- ❖ To interconnect terminals with mainframes
- ❖ LAN/ WLAN: Possible, but not used
- ❖ GSM: Cell phones use this protocol to request a channel from the base stations
- ❖ Two types
 - Pure Aloha (Continuous time)
 - Slotted Aloha

Pure Aloha (Sender)

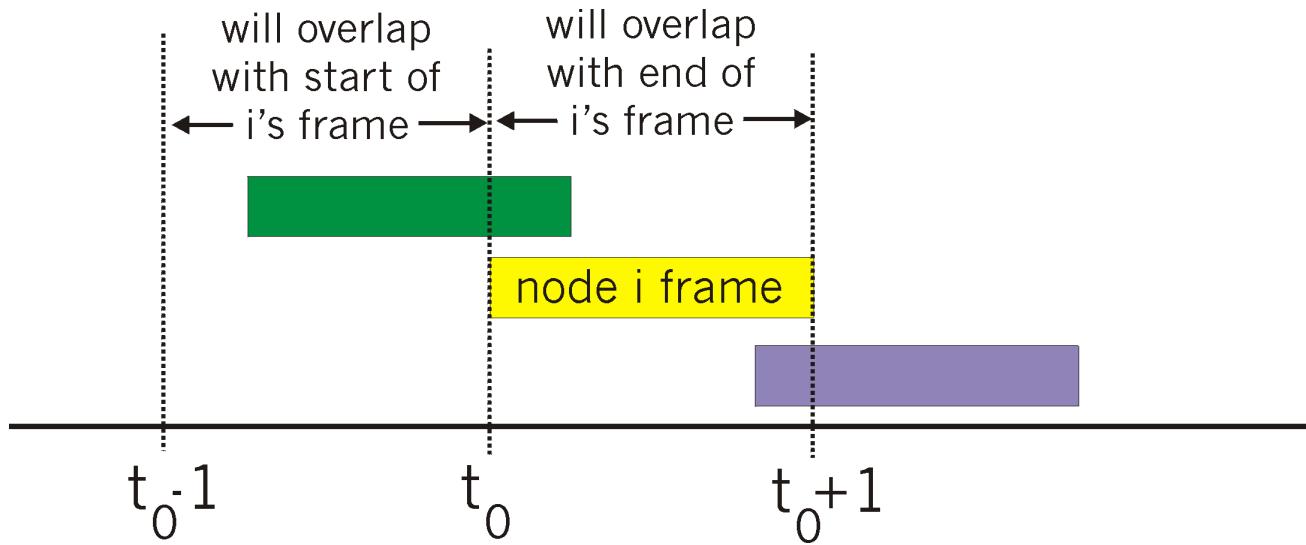


Pure Aloha

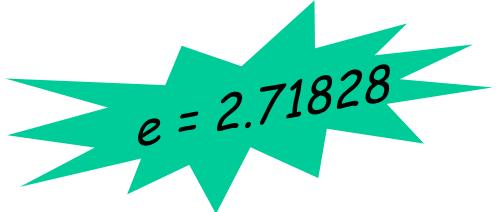


Pure (unslotted) ALOHA

- ❖ unslotted Aloha: simpler, no synchronization
- ❖ when frame first arrives
 - transmit immediately
- ❖ **collision probability increases:**
 - frame sent at t_0 collides with other frames sent in $[t_0-1, t_0+1]$



Pure Aloha efficiency


$$e = 2.71828$$

Throughput of Pure Aloha =

Total input rate (G) * Prob. of successful Packet Trans.

$$= G * e^{-2G}$$

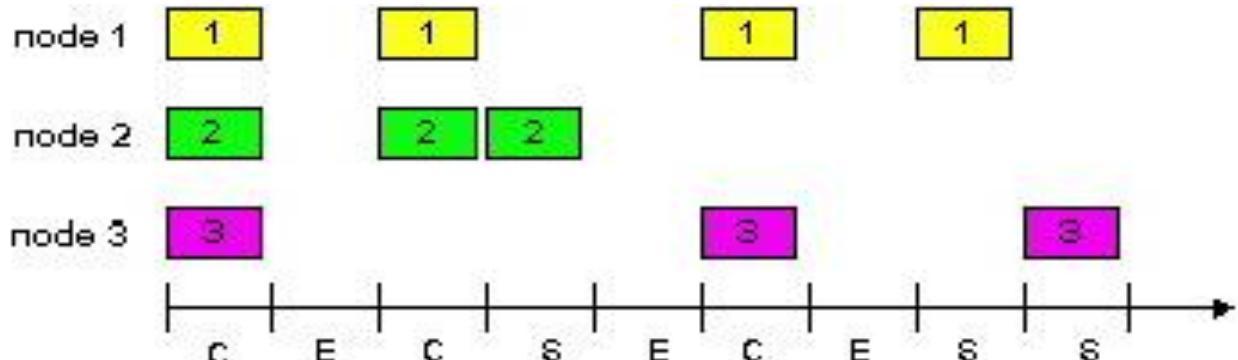
G : Number of packets in X sec.

where X is the packet Tx time.

Max throughput occurs at $G = 0.5$.

$$\text{Max Throughput} = 1/(2e) = .18$$

Slotted ALOHA



Assumptions:

- ❖ all frames are **same size**
- ❖ time divided into **equal size slots** (time to transmit 1 frame)
- ❖ nodes start to **transmit only at slot beginning**
- ❖ nodes are **synchronized**

Operation:

- ❖ when node obtains fresh frame, **transmit in next slot**
- ❖ if Tx is not successful, retransmit in following slots after some wait..

Slotted ALOHA

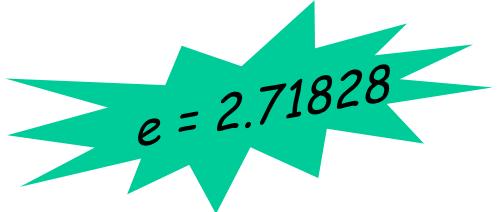
Pros

- ❖ single active node can continuously transmit at full rate of channel
- ❖ partly decentralized: only slots in nodes need to be in sync

Cons

- ❖ collisions, wasting slots
- ❖ idle slots
- ❖ Keep transmitting even if there is collision
- ❖ clock synchronization among nodes

Slotted Aloha efficiency


$$e = 2.71828$$

Throughput of Slotted Aloha =

Total input rate (G) * Prob. of successful Packet Trans.

$$= G * e^{-G}$$

G : Number of packets in X sec.

where X is the packet Tx time.

Max throughput occurs at $G = 1$.

$$\text{Max Throughput} = 1/e = .37$$

CSMA (Carrier Sense Multiple Access)

CSMA: listen before you transmit:

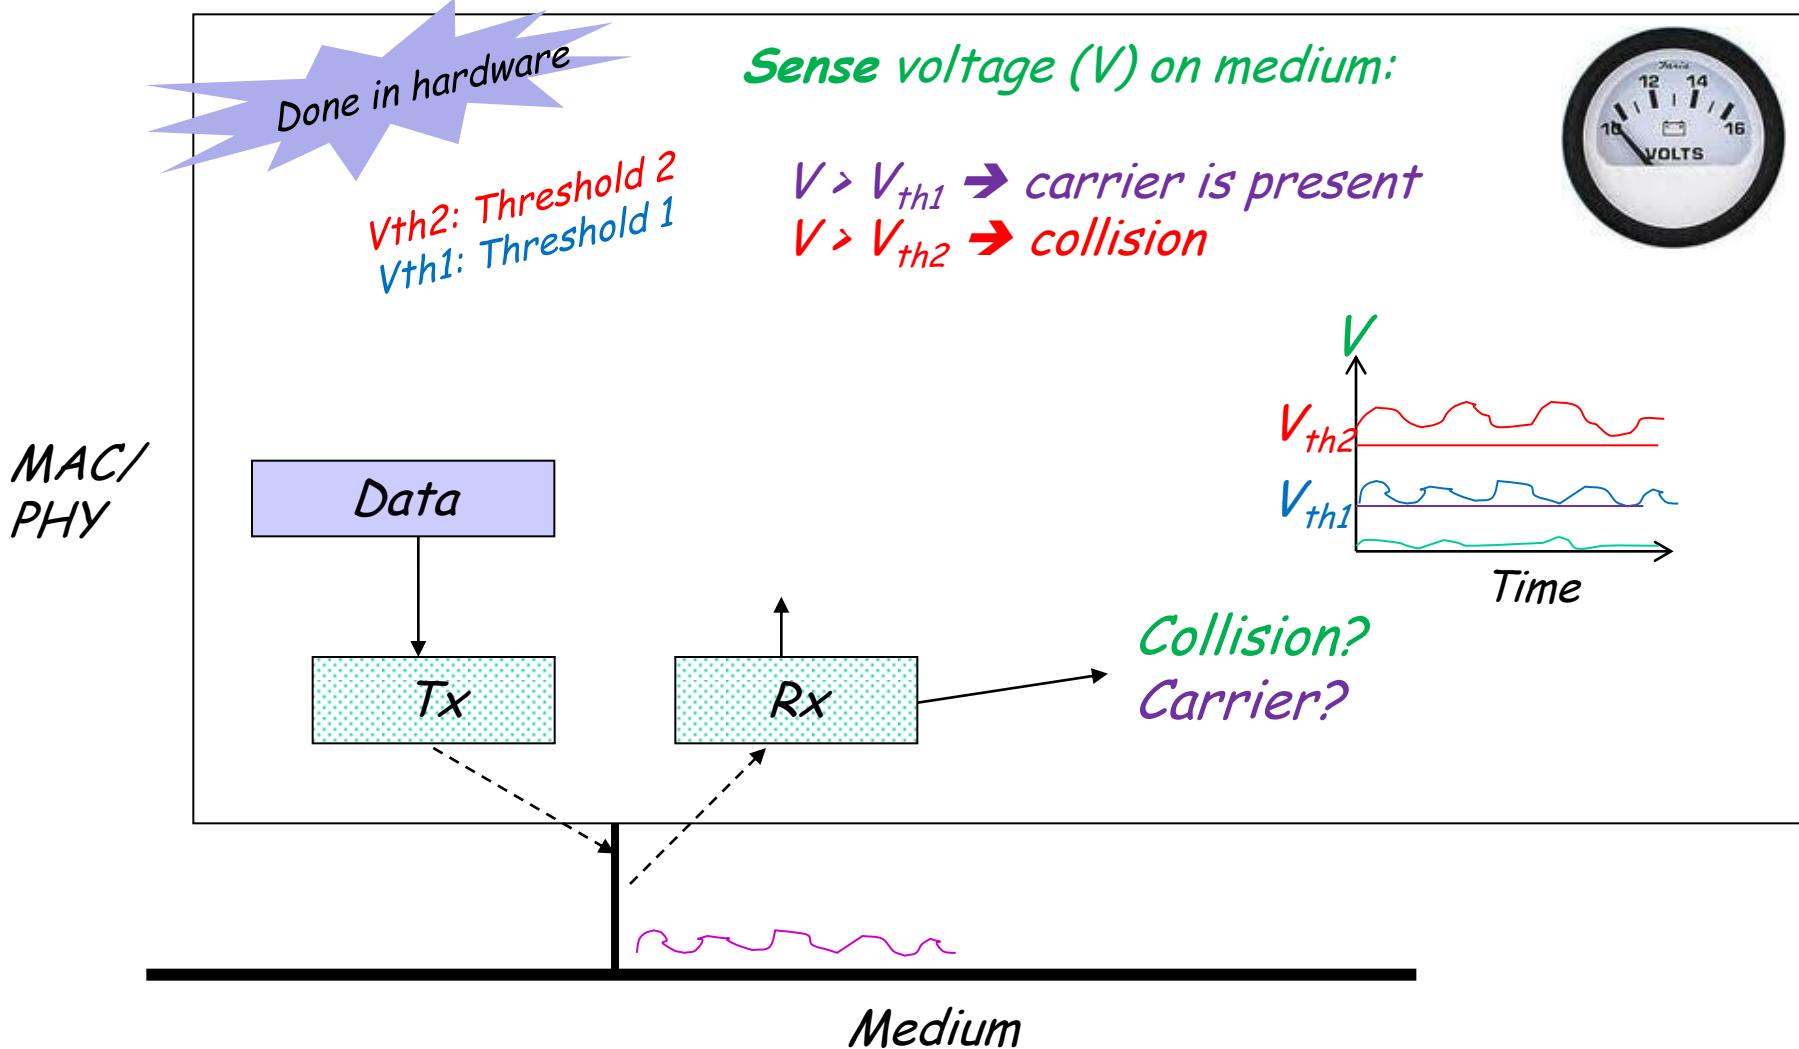
- ❖ If channel sensed idle: transmit entire frame
- ❖ If channel sensed busy, defer transmission

A cartoon illustration of a character with a single hair tuft, a small nose, and a wide, toothy grin showing many sharp, jagged teeth. A large, light blue speech bubble originates from the character's mouth, containing the text "Don't interrupt others..." in a cursive font.

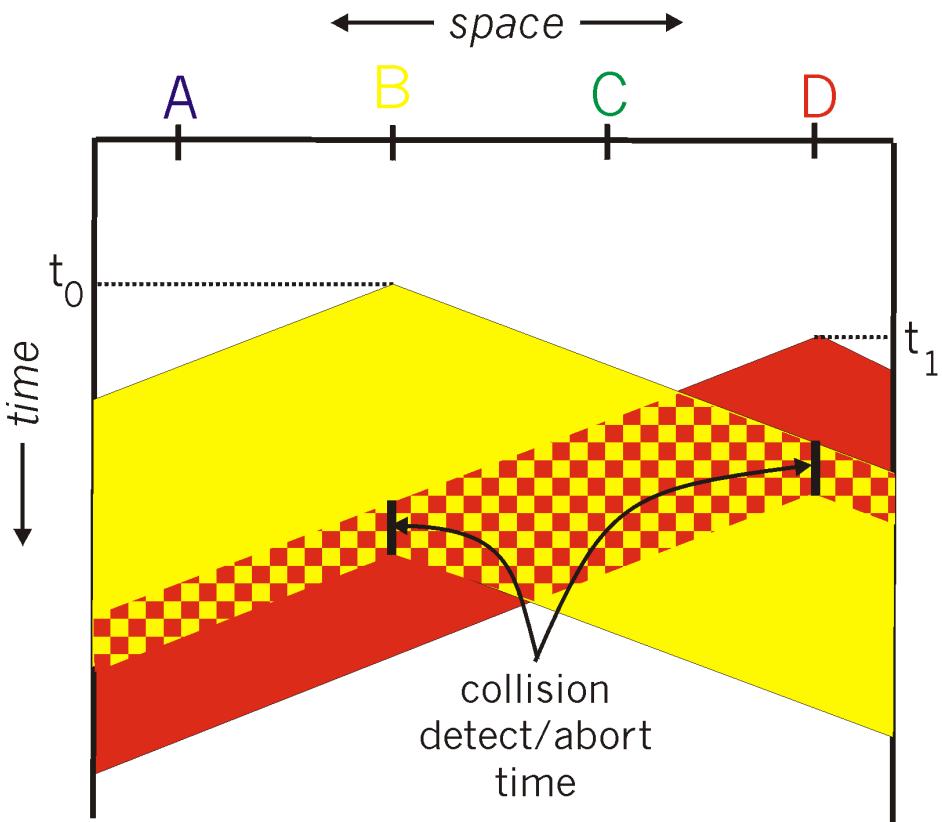
Don't interrupt others...



CSMA/CD Concepts of Carrier Sense and Collision Detection



CSMA/CD collision detection



*Signal from all nodes
must reach all others*



*Therefore,
collision occurring at
a receiving node is
detected by the frame's
sender.... after some delay..*

CSMA/CD (Collision Detection)

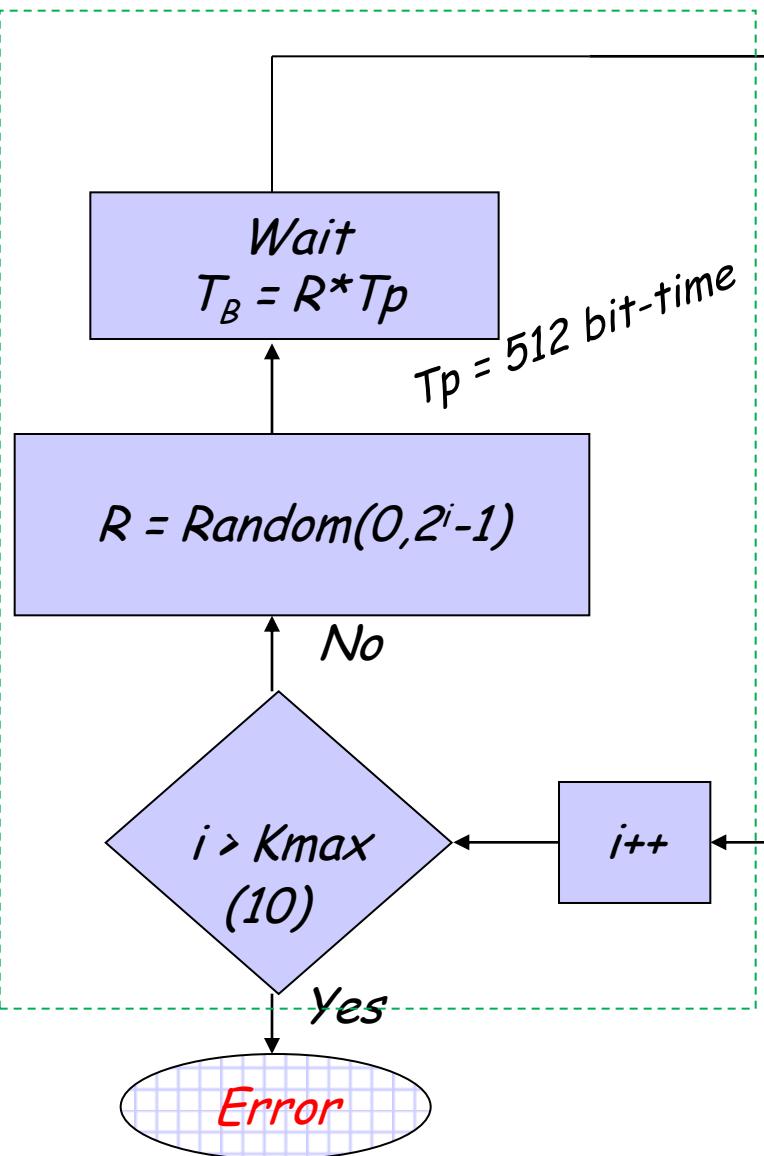
CSMA/CD: Carrier Sense Multiple Access with CD

- Tx attempts multiple times (multiple access)
- Collisions are *detected* within short time
- A colliding transmission is aborted ← don't keep wasting channel resource

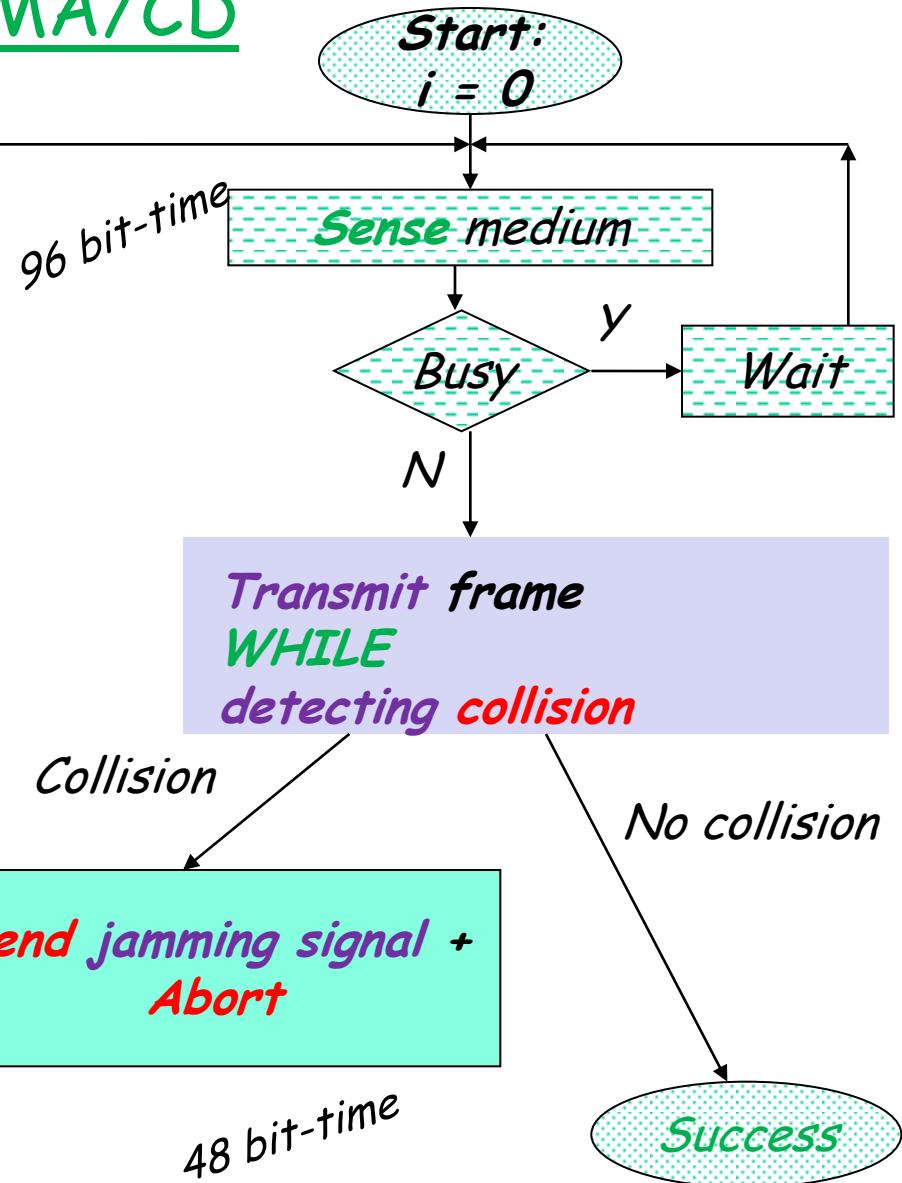
❖ Collision detection:

- Easy in wired LANs: explained before ...
- Not possible in wireless LANs: will be explained ...

Exponential backoff



CSMA/CD



CSMA/CD

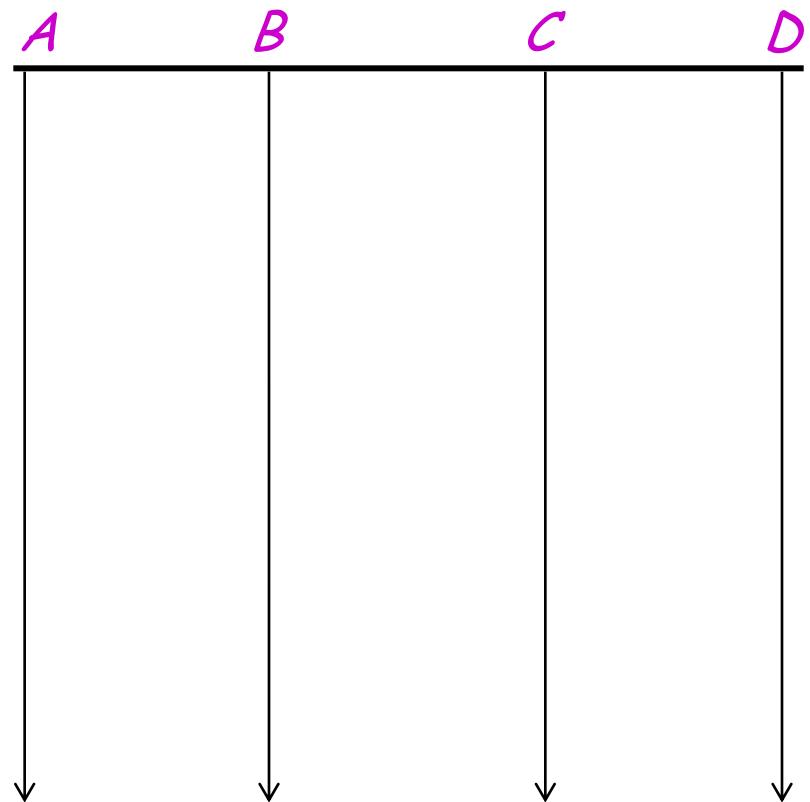
- Medium sensing is done for **96** bit-times.
- Jamming signal length is **48** bits. Jamming signal **creates enough energy** on the medium for collision detection.
- T_p is equated with **512** bit-times.
- “ i ” saturates at **10**.



I started transmitting
after sensing no-carrier....
Why did it collide with another frame.....?
Are they not listening...



Past final
exam





I have learned CSMA/CD...
What kinds of questions should I
expect on the exams

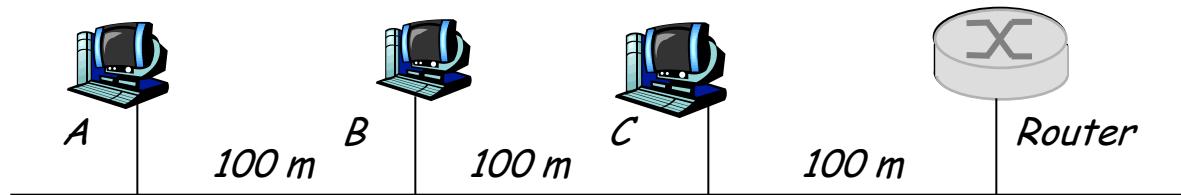


F12 Mid-term



Computers A, B, and C and one router are connected to an Ethernet bus 100 metres apart to make a LAN using the CSMA/CD protocol.

Assume that signal propagates on the Ethernet bus at a speed of 2×10^8 m/sec, and all nodes can transmit data at the rate of 100 Mbps.



- After finding the bus to be idle for a little while, A and the router start transmitting their frames exactly at the same time. How many bits of data can node A transmit before detecting collision?
- What is the maximum time gap between the start of transmission and detection of collision by A?
- What is the length of the smallest frame that A can transmit while knowing whether or not it collided with a transmission from the router?



Another question



Justify the use of
binary exponential backoff
in MAC protocols.

Mid-term F'12



Compare the slotted Aloha protocol
with the CSMA/CD protocol by
identifying five comparison criteria.

Criteria	Slotted Aloha	CSMA/CD

"Taking Turns" MAC protocols

channel partitioning MAC protocols:

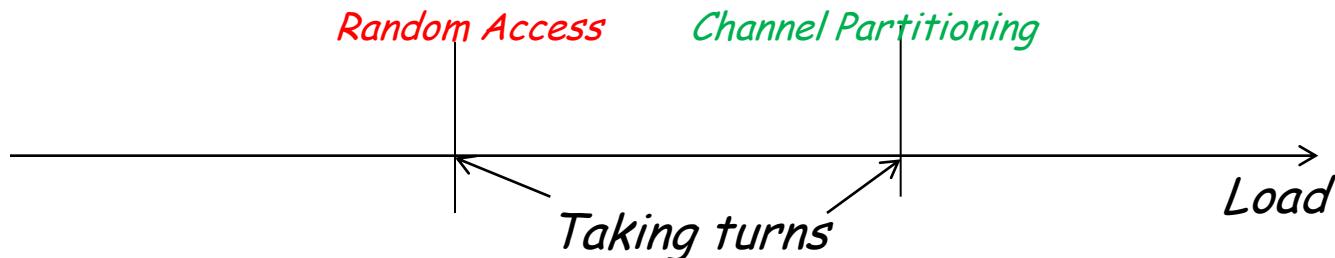
- share channel *efficiently* and *fairly* at high load
- *inefficient* at low load: $1/N$ bandwidth allocated even if only 1 active node!

random access MAC protocols:

- *efficient* at low load: single node can fully utilize channel
- Much *collision overhead* at high load

"taking turns" protocols

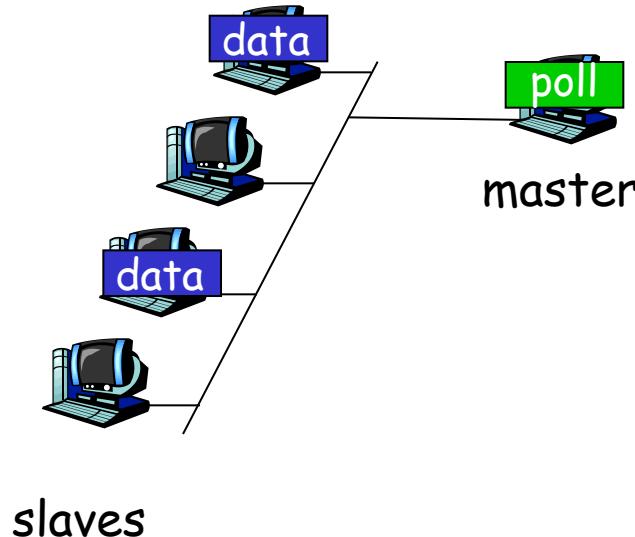
look for best of both worlds!



"Taking Turns" MAC protocols

Polling:

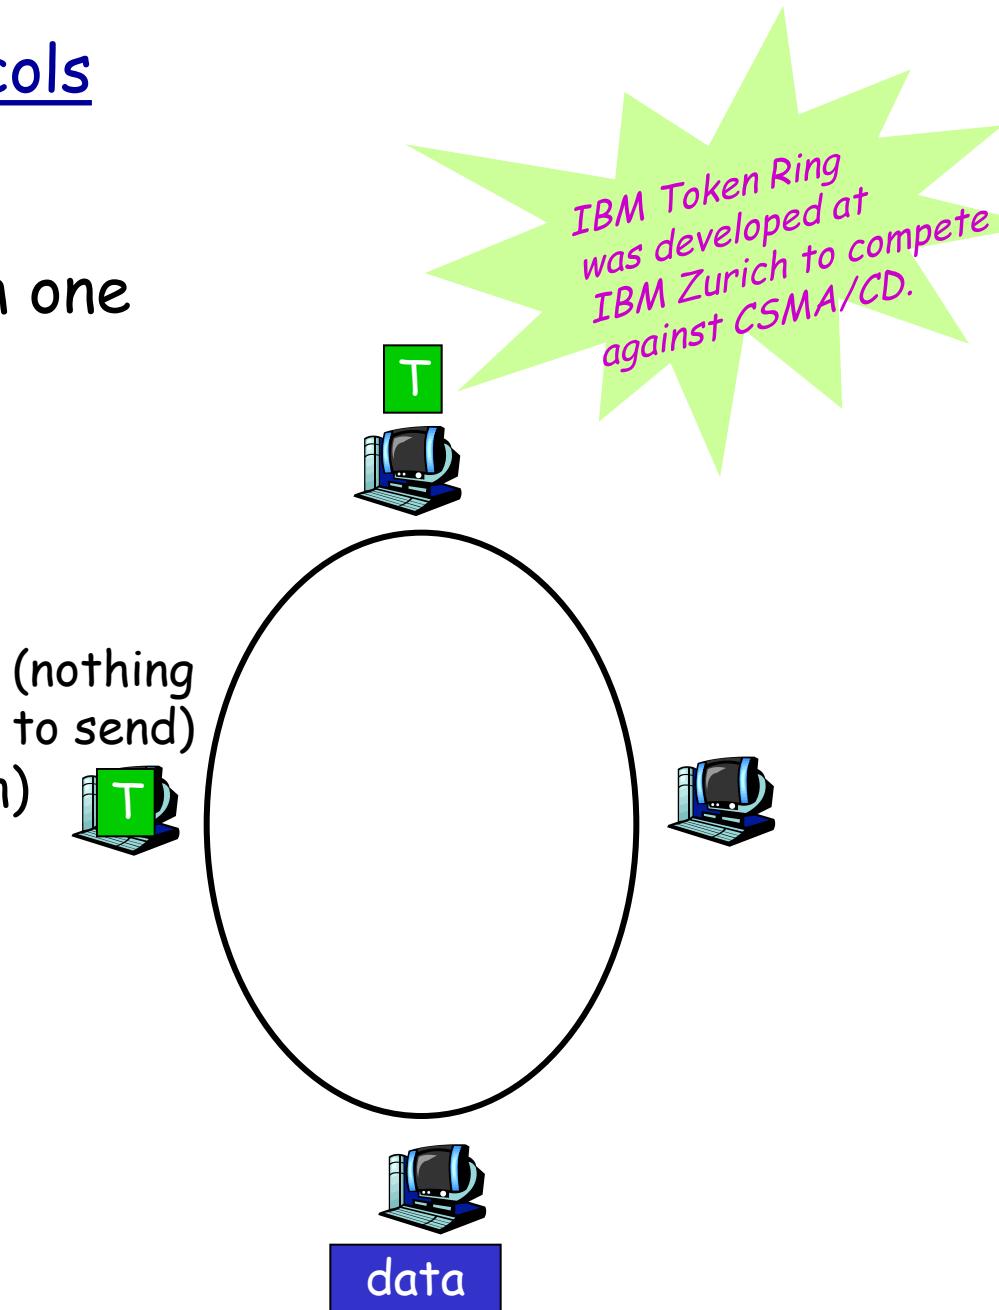
- ❖ master node
 - “invites” slave nodes to transmit in turn
- ❖ typically used with “dumb” slave devices
- ❖ concerns:
 - polling overhead
 - latency
 - single point of failure (master)



"Taking Turns" MAC protocols

Token passing:

- ❖ control **token** passed from one node to next sequentially.
- ❖ token message
- ❖ **concerns:**
 - token overhead
 - latency
 - single point of failure (token)





Where are all those
different kinds of MAC
protocols used



Ch. Partitioning MAC

- Cellphone networks
- Bluetooth

Random access MAC

- LAN (Local Area Network)
- WLAN (Wireless LAN)

Taking turns

- Bluetooth

Link Layer

5.1 Introduction and services

5.2 Error detection and correction

5.3 Multiple access protocols

5.4 Link-Layer Addressing

5.5 Ethernet

5.6 Link-layer switches

5.7 PPP

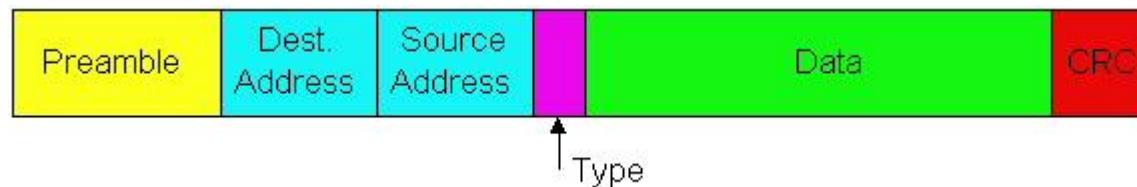
5.8 Link virtualization:
MPLS

5.9 A day in the life of a web request

3.4 Reliable data transfer

Ethernet Frame Structure

Sending adapter encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**

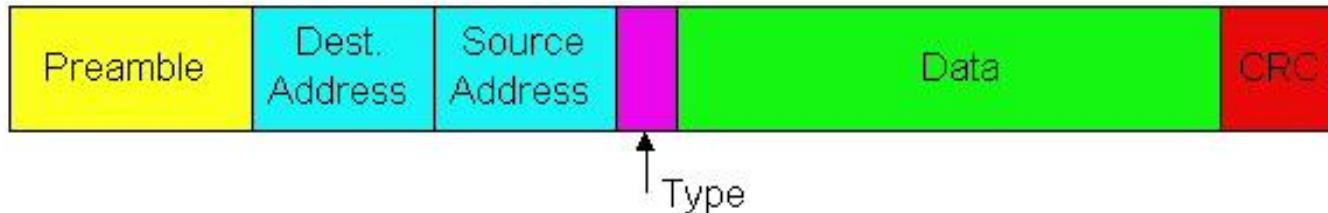


Preamble:

- ❖ 7 bytes with pattern **10101010** followed by one byte with pattern **10101011**
- ❖ used to synchronize receiver, sender clock rates

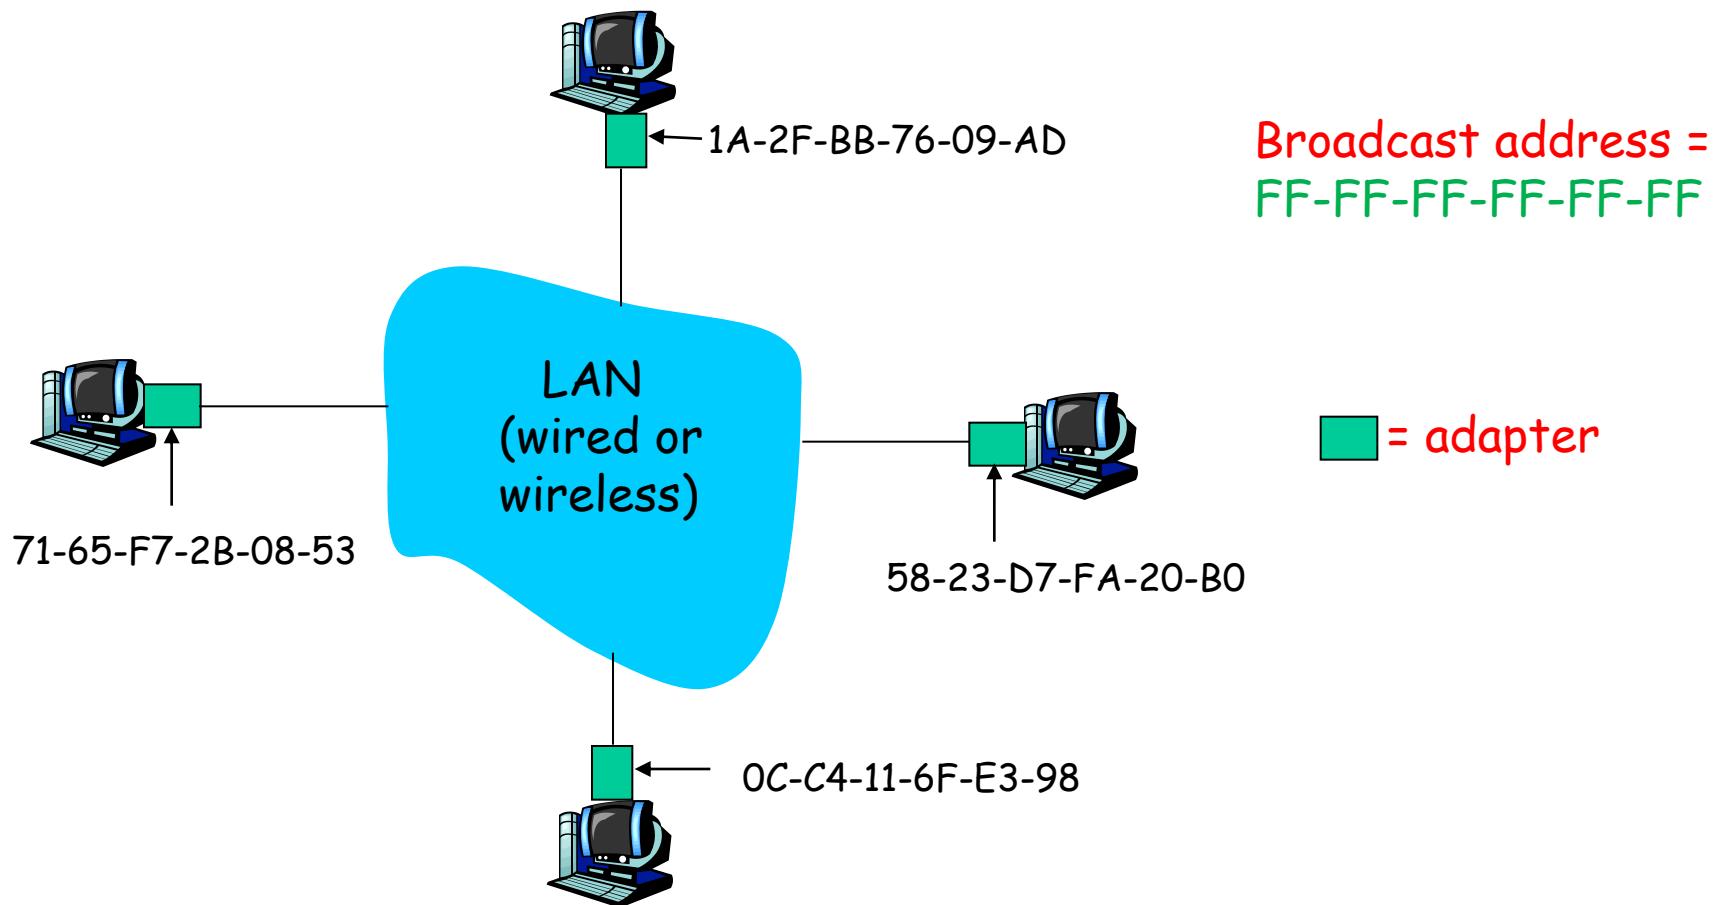
Ethernet Frame Structure (more)

- ❖ **Addresses:** 6 bytes
 - if adapter receives frame with matching destination address, or with broadcast address (e.g. ARP packet), it **passes data** in frame to upper layer protocol
 - otherwise, adapter **discards** frame
- ❖ **Type:** indicates higher layer protocol (mostly IP but others possible, e.g., Novell IPX, AppleTalk)
- ❖ **CRC:** checked at receiver, if error is detected, frame is dropped



LAN Addresses and ARP

Each adapter on LAN has unique LAN address

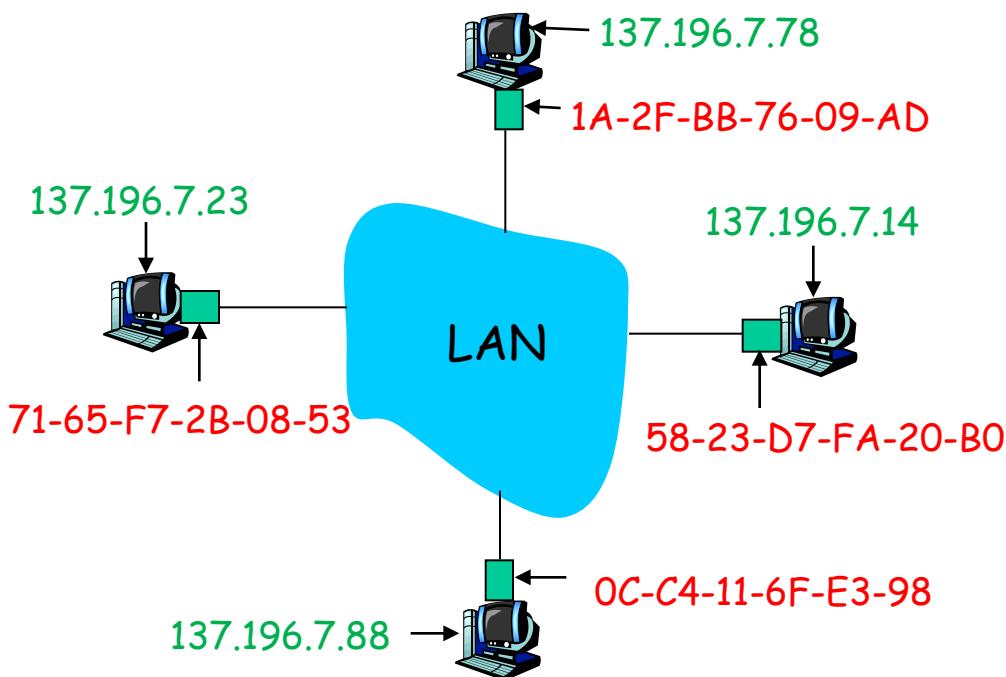


ARP: Address Resolution Protocol

Question:

how to determine B's **MAC address** knowing B's **IP address**?

- ❖ Each IP node (host, router) on LAN has **ARP table**
- ❖ **ARP table: IP/MAC addr. mappings for LAN nodes**



< **IP address; MAC address; TTL** >

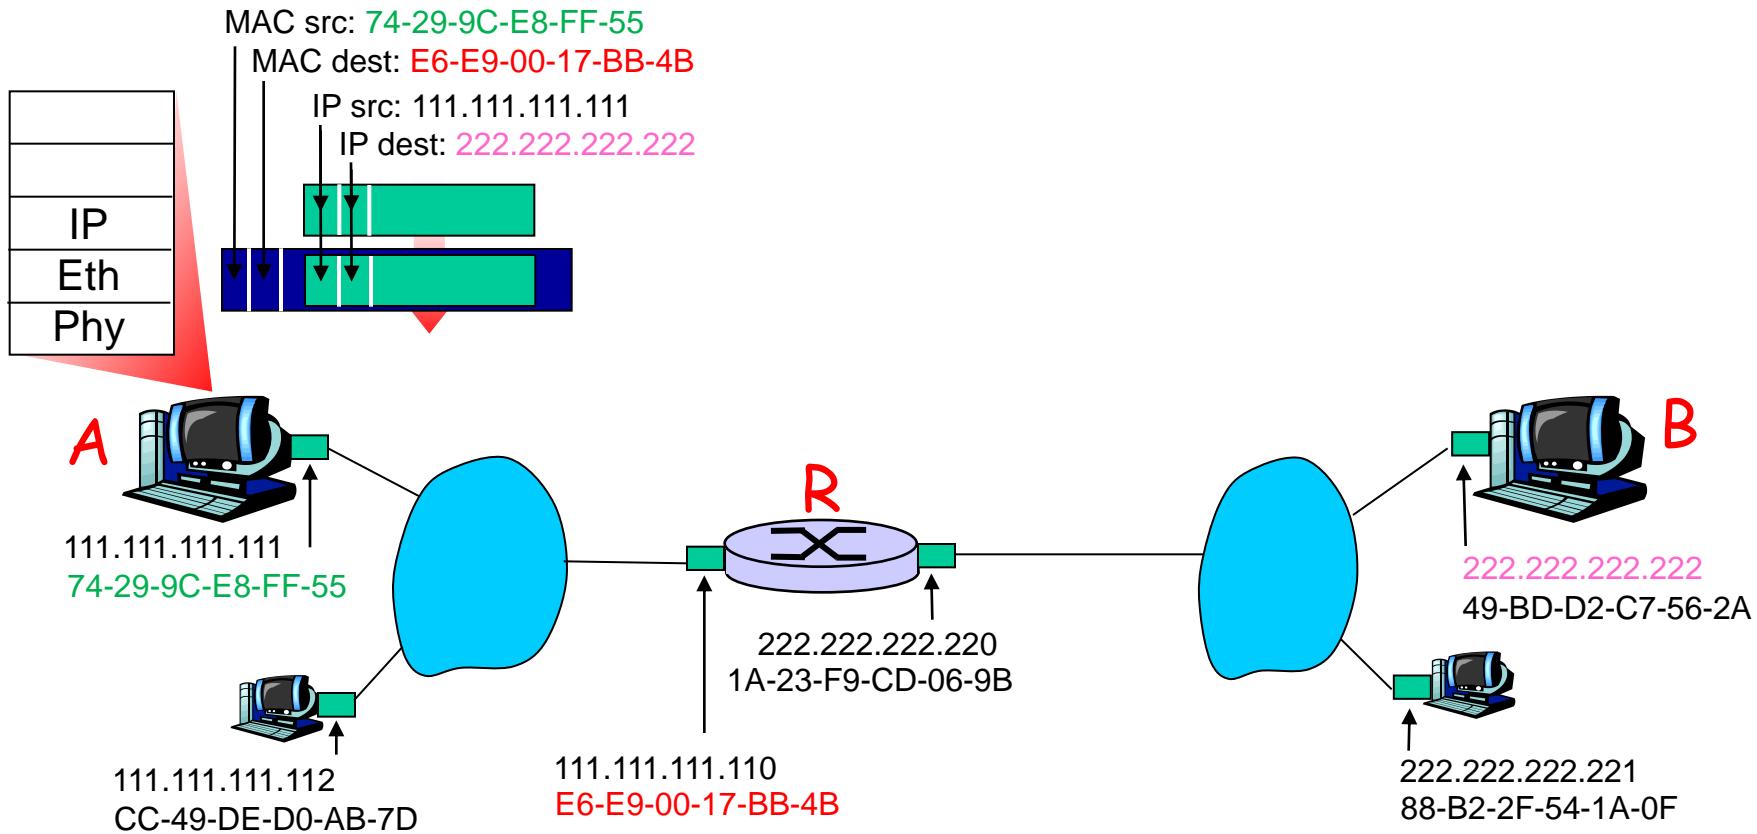
- **TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)**

ARP protocol: Same LAN (network)

- ❖ A wants to send datagram to B, and B's MAC address **not** in A's ARP table.
- ❖ A broadcasts ARP query pkt containing B's IP addr
 - dest MAC address = FF-FF-FF-FF-FF-FF
 - all machines on LAN receive ARP query
- ❖ B receives ARP query, replies to A with its (B's) MAC addr
 - frame sent to A's MAC address (unicast)
- ❖ A caches (saves) IP-to-MAC addr pair **in its ARP table** until this becomes old
- ❖ ARP is "plug-and-play":
 - nodes create their ARP tables without intervention from net administrator

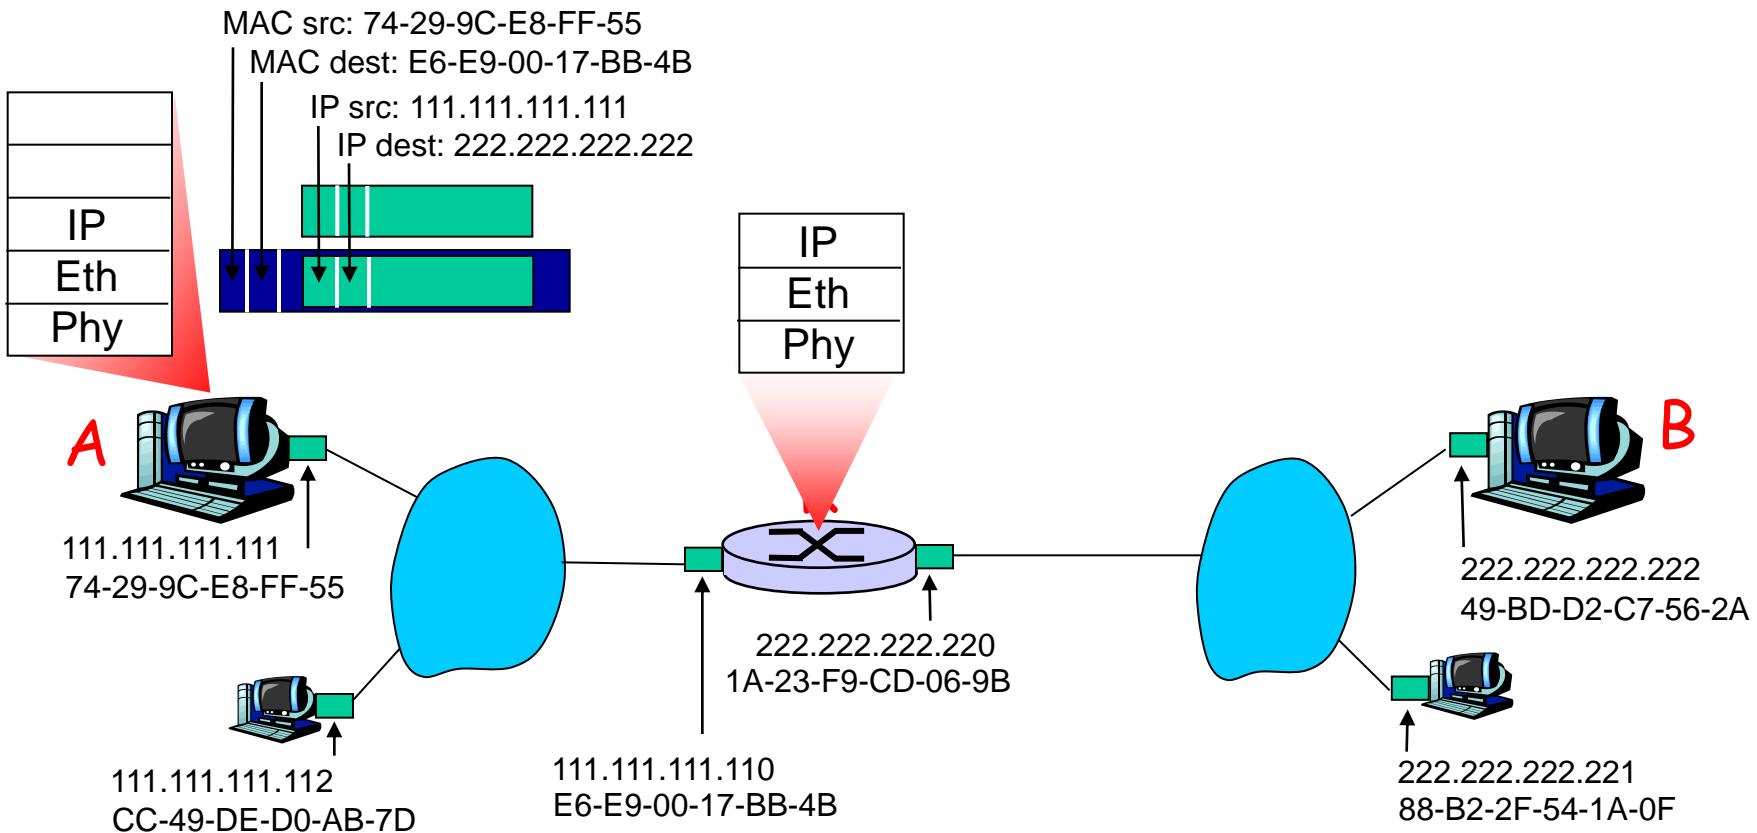
Addressing: routing to another LAN

- ❖ A creates IP datagram with IP source A, destination B
- ❖ A creates link-layer frame with R's MAC address as dest, frame contains A-to-B IP datagram



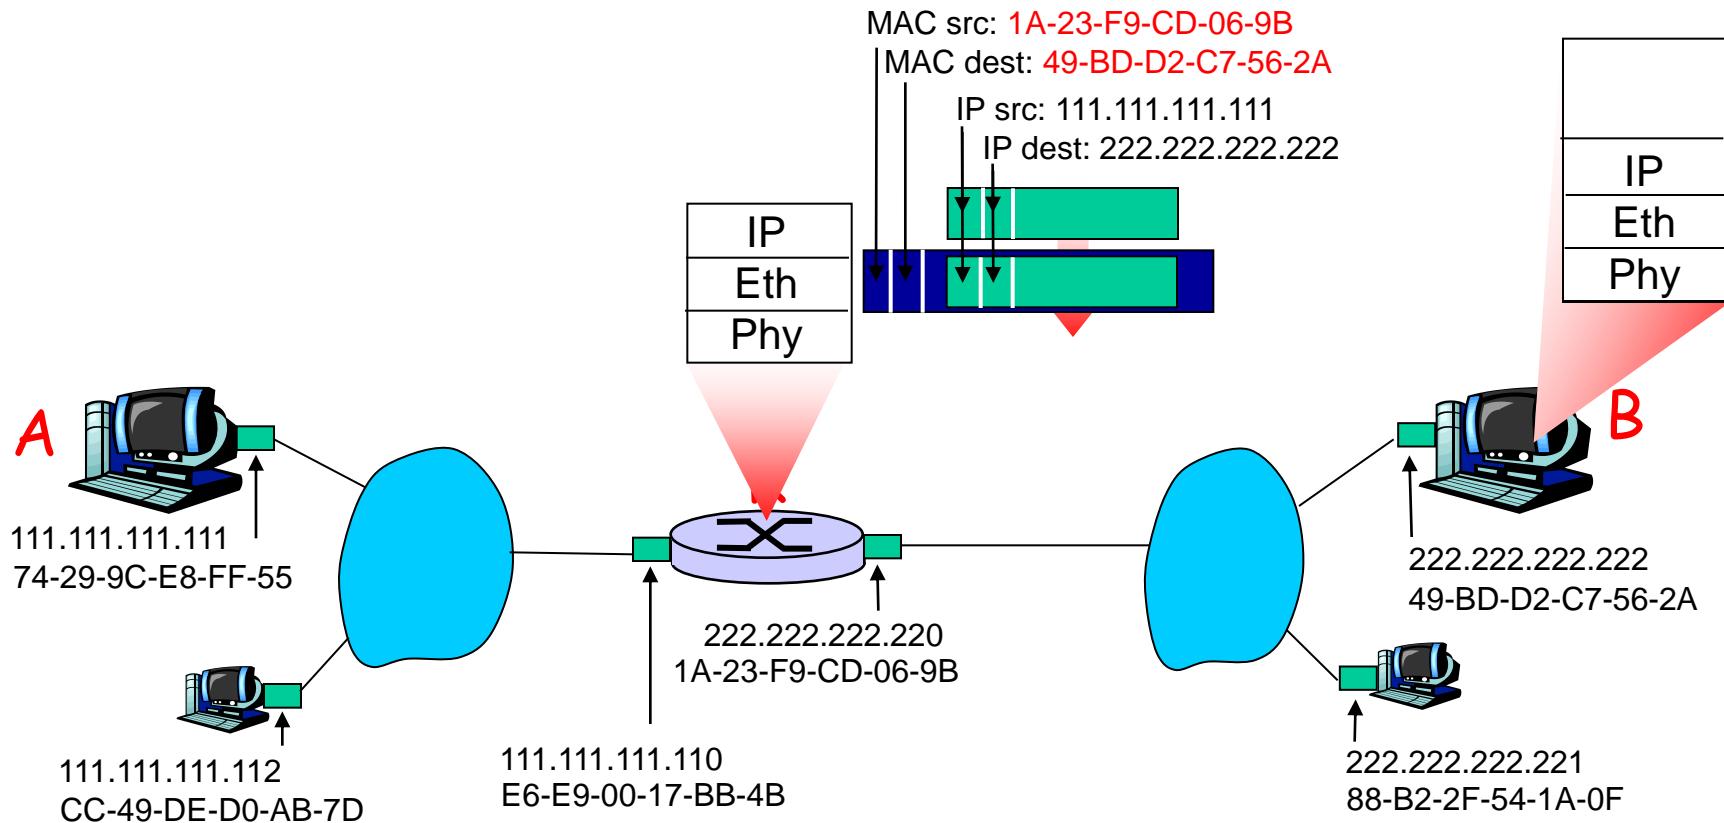
Addressing: routing to another LAN

- ❖ frame sent from A to R
- ❖ frame received at R, datagram removed, passed up to IP



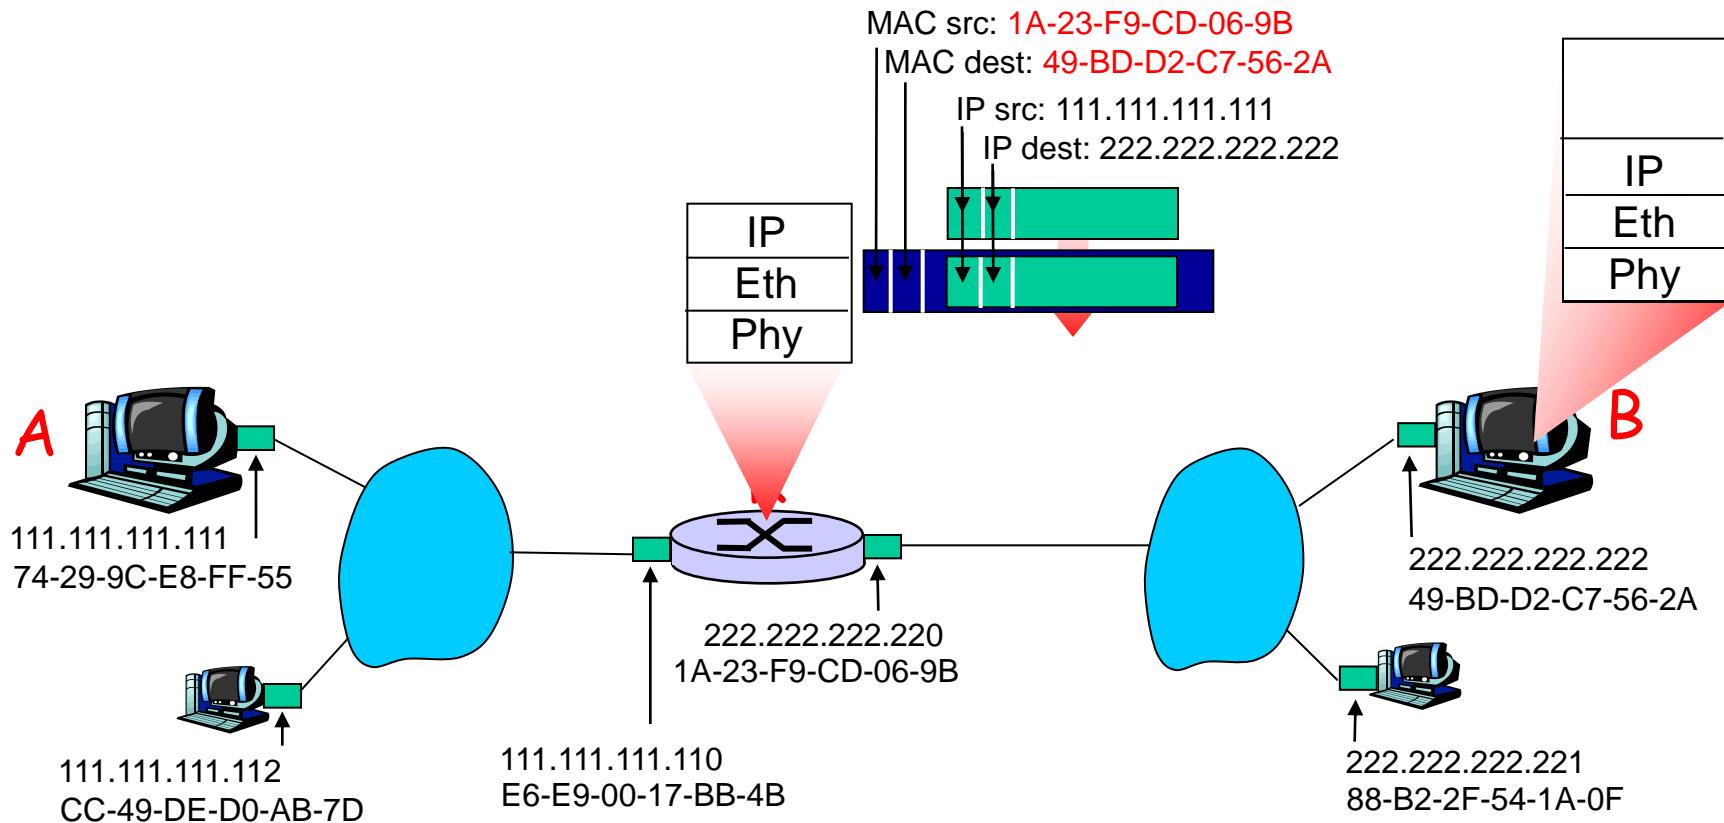
Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



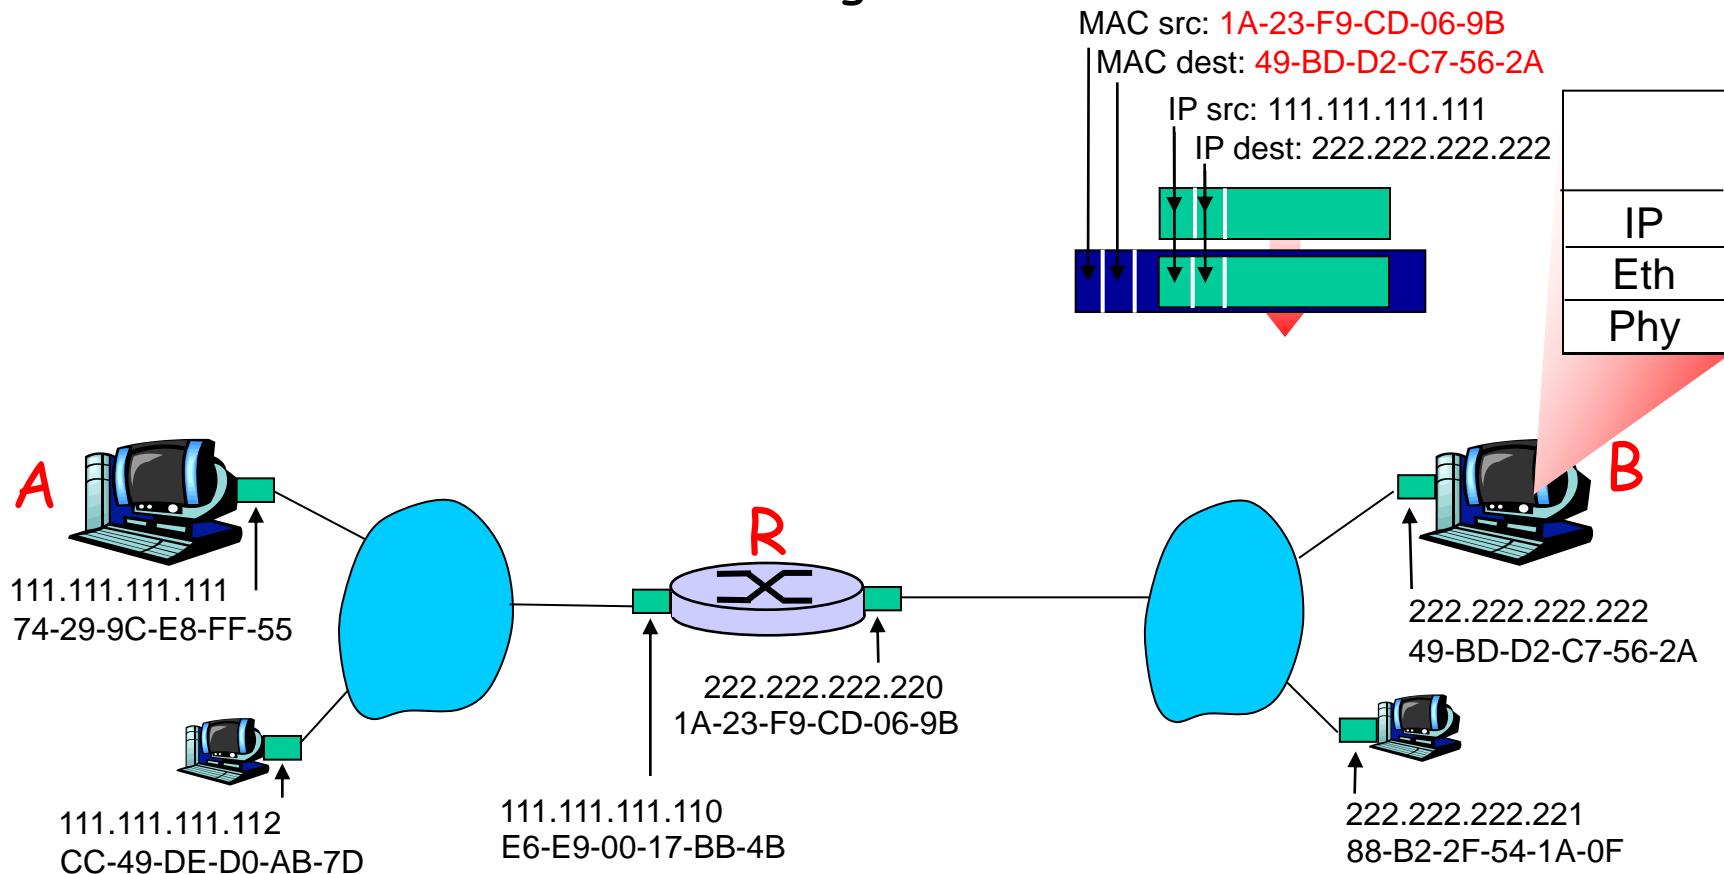
Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



Addressing: routing to another LAN

- ❖ R forwards datagram with IP source A, destination B
- ❖ R creates link-layer frame with B's MAC address as dest, frame contains A-to-B IP datagram



Link Layer

5.1 Introduction and services

5.2 Error detection and correction

5.3 Multiple access protocols

5.4 Link-layer Addressing

5.5 Ethernet

5.6 Link-layer switches, LANs, VLANs

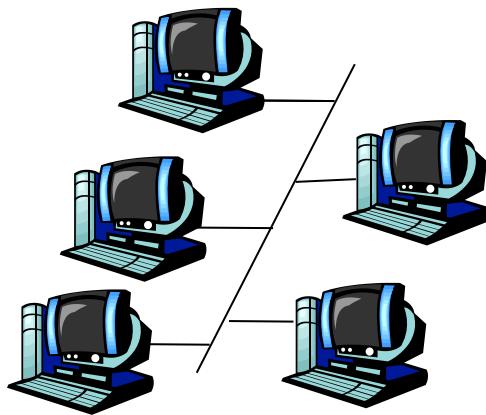
5.7 PPP

5.8 Link virtualization:
MPLS

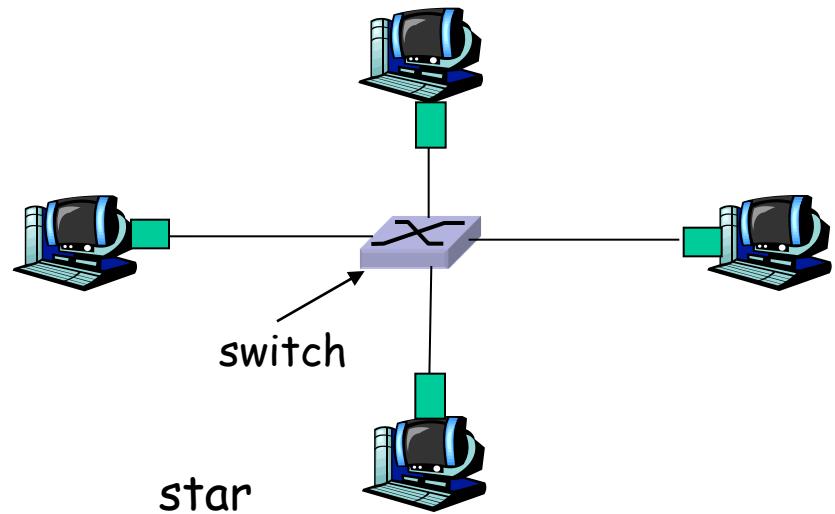
5.9 A day in the life of a web request

LAN with star topology

- ❖ bus topology popular through mid 90s
 - all nodes in same collision domain (can collide with each other)
- ❖ today: star topology prevails
 - active switch in center
 - each "spoke" runs a (separate) Ethernet protocol (nodes do not collide with each other)



bus: coaxial cable

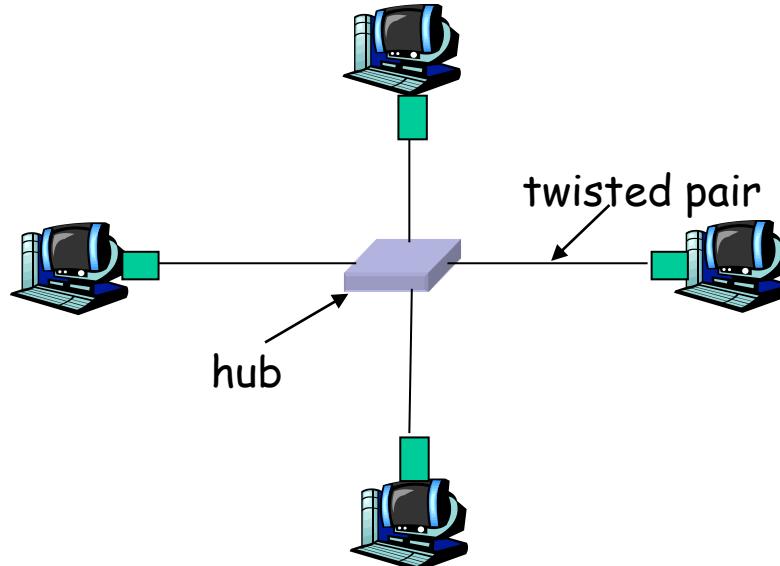


star

Hubs

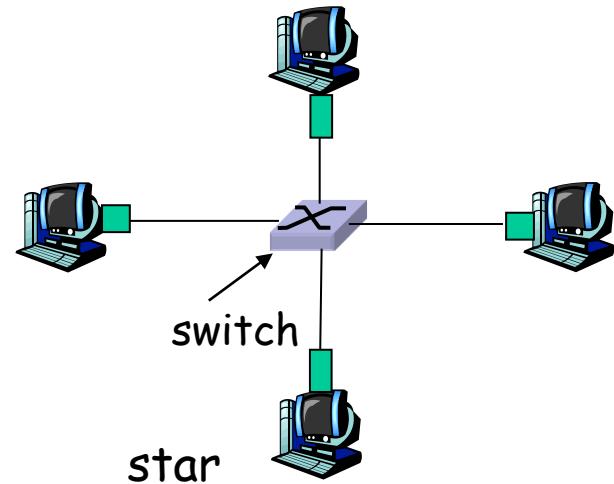
... physical-layer ("dumb") repeaters:

- bits coming in on one link go out *all* other links at same rate
- all nodes connected to hub can collide with one another
- no frame buffering
- no CSMA/CD at hub: host NICs detect collisions



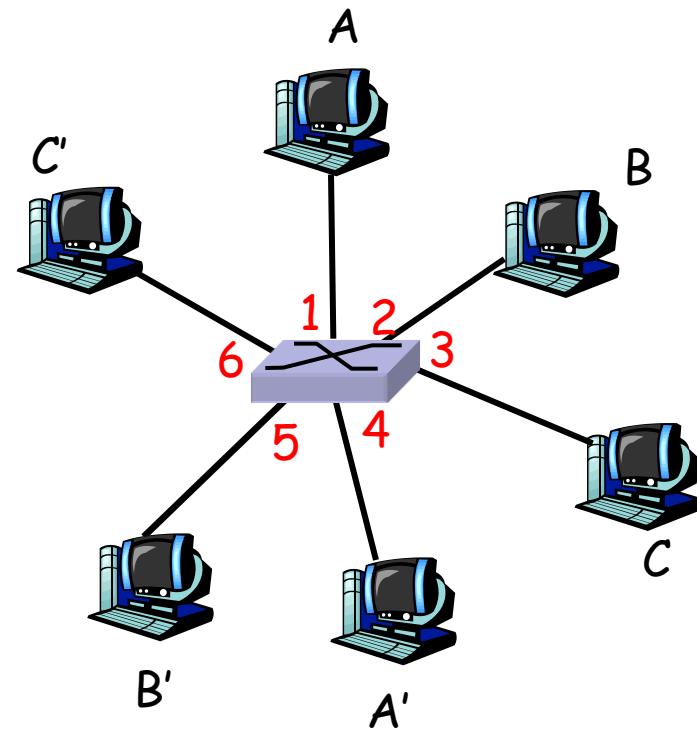
Switch

- ❖ smarter than hubs, take active role
 - Store and forward frames
 - Examine incoming frame's MAC address, selectively forward frame to one-or-more outgoing links when frame is to be forwarded on segment, using CSMA/CD to access segment
- ❖ transparent hosts are unaware of presence of switches
- ❖ plug-and-play, self-learning switches do not need to be configured



Switch: allows *multiple* simultaneous transmissions

- ❖ hosts have dedicated, direct connection to switch
- ❖ switches buffer packets
- ❖ Ethernet protocol used on each link
 - each link is its own collision domain
- ❖ *switching: A-to-A' and B-to-B' simultaneously, without collisions*
 - not possible with dumb hub

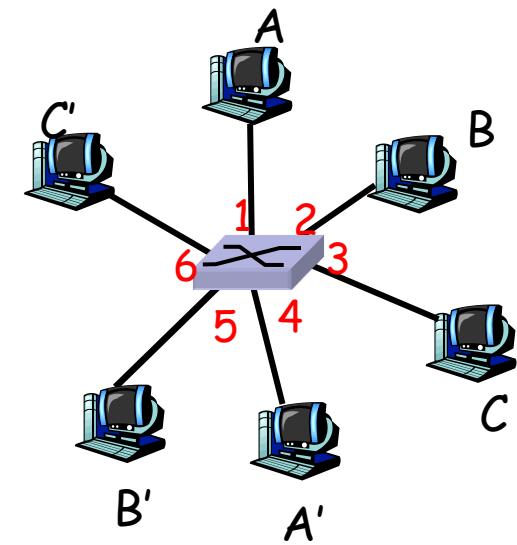


*switch with six interfaces
(1,2,3,4,5,6)*

Switch Table

- ❖ Q: how does switch know that A' reachable via interface 4, B' reachable via interface 5?
- ❖ A: each switch has a switch table

MAC addr of host	Interface #	Time To Live (TTL)



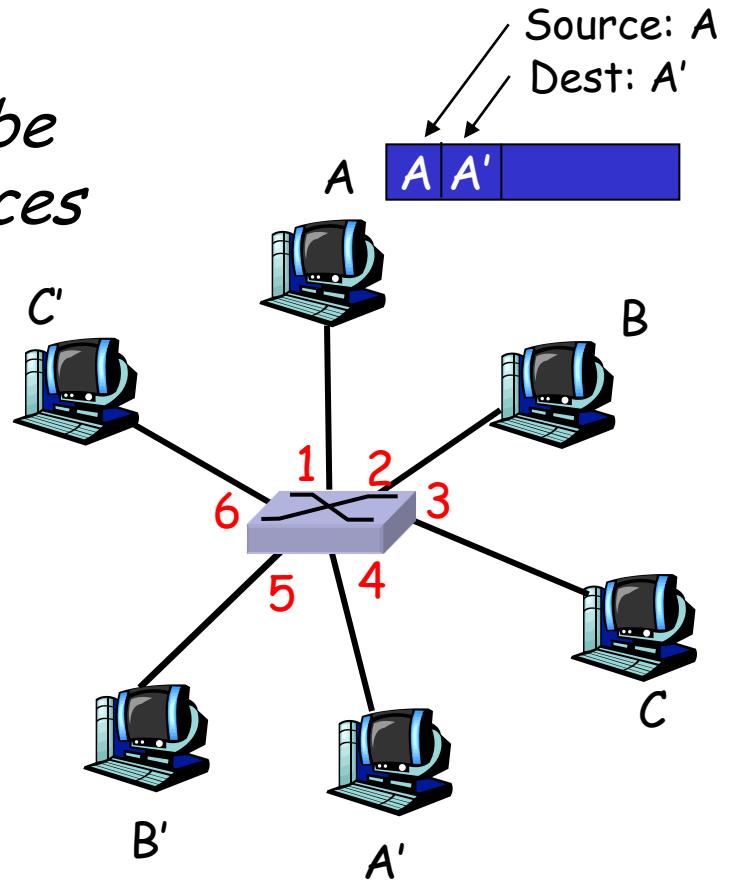
switch with six interfaces
(1,2,3,4,5,6)

- ❖ Q: How are entries created and maintained in switch table?

Switch: self-learning

- ❖ Switch **learns** which hosts can be reached through which interfaces

- ❖ When frame received
 - Switch “learns” location of sender: incoming LAN segment
 - Records sender/location pair in switch table



MAC addr	interface	TTL
A	1	60

*Switch table
(initially empty)*

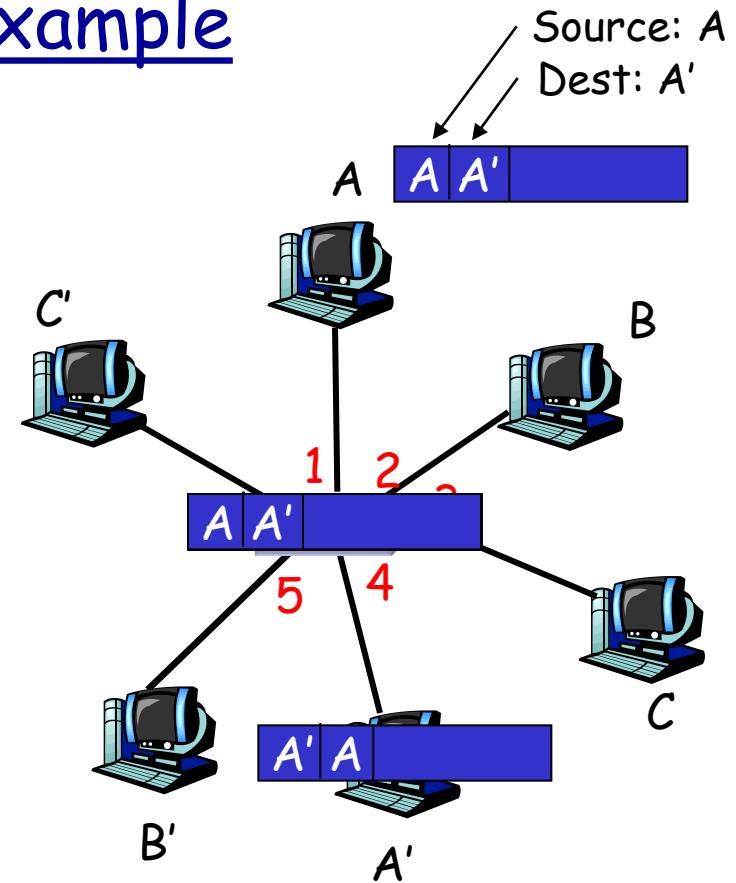
Switch: frame filtering/forwarding

When frame received:

1. record link associated with sending host
 2. index switch table using MAC dest address
 3. if entry found for destination then {
 if dest on segment from which frame arrived
 then drop the frame
 else forward the frame on interface indicated
}
- else flood
- forward on all but the interface
on which the frame arrived

Self-learning, forwarding: example

- ❖ frame destination unknown: *flood*
- ❖ destination A location known: *selective send*

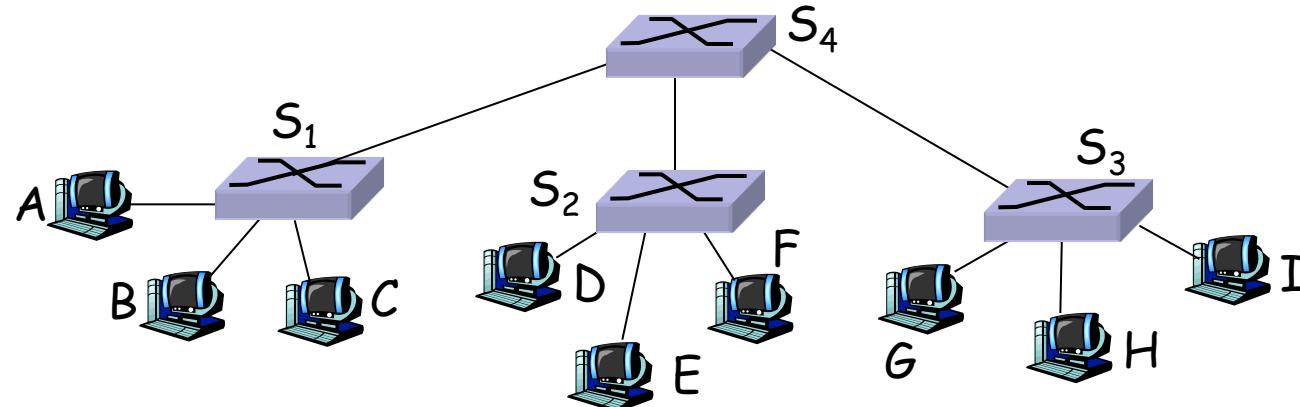


MAC addr	interface	TTL
A	1	60
A'	4	60

*Switch table
(initially empty)*

Interconnecting switches

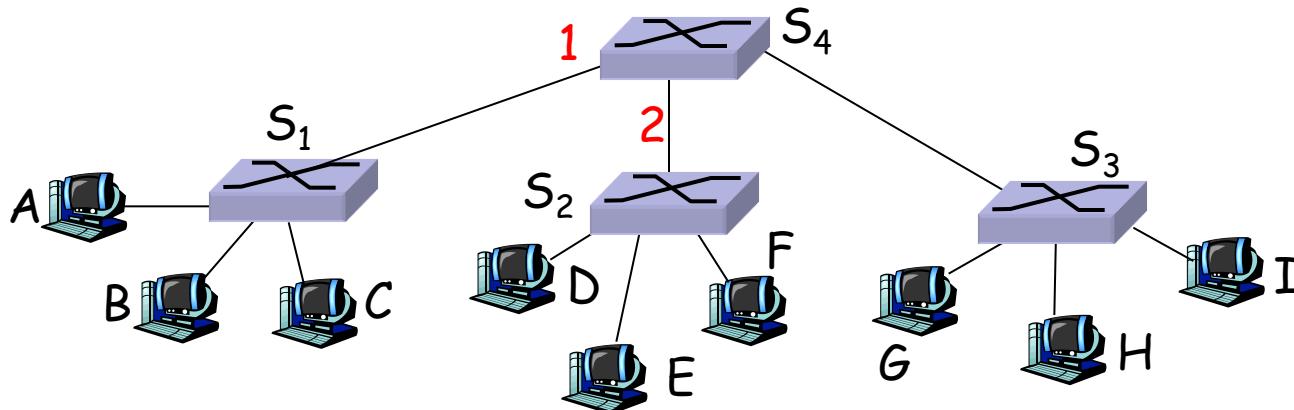
- ❖ switches can be connected together



- ❖ **Q:** sending from A to G - how does S₁ know to forward frame destined to G via S₄ and S₃?
- ❖ **A:** self learning! (works exactly the same as in single-switch case!)

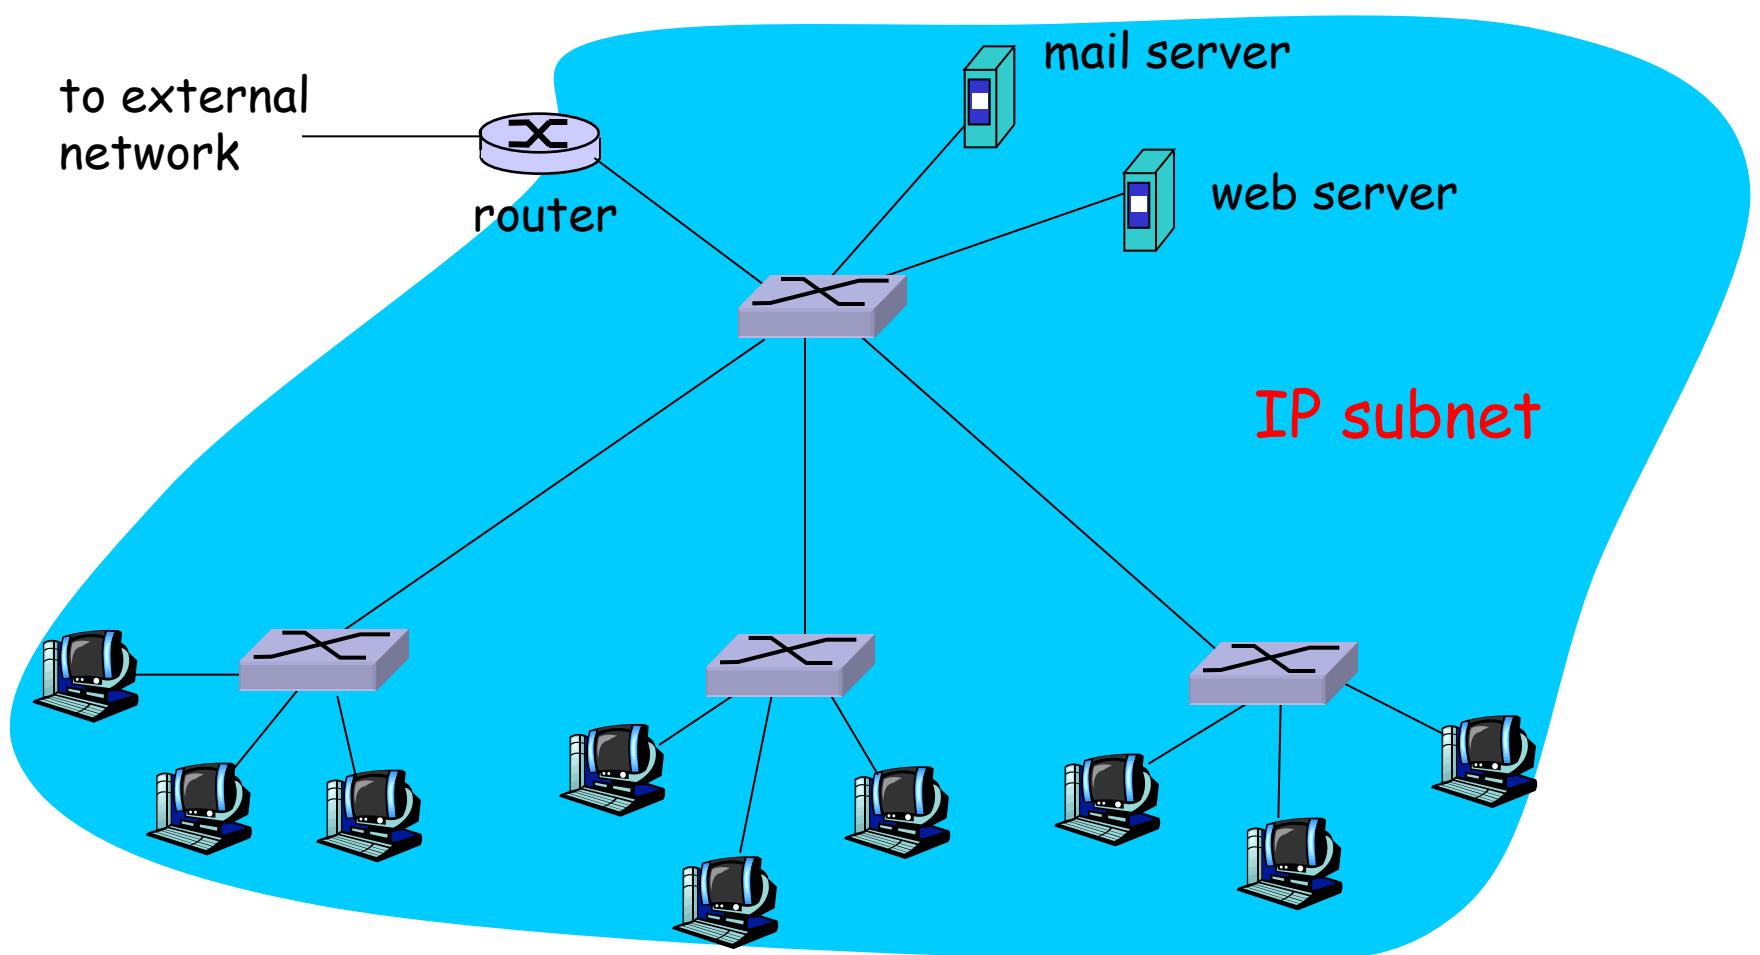
Self-learning multi-switch example

Suppose C sends frame to I, I responds to C



- ❖ **Q:** show switch tables and packet forwarding in S_1 , S_2 , S_3 , S_4

Institutional network



Fall 2012 Final Exam Question

[3] Consider the institutional local area network shown in Figure 1.

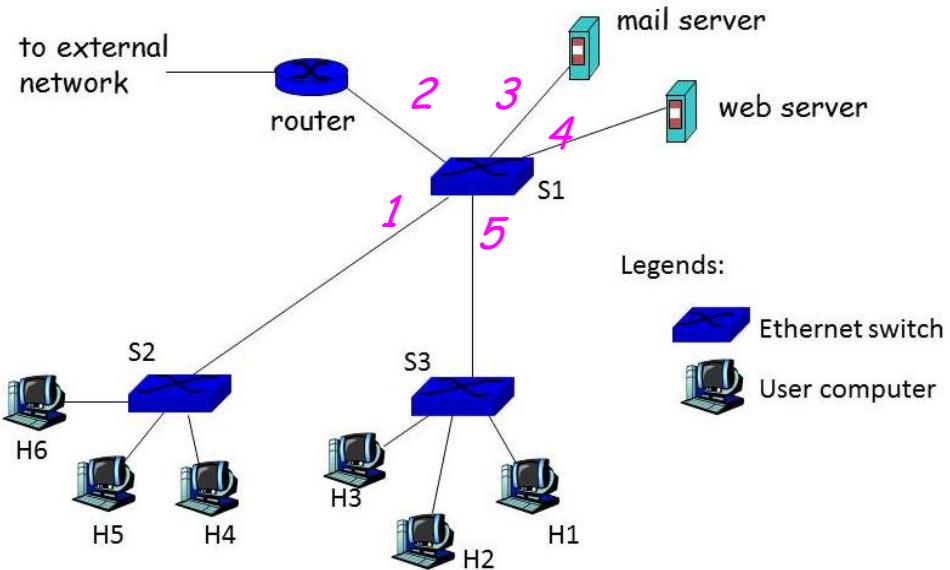


Figure 1: An institutional network

All the computing entities, except H1 and H4, are **fully functional and sending and receiving data packets**.

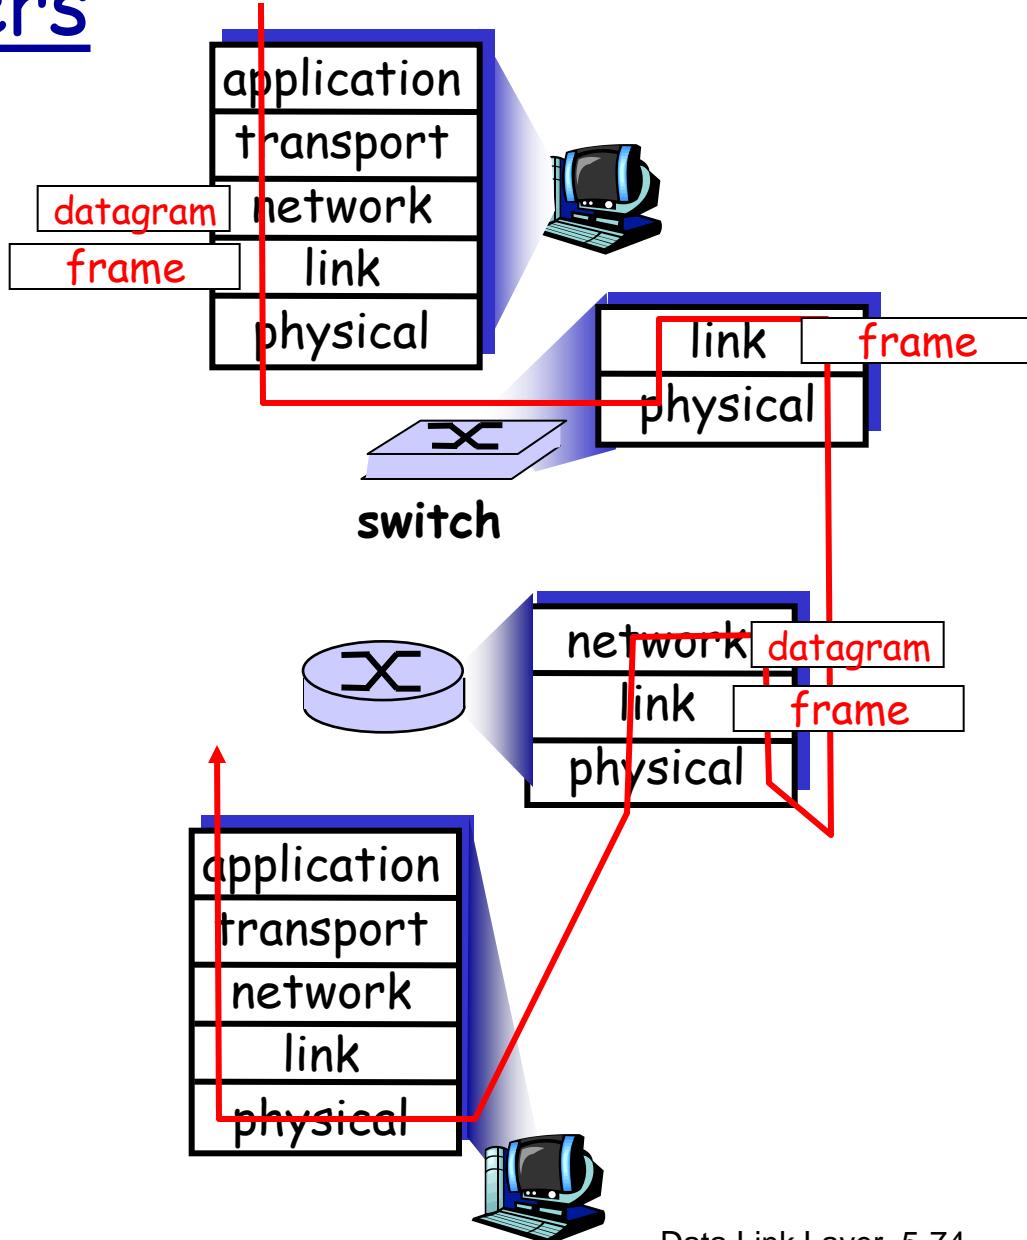
H1 is **switched off**, whereas H4's Ethernet interface is **broken**.

Complete the following switching table at Ethernet switch S1.

An arbitrary number of rows have been given. Add more rows, if you need them. Make assumptions if there is a need.

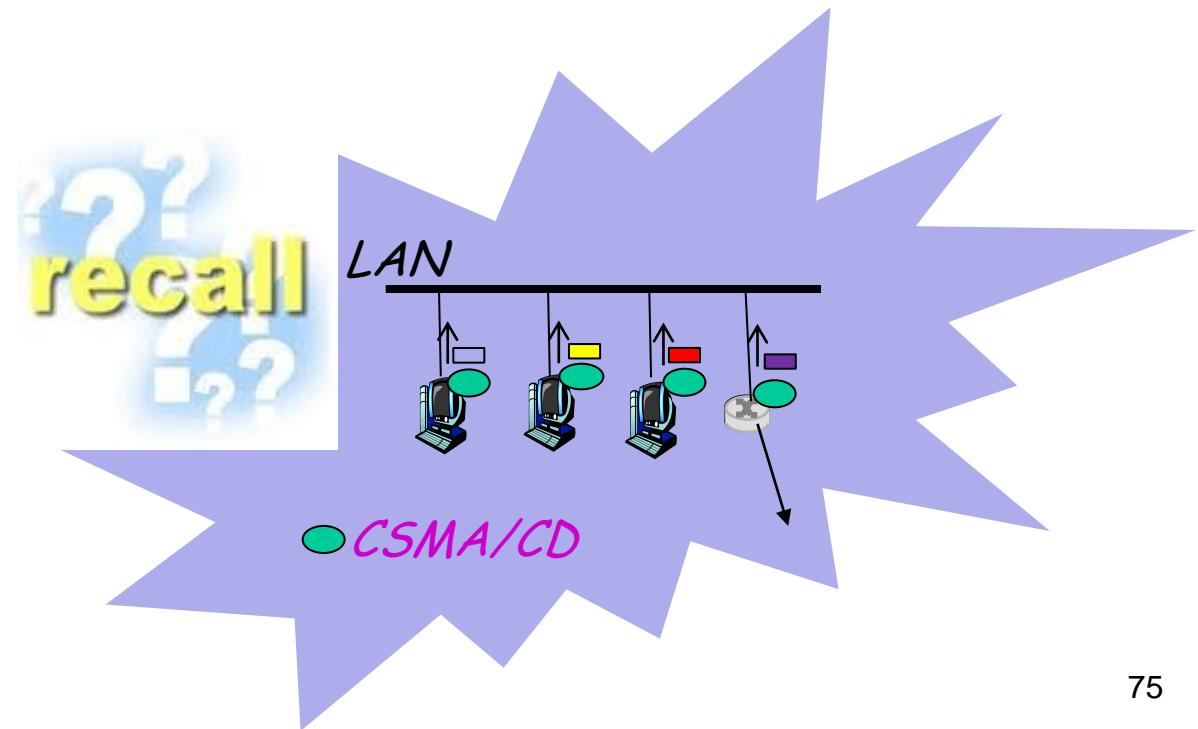
Switches vs. Routers

- ❖ both store-and-forward devices
 - routers: network-layer devices (examine network-layer headers)
 - switches are link-layer devices (examine link-layer headers)
- ❖ routers maintain routing tables, implement routing algorithms
- ❖ switches maintain switch tables, implement filtering, learning algorithms



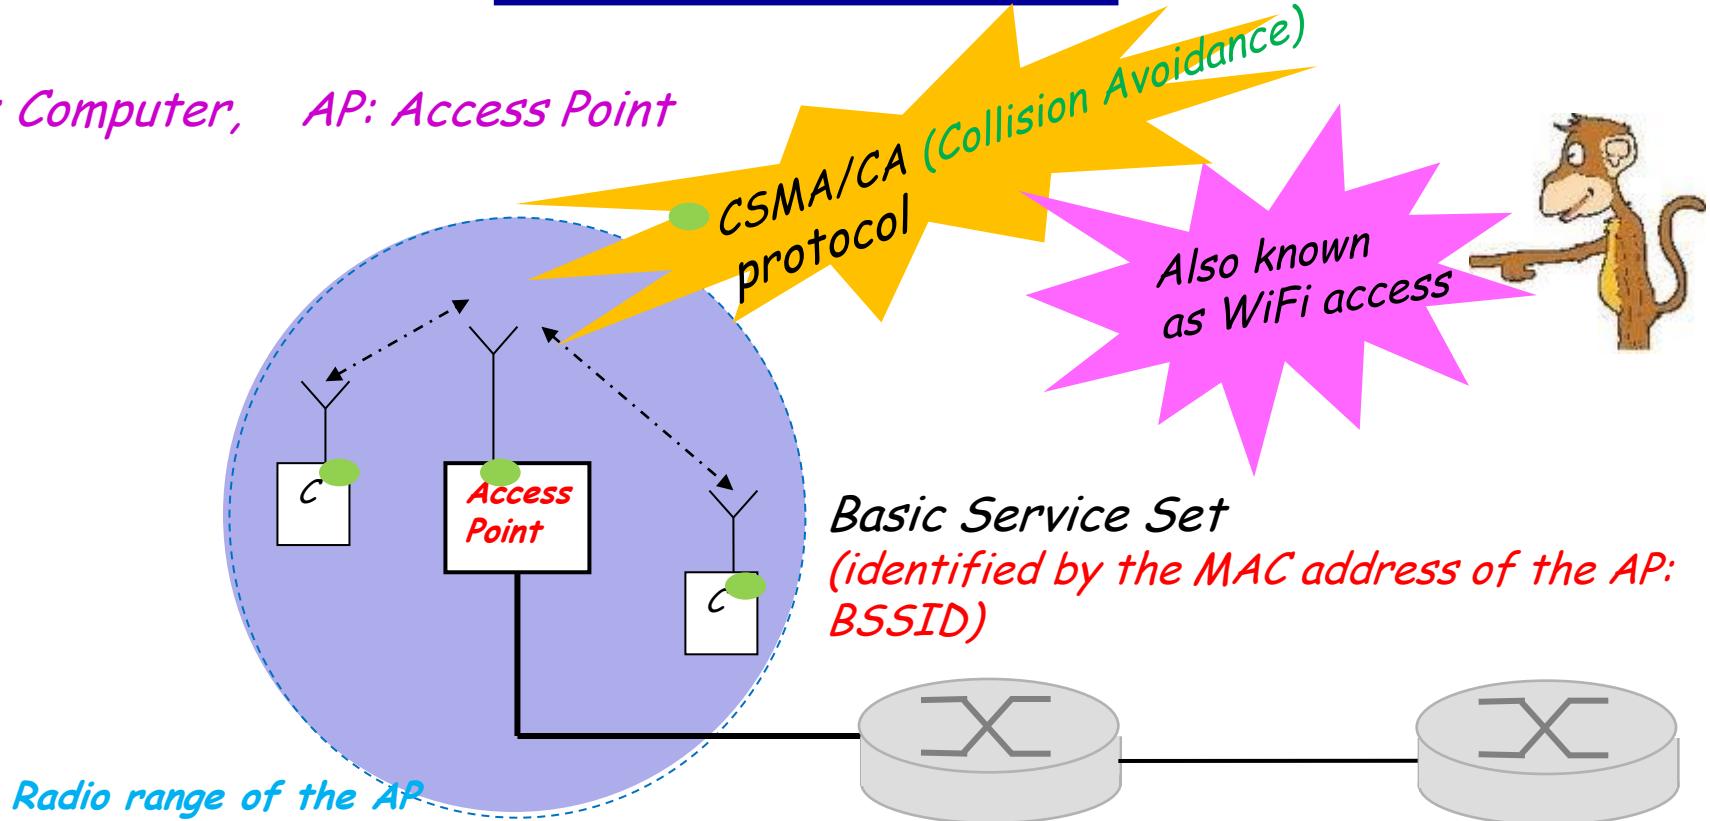
Wireless LAN

Standard: IEEE 802.11/a/b/g



WLAN: Basic idea

C: Computer, AP: Access Point



Independent BSS (IBSS)= BSS - AP

Extended Service Set (ESS): A collection of BSS connected by a Distribution System

Example: The UW WiFi network is one ESS.

IEEE 802.11/a/b/g/n Family

IEEE	Signal Transmission	Frequency Band	Rate (Mbps)
802.11	DSSS	2.4 GHz	1 and 2
	FHSS	2.4 GHz	1 and 2
802.11a	OFDM	5 GHz	6--54
802.11b	DSSS	2.4 GHz	5.5 and 11
802.11g	OFDM	2.4 GHz	22 and 54
802.11n	OFDM	2.4/5 GHz	72 and <u>150</u>
802.11ac (Draft/Nov. 2011)	OFDM	5 GHz	<u>6.9 Gbps</u>

DSSS: Direct Sequence Spread Spectrum

FHSS: Frequency Hopping Spread Spectrum

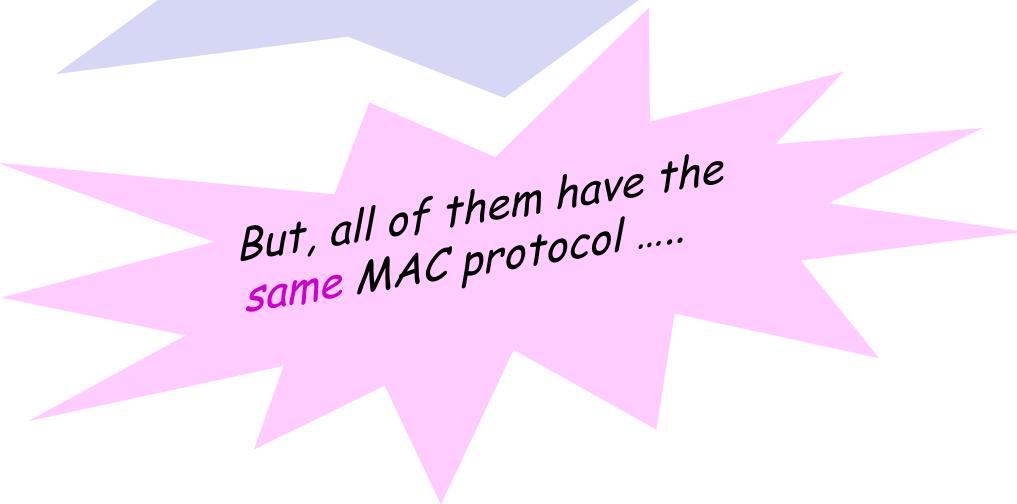
OFDM: Orthogonal Frequency Division Multiplexing



Are there
different MAC protocols in
IEEE 802.11/a/b/g/n standards



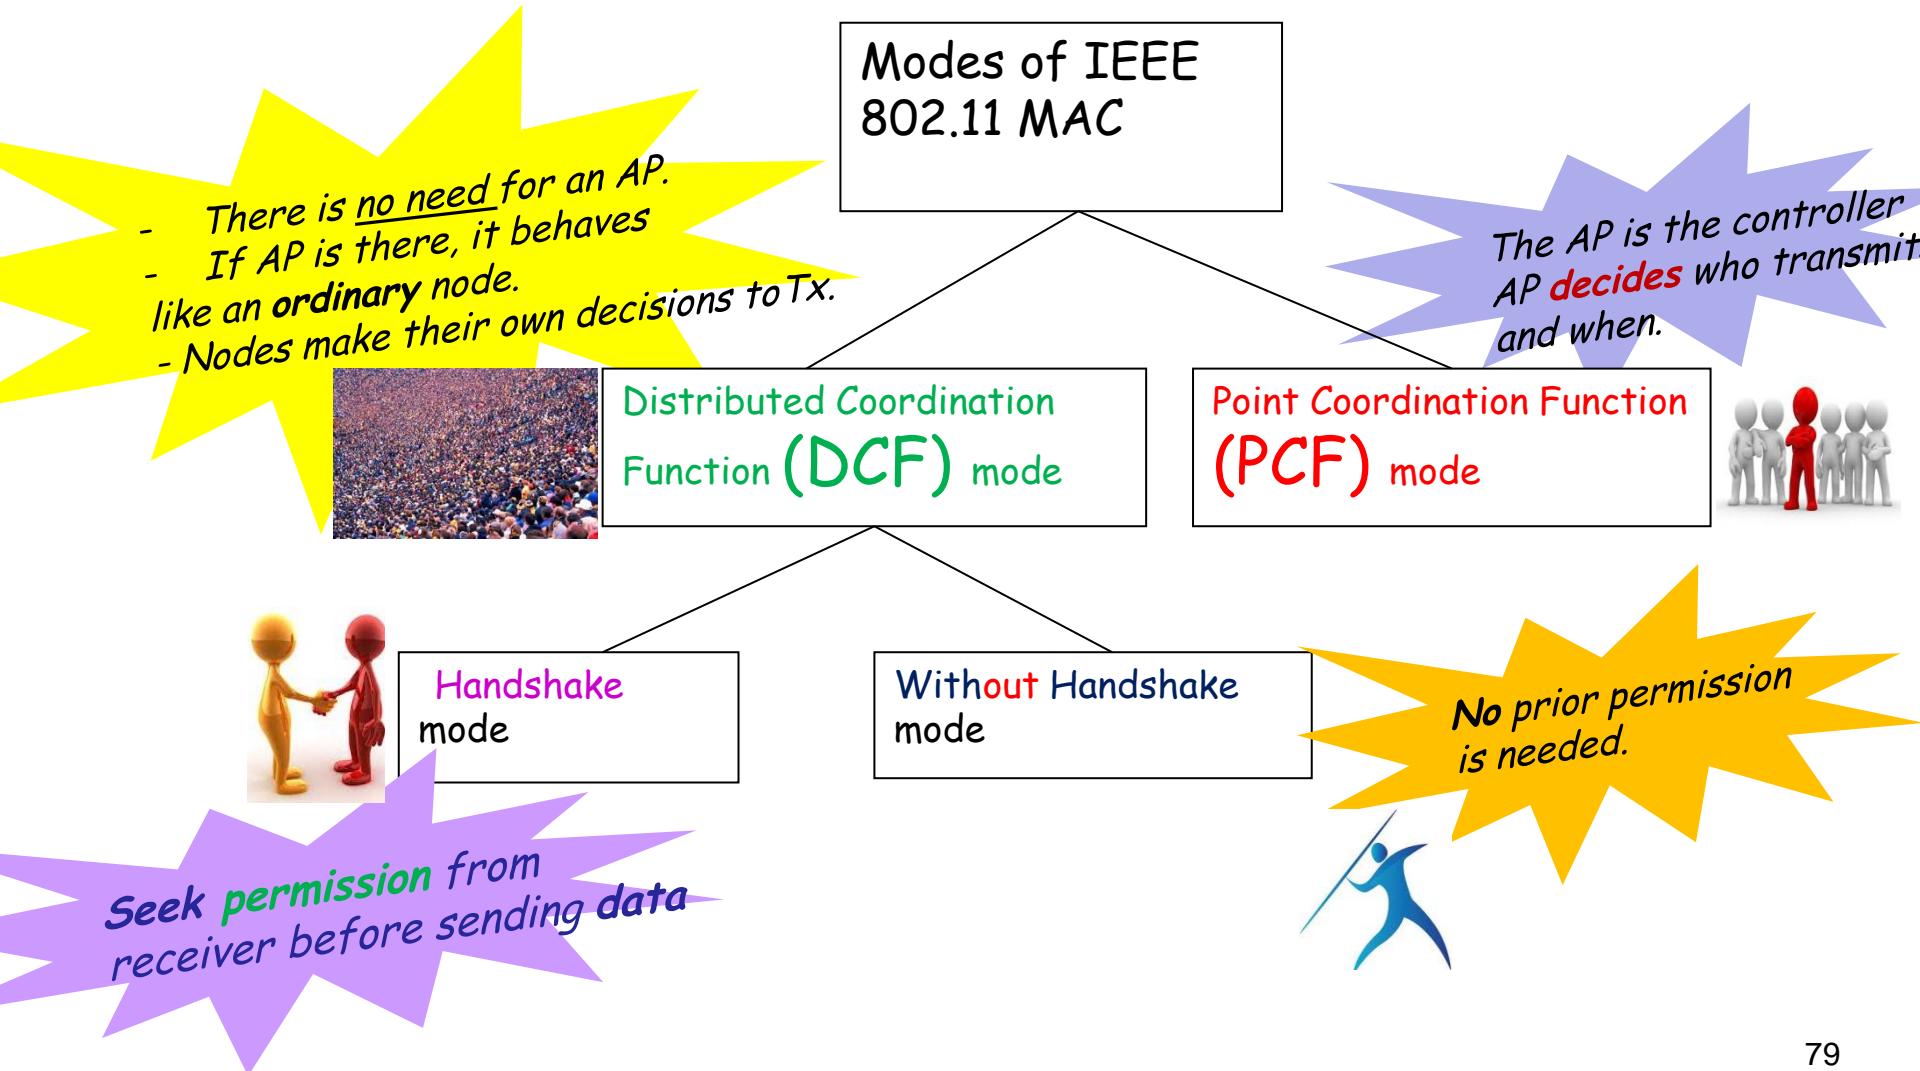
They have all different
PHY layers



But, all of them have the
same MAC protocol

First

Different Modes of Operation of MAC in IEEE 802.11





PCF Mode



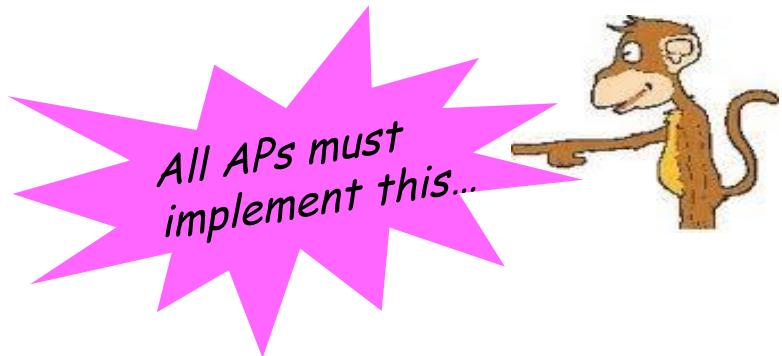
❖ The AP

- Operates as the **central controller** in the BSS.
- Decides who transmits and when.
- There is no contention for medium access
- Can follow a round-robin policy to allocate slots.

❖ This mode

- Leads to **waste of bandwidth** if a scheduled node has no traffic.

DCF Mode



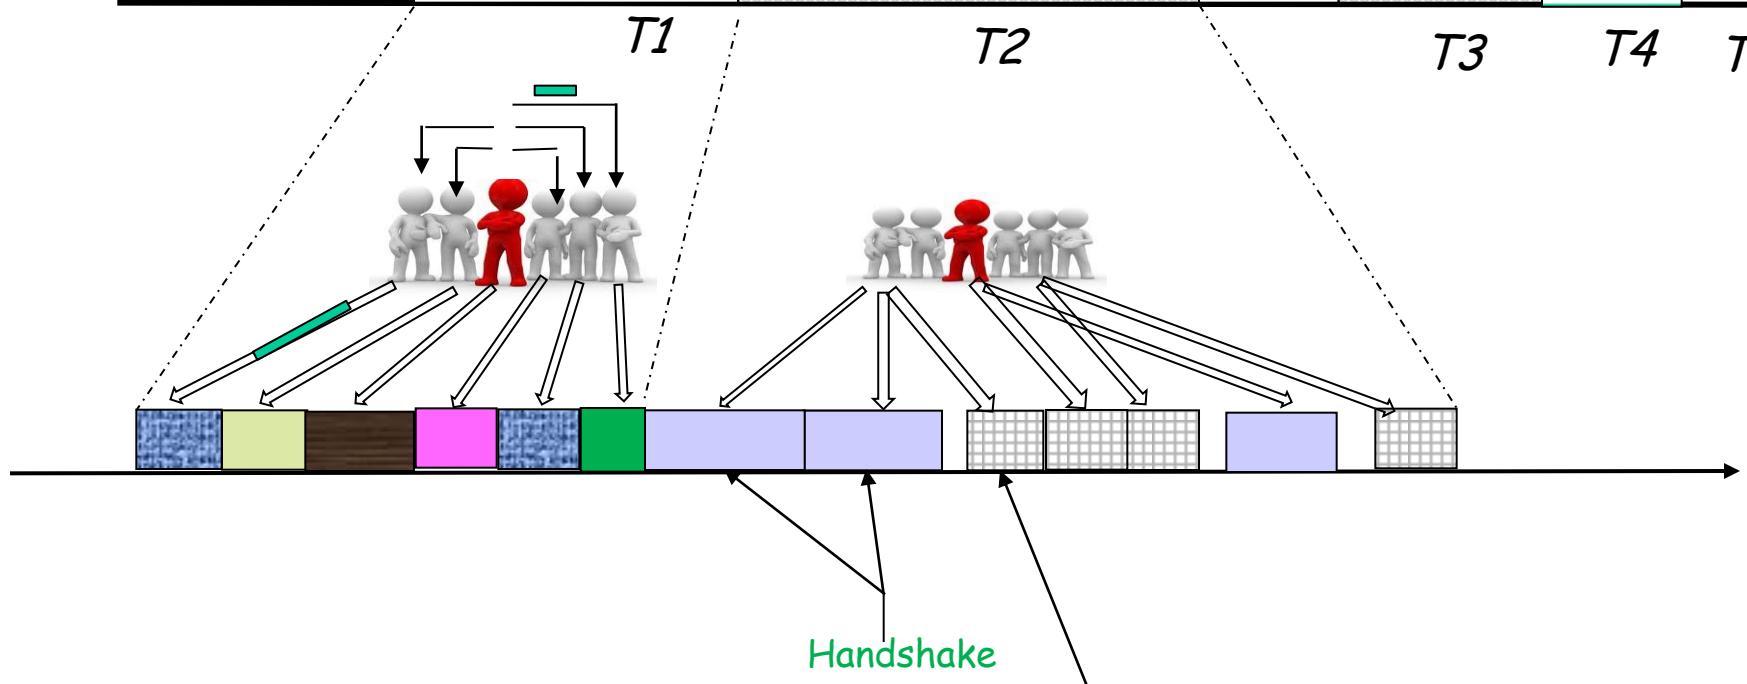
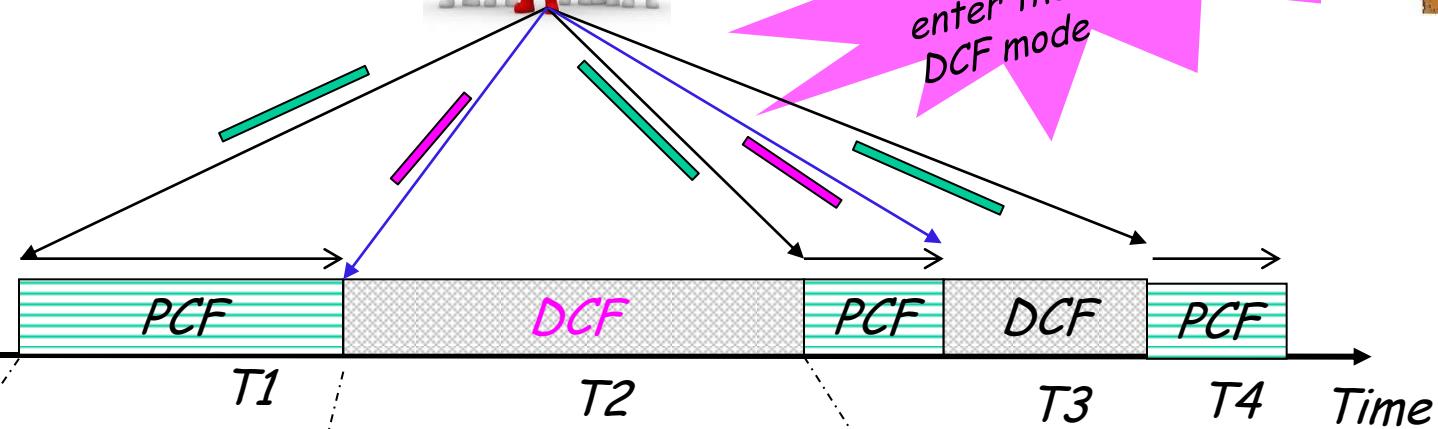
- ❖ An AP
 - Need not be used.
 - Computers can directly communicate among themselves <= Ad hoc.
 - Is used to provide connectivity to the Internet.

- ❖ In DCF
 - All nodes, including the AP, compete for medium access.
 - The AP does not operate as a central controller.

How are all those modes related?



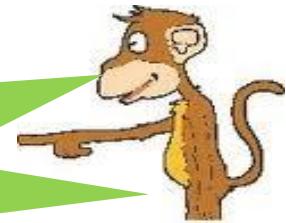
The AP tells all when to enter the PCF mode and DCF mode



How do I know when to use
handshake and when to use
without-handshake



The MAC management database
holds a variable: **dotRTSThreshold**
(integer in bytes)



Handshake mode

Frame length \geq dotRTSThreshold

Without-Handshake mode

Frame length $<$ dotRTSThreshold

For "long" frames, you want
to reduce the prob. of collision...
with some overhead

DCF with handshake

How is handshake implemented

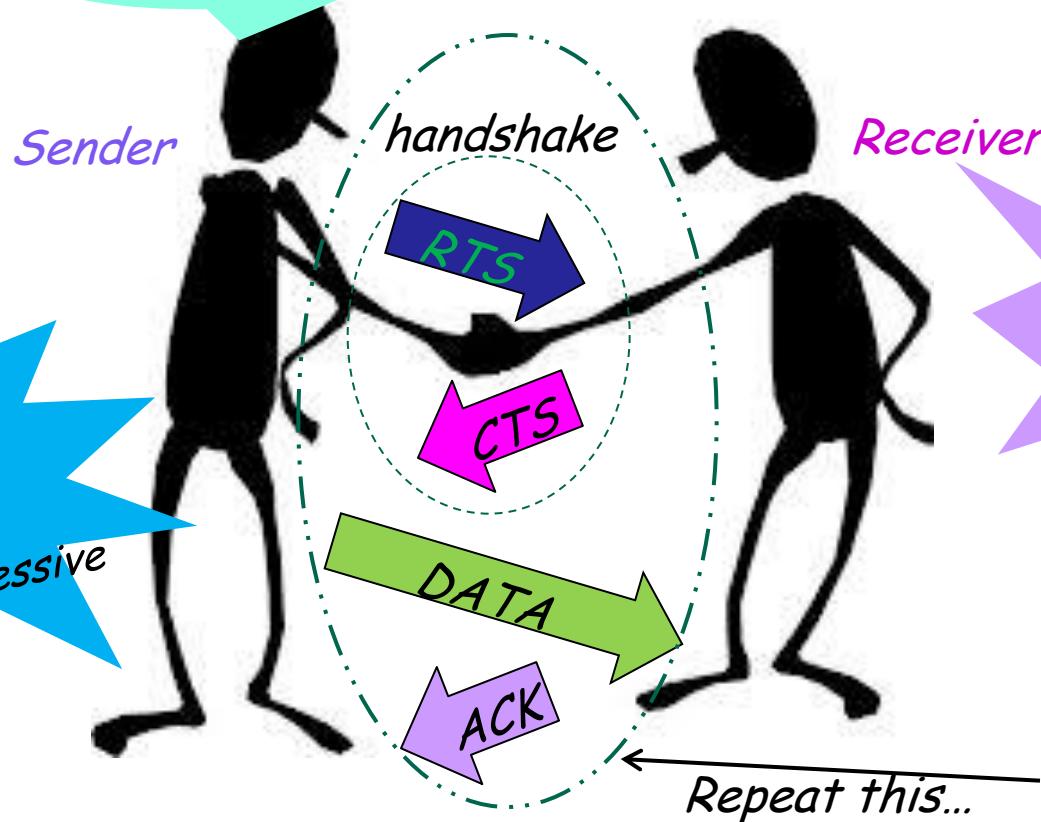


Under a certain condition
(to be explained soon)

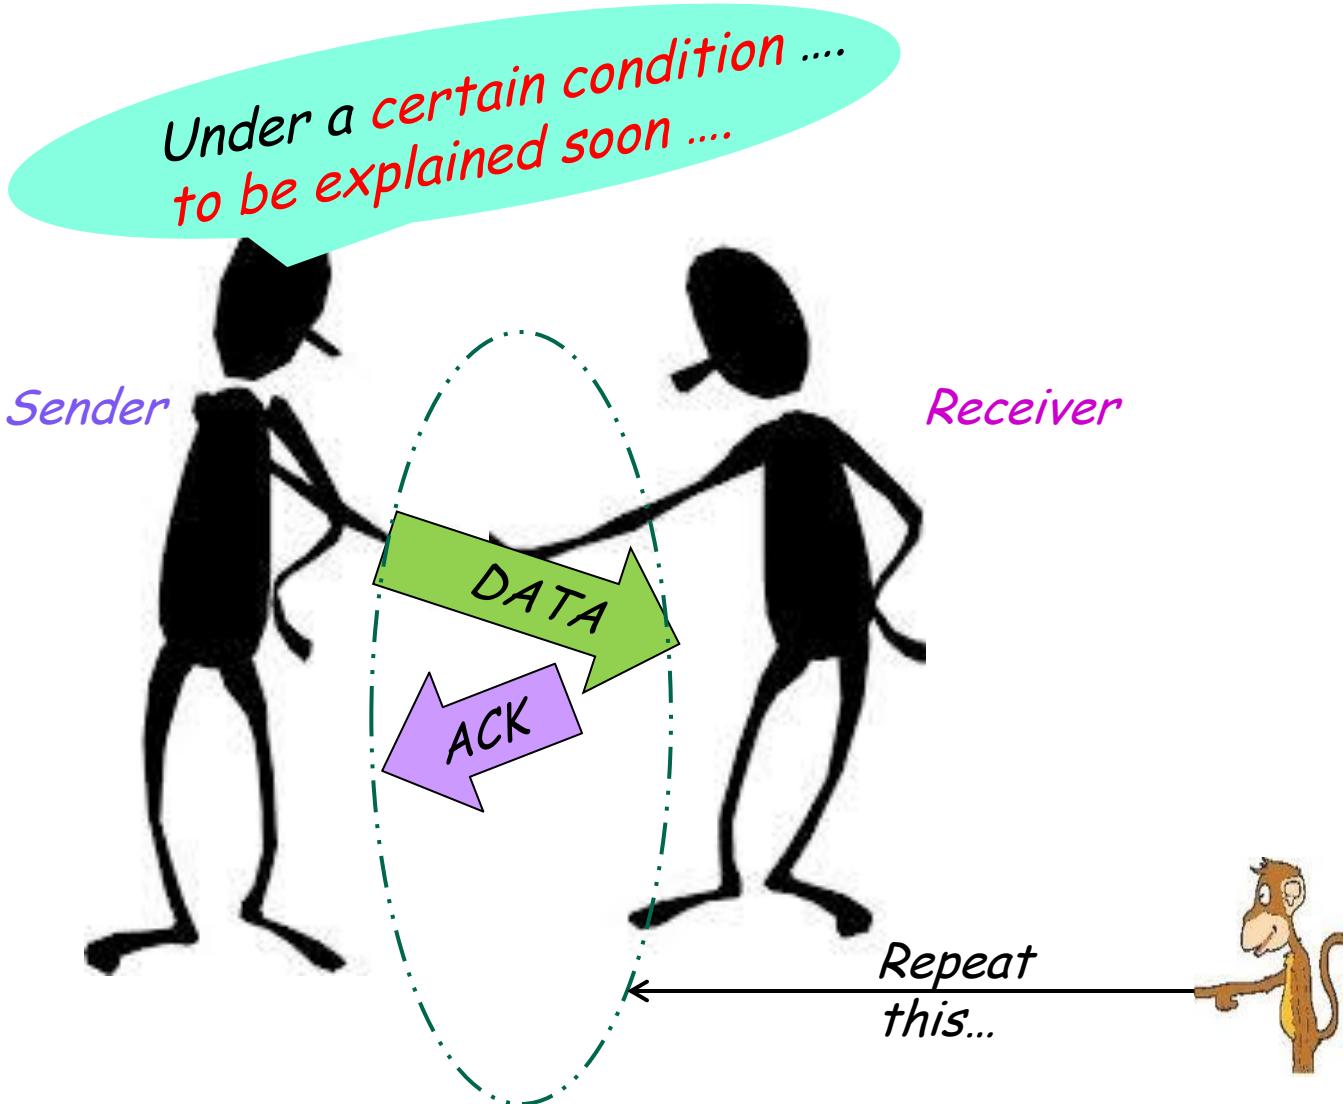
Control Frames:

RTS: Request To Send

CTS: Clear To Send



DCF without handshake



What is the big deal in a WLAN?

What happens to my RTS, CTS, DATA, ACK,.....
Do they **collide** with Tx from other nodes

Note the diff. between
a **wired** medium & a **wireless** medium



**Signal from all nodes
must reach all others**

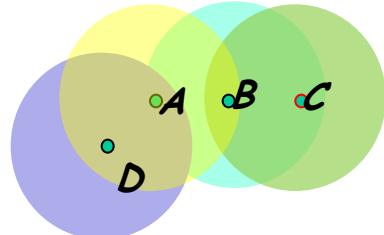


Wired

Remember this diff...

Wireless

**Signal from some nodes
may not reach all others**

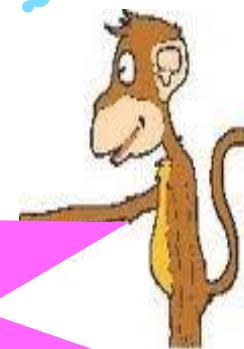


This diff. leads
to some problems

What problems can arise
in a wireless net



Collision detection is not always
possible in a wireless LAN.



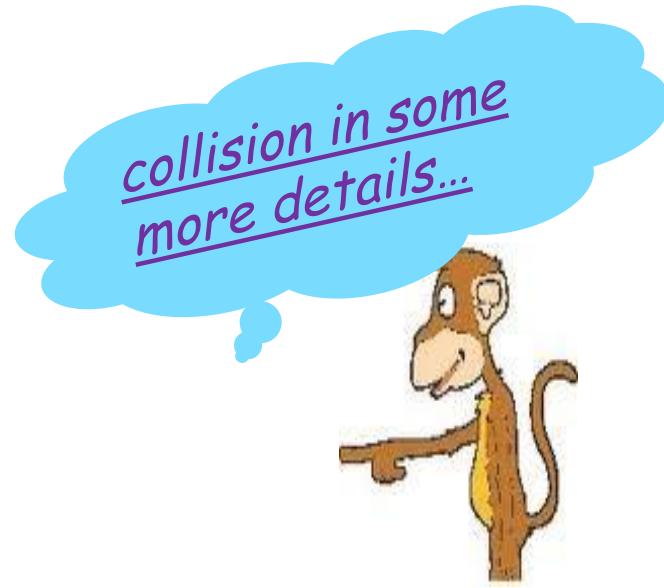
You do not receive
the **signal** received by
the intended receiver of
your frame.



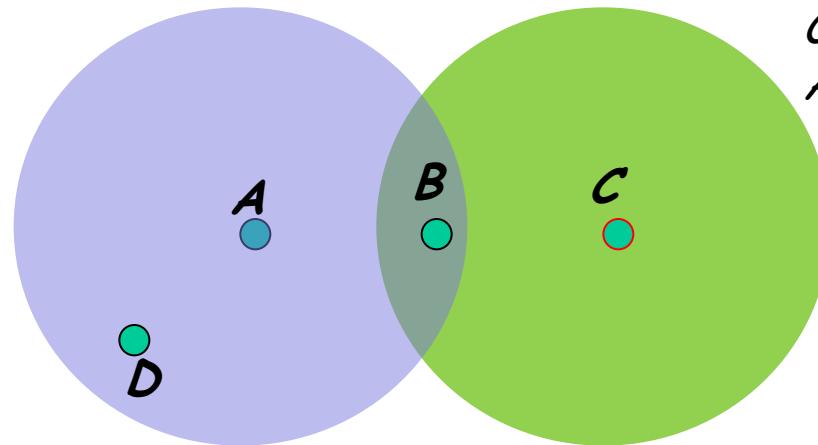
As a result, if your frame
encountered collision at the
receiver, you are **NOT** aware
of that.

Two Problems in WLAN

- ❖ Hidden Terminal Problem
- ❖ Exposed Terminal Problem



Hidden Terminal Problem



*C is transmitting a frame to B.
A is unaware of C's Tx.*

Now, if A transmits, A's frame will collide with C's at B

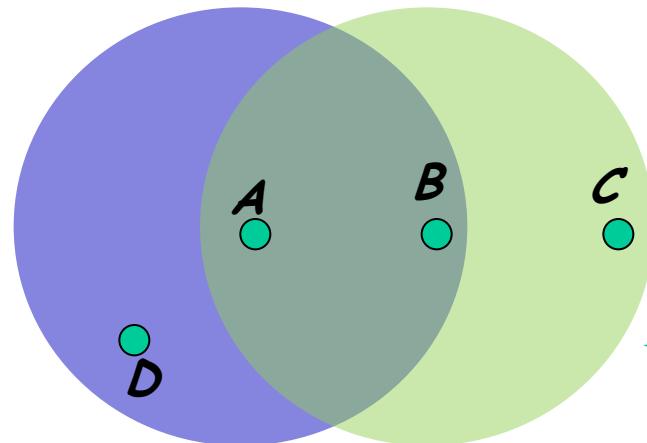
- This problem is due to C being **hidden** from A.
 - *Hidden means being "far away" ...*



*Collision occurs at B,
but A cannot detect it...*

*Solution exists:
CSMA/CA
(Collision Avoidance)*

Exposed Terminal Problem



A is *transmitting* a frame to D.

B knows that someone is transmitting.

If B transmits a frame to C, it does not collide with A's at D.

However, B *does not transmit* because it is unaware of D's location.

Problem: Loss of opportunity to transmit → Loss of bandwidth

- The above problem is due to B being *exposed* to A's Tx.

WLAN MAC: CSMA/CA (basic idea)

- ❖ Collision is avoided by using two techniques:
 - **PHY-level carrier sensing:** Done in receiver hardware
 - **Virtual carrier sensing:**
 - An **integer variable** in each node - called **NAV** (Network Allocation Vector) -- indicates whether or not a nearby node is likely to be transmitting.
 - **NAV > 0** → Most likely another node is transmitting.
 - **NAV = 0** → No one nearby is transmitting.
 - **NAV < 0**: This condition does not occur.
- ❖ **Transmit condition:** *Medium is idle at the receiver.*

(Carrier is absent) AND (NAV = 0)



Time
91

CSMA/CA: NAV in detail

- Recall the handshake mechanism with *RTS* and *CTS*
- A *duration* field in frame header indicates the length of time the sender of the frame may use the medium.
- All nodes which receive *RTS* and *CTS* (and *DATA* and *ACK*) update their *NAV* as follows
 - Events that update *NAV*: Initially $NAV = 0$

With each passing μs (micro second)

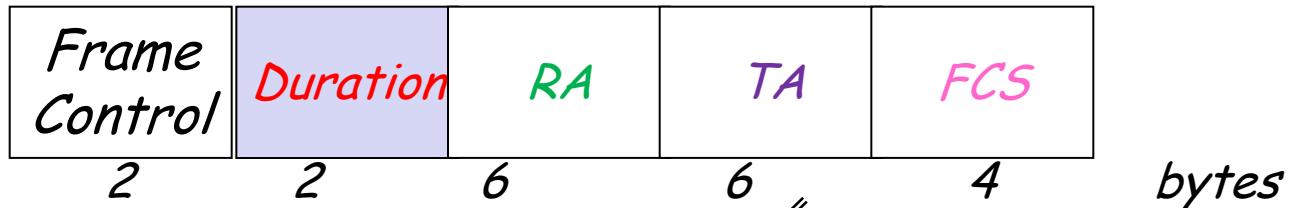
If $(NAV > 0)$ $NAV = NAV - 1$

Else stop decrementing *NAV*

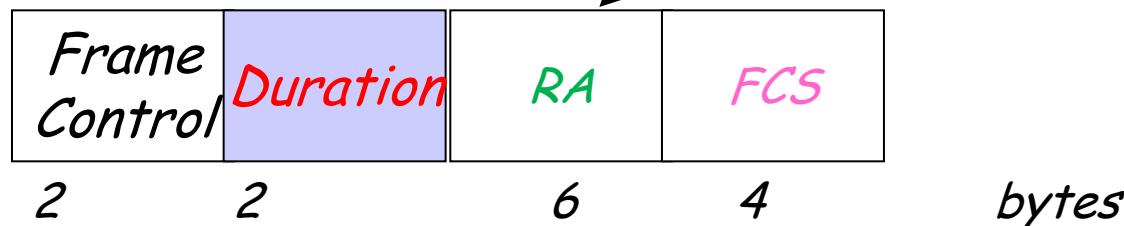
If a frame (e.g. *RTS*, *CTS*) with *duration* field is received
 $NAV = \text{Max}(NAV, \text{duration})$

RTS and CTS Frames

RTS



CTS/
ACK

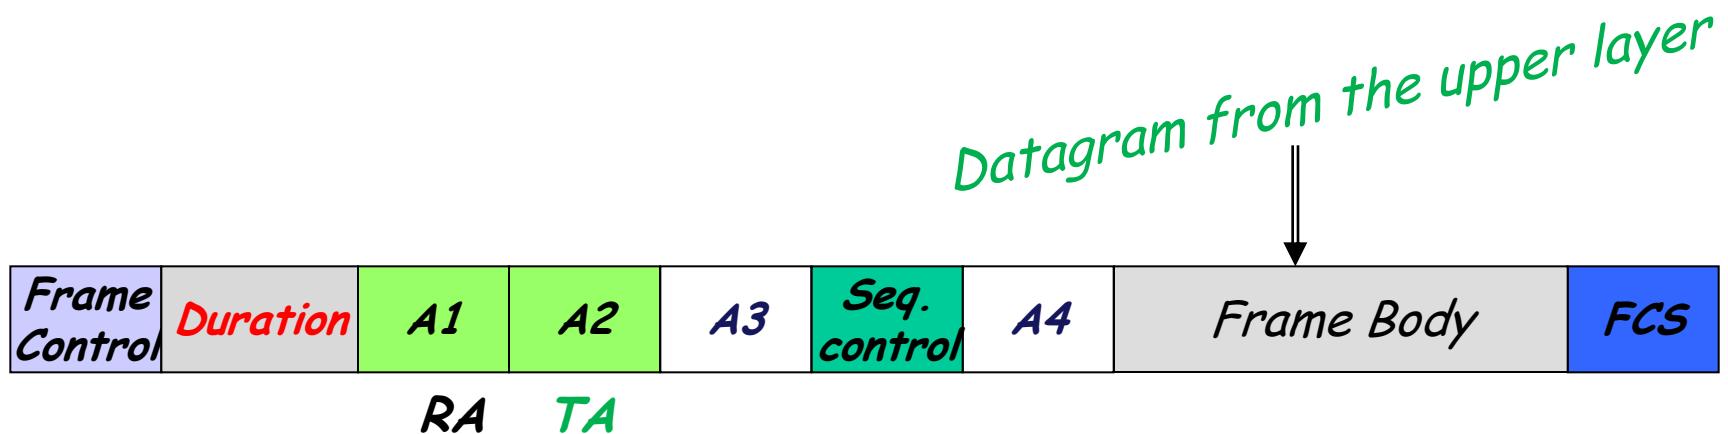


FCS: Frame Check Sequence ← CRC

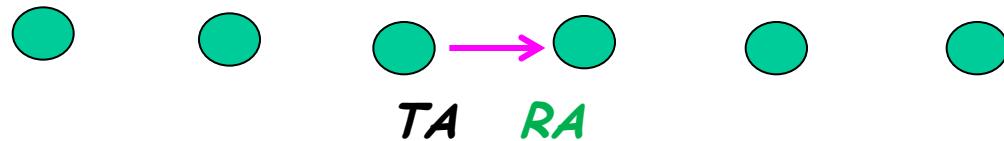
RA: Receiver Address

TA: Transmitter Address

DATA Frame



TA: Physically transmitting the frame.



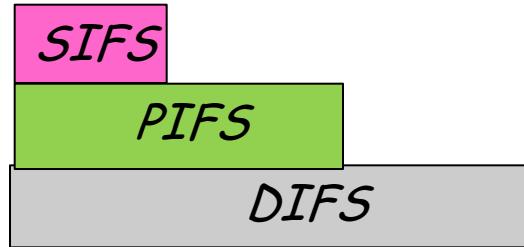
Timing Intervals

- ❖ The IEEE 802.11 MAC defines 4 timing intervals
 - 2 at the PHY level
 - *SIFS*: Short Inter-Frame Space (10 micro-sec) between successive frames (RTS, CTS, DATA, ACK, ...)
 - *aSlot* (20 micro-sec)
 - 2 at the MAC level
 - *PIFS*: Priority (in PCF) IFS (*SIFS* + *aSlot*)
 - *DIFS*: Distributed IFS (*PIFS* + *aSlot*)

aSlot is chosen s.t. a node can determine if another node initiated a Tx *aSlot* time before (= propagation time + some hardware delay)



Timing intervals are used to control priority



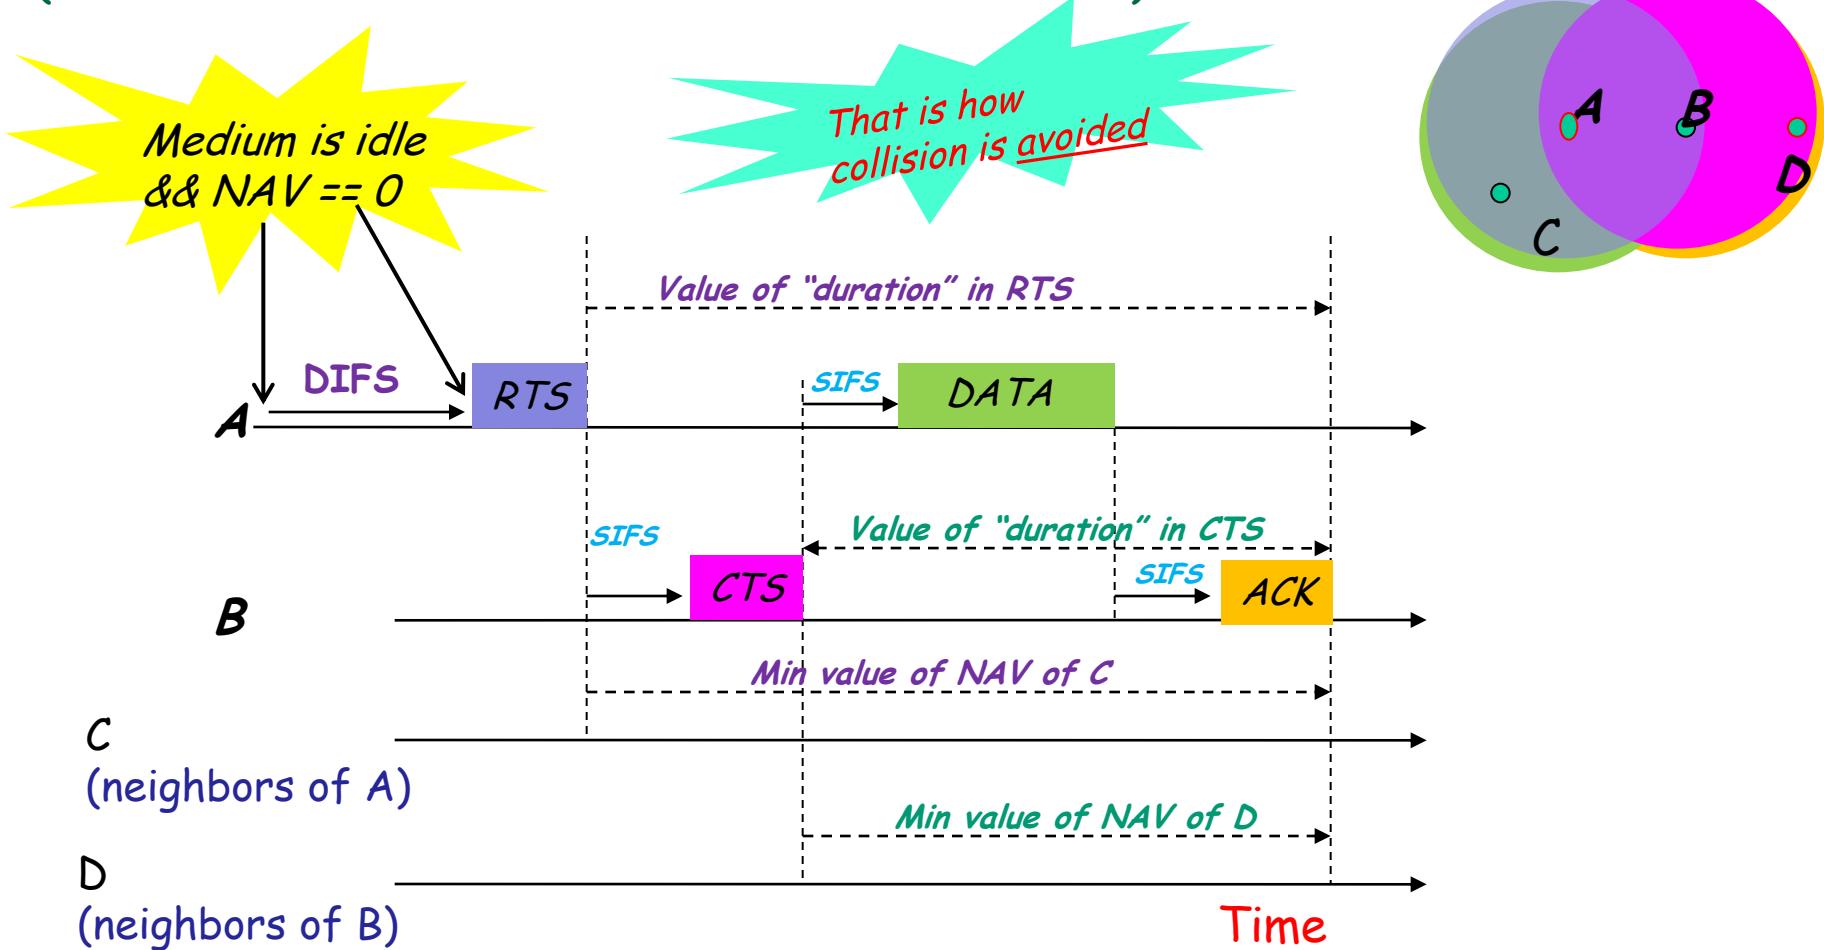
By keeping SIFS shortest, it is ensured that an ongoing cycle of Tx (with handshake or without handshake) is not disturbed

No one else can Tx between RTS, CTS, DATA, ACK.

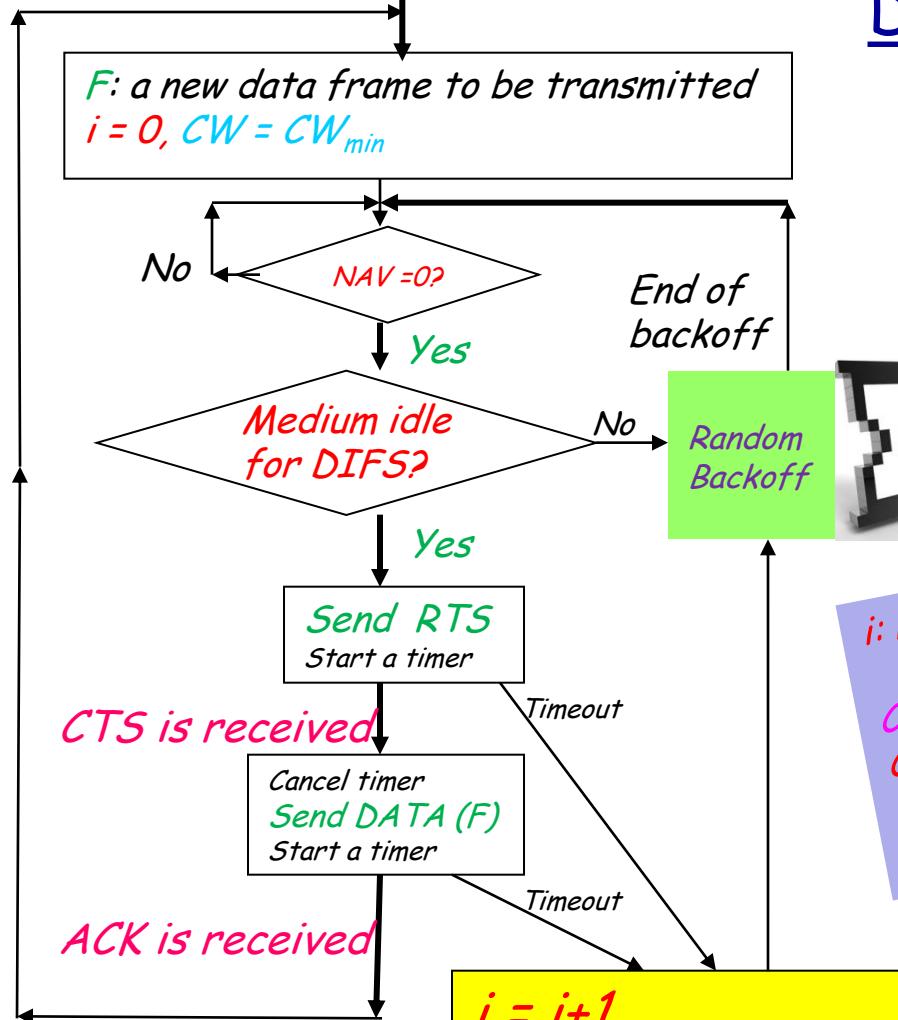
PIFS < DIFS enables an AP to become the controller of a BSS ...

Handshake using RTS/CTS (A wants to send data to B)

(Note: "duration" info in DATA will not be shown.)



DCF with Hand-shake: Tx



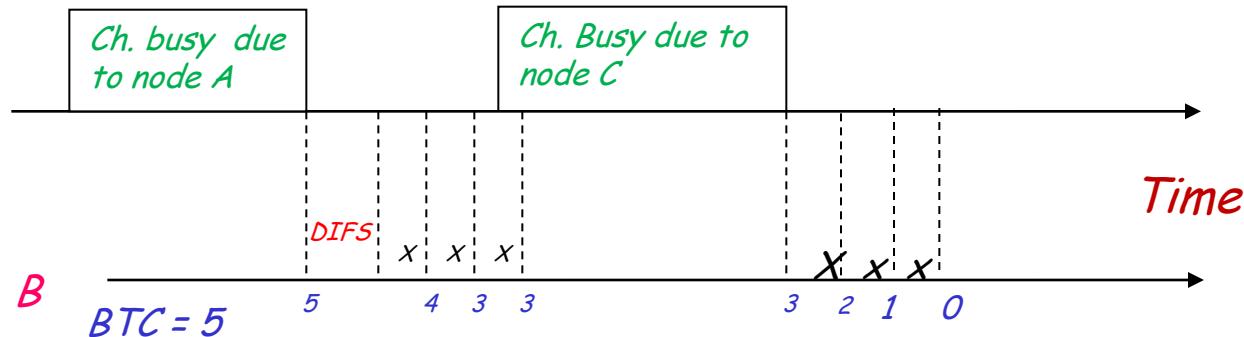
i : Retry count, CW : Contention Window
 CW_{min} : Min. value of CW (typical: 32)
 CW_{max} : Max. value of CW (typical : 256)
DIFS: Distributed Interframe Space

$CW = CW_{min} * 2^i$
 $(CW \text{ saturates at } CW_{max})$

Backoff Mechanism

- ❖ Initialize a counter: Backoff Time Counter (BTC)
 - $BTC = \text{Random}(0, CW-1)$ time unit of BTC is aSlot
- ❖ As time passes, BTC is decremented as follows
 - At the start, let ch. remain idle for DIFS.
 - Next, if ch. is idle for aSlot: $BTC = BTC - 1$ \leftarrow Repeat this
 - Anytime the medium is busy: Freeze BTC
- ❖ $BTC == 0 \rightarrow$ End of backoff

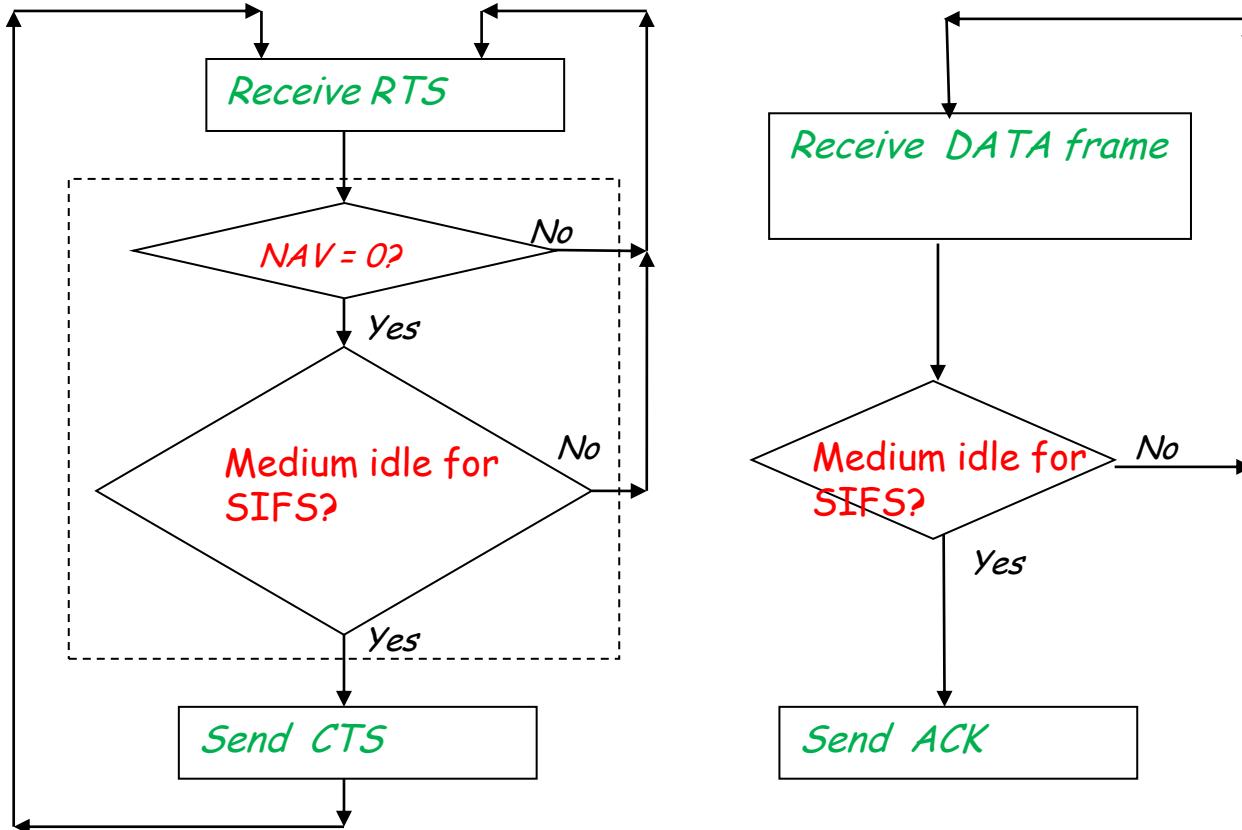
Backoff Mechanism



B is executing *backoff*

$$X = aSlotTime$$

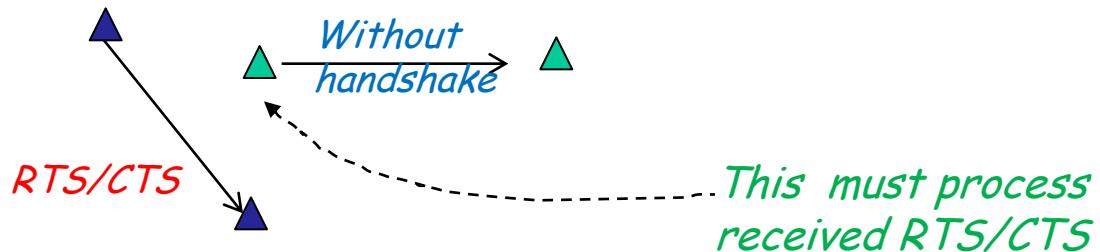
DCF with Handshake: Rx



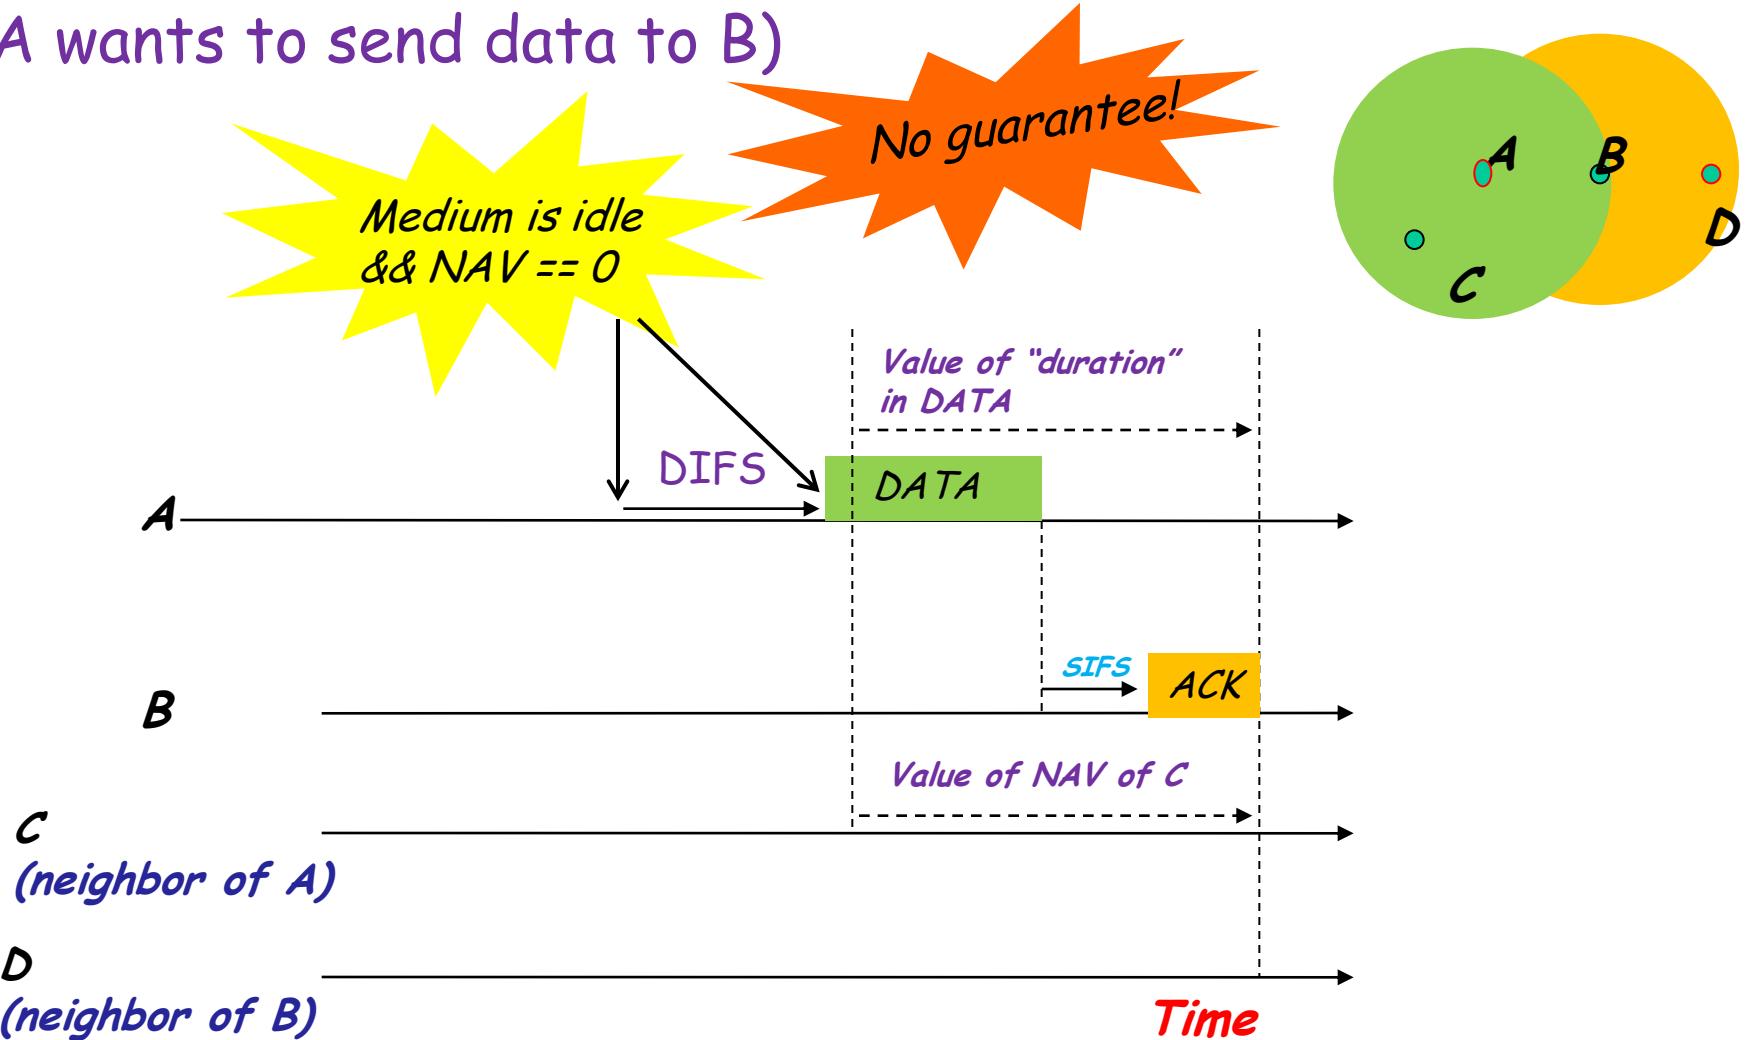
Note: The above two fragments of flow-charts can be easily merged.

DCF Mode **without** Handshake

- ❖ A special case of DCF with handshake
 - RTS/CTS frames are **not** exchanged.
- ❖ The idea of **NAV** is still used in this mode
 - Recall: All nodes process the received **RTS/CTS** of others

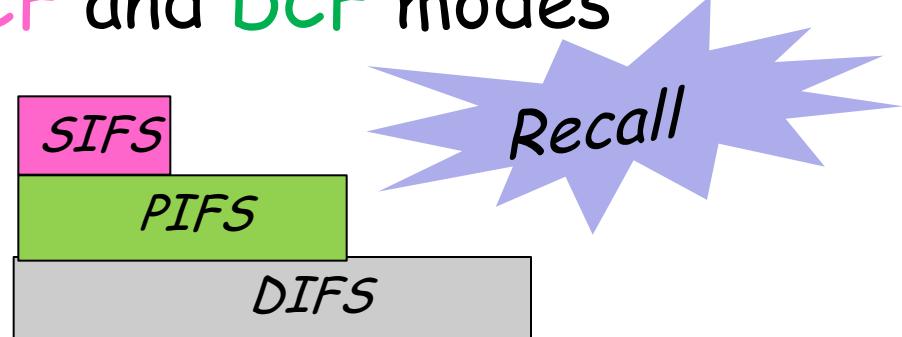


Data transmission without handshake (A wants to send data to B)



PCF Mode: AP becomes the controller: How?

- ❖ AP alternates between PCF and DCF modes



- ❖ AP operates as the controller as follows

If AP finds medium to be idle for PIFS, it transmits a beacon frame

Beacon contains a CFPMaxDuration field

Nodes receiving a beacon update their NAV to CFPMaxDuration

These nodes perceive the medium to be busy for CFPMaxDuration

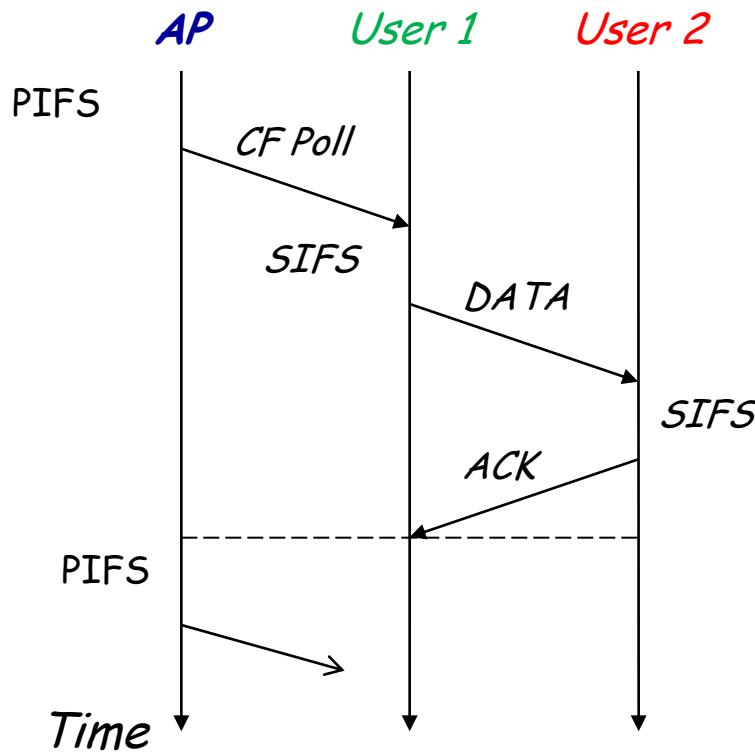
They do not transmit unless asked to do so by the AP.

PCF Mode of Operation (Contd.)

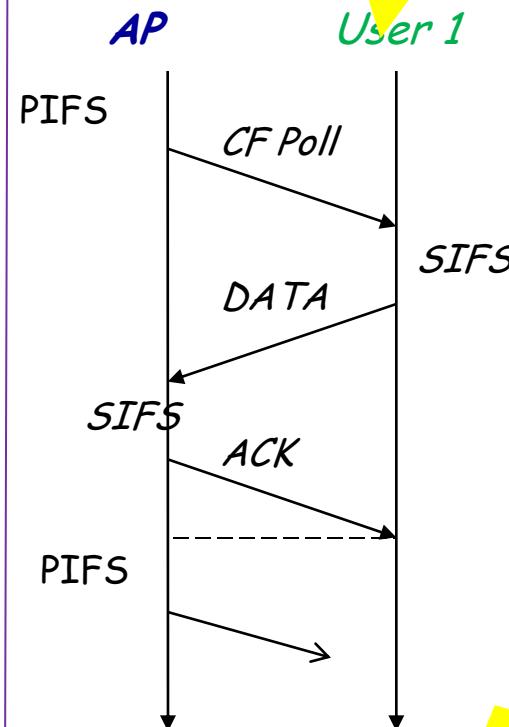
- After transmitting a beacon, AP waits for PIFS before transmitting one of the following
 - DATA frame
 - CF Poll frame **Contention Free**
 - DATA+CF Poll frame
 - ACK frame
 - CF End frame

PCF Mode of Operation (Contd.)

CF Poll frame



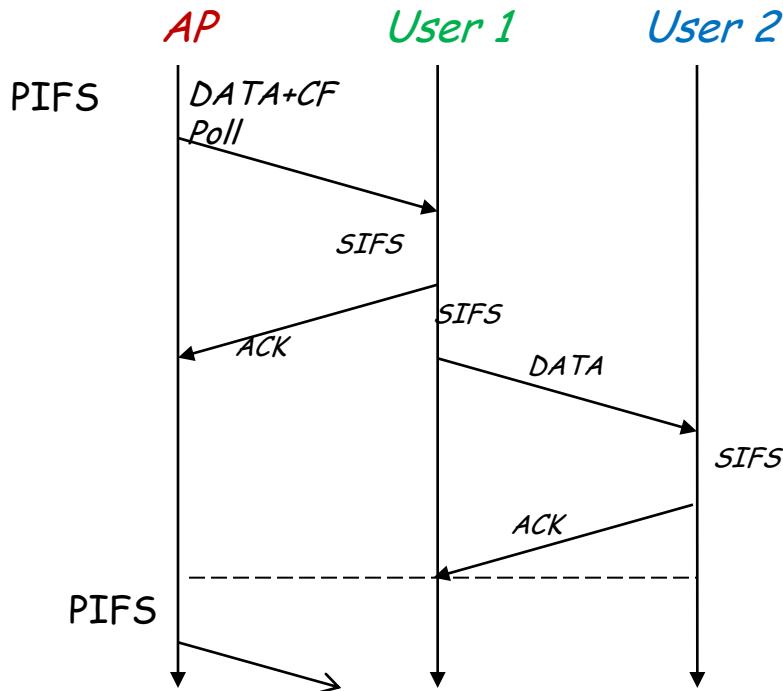
If a user does **not** have data, it sends a **null** DATA frame (A DATA frame with no actual data)



If ACK is not received, user 1 waits until it is polled again.

PCF Mode of Operation (Contd.)

DATA + CF Poll frame



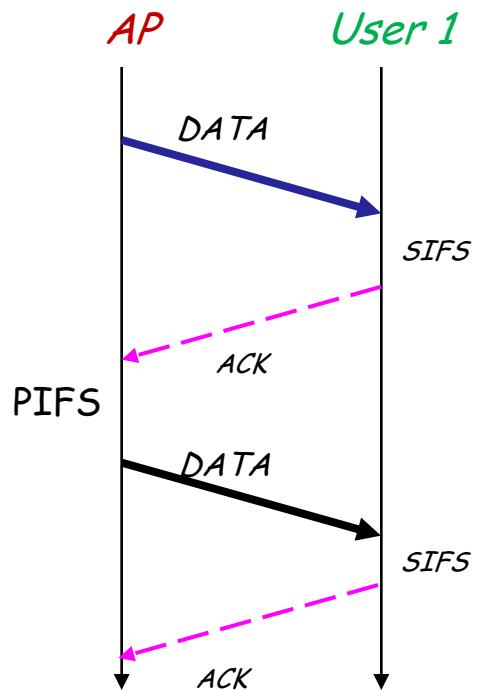
Note 1: If AP does **not** receive an ACK, it **retransmits** data after PIFS.

Note 2: If User 1 does not receive ACK, it does **not** retransmit data.

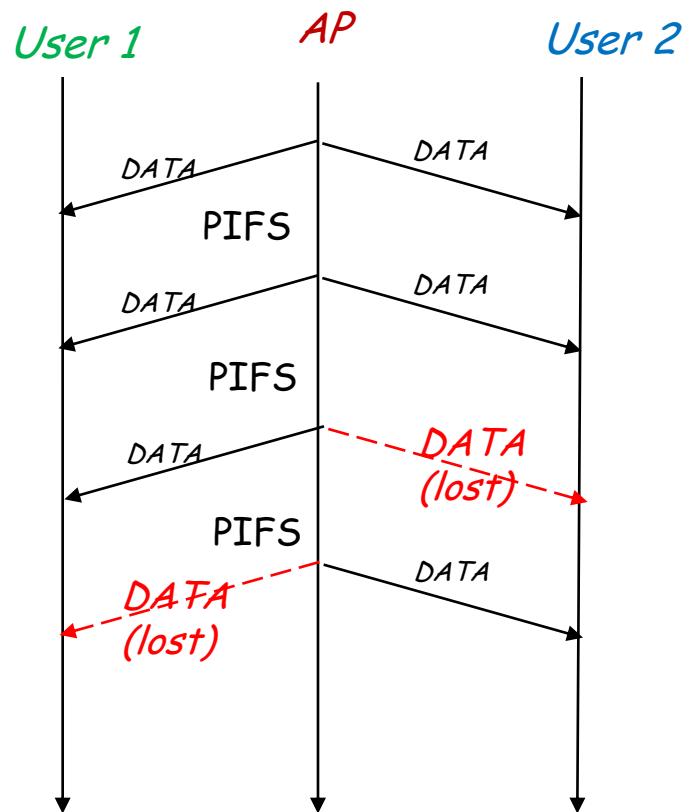
The polled user receives data from the AP and sends data to another user.

PCF Mode of Operation (Contd.)

DATA from AP (unicast): $1 \rightarrow 1$



DATA from AP (broadcast): $1 \rightarrow \text{all}$

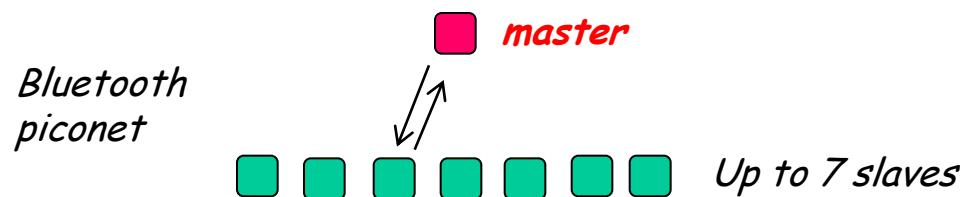


PCF Mode of Operation (Contd.)

- ❖ CF End frame
 - Identifies the **end of CF period**
 - Receiving nodes set **NAV = 0.**

Summary of MAC protocols

- ❖ *channel partitioning* by time, frequency or code
 - Time Division, Frequency Division, Code Division
- ❖ *random access* (dynamic),
 - ALOHA, S-ALOHA, CSMA/CD, CSMA/CA
 - carrier sensing: easy in some technologies (wire), hard in wireless
 - CSMA/CD used in Ethernet
 - CSMA/CA used in 802.11
- ❖ *taking turns*
 - polling from central site, token passing
 - **Bluetooth**, IBM Token Ring



Could I see some past exam questions ...

F'12 Mid-term



Identify five differences between the CSMA/CD and the CSMA/CA protocols by identifying five comparison criteria.

Use the following table to answer the question.



Criteria	CSMA/CD	CSMA/CA

More past exam questions ...

F'12 Mid-term



Clearly explain why collision detection is not possible in wireless local area networks.

Clearly explain why the timing intervals SIFS, PIFS, and DIFS have been ordered as SIFS < PIFS < DIFS in the CSMA/CA protocol.

Clearly explain the hidden terminal problem

S'12 Mid-term

Link Layer

5.1 Introduction and services

5.2 Error detection and correction

5.3 Multiple access protocols

5.4 Link-Layer Addressing

5.5 Ethernet

5.6 Link-layer switches

5.7 PPP

3.4 Reliable Data Transfer

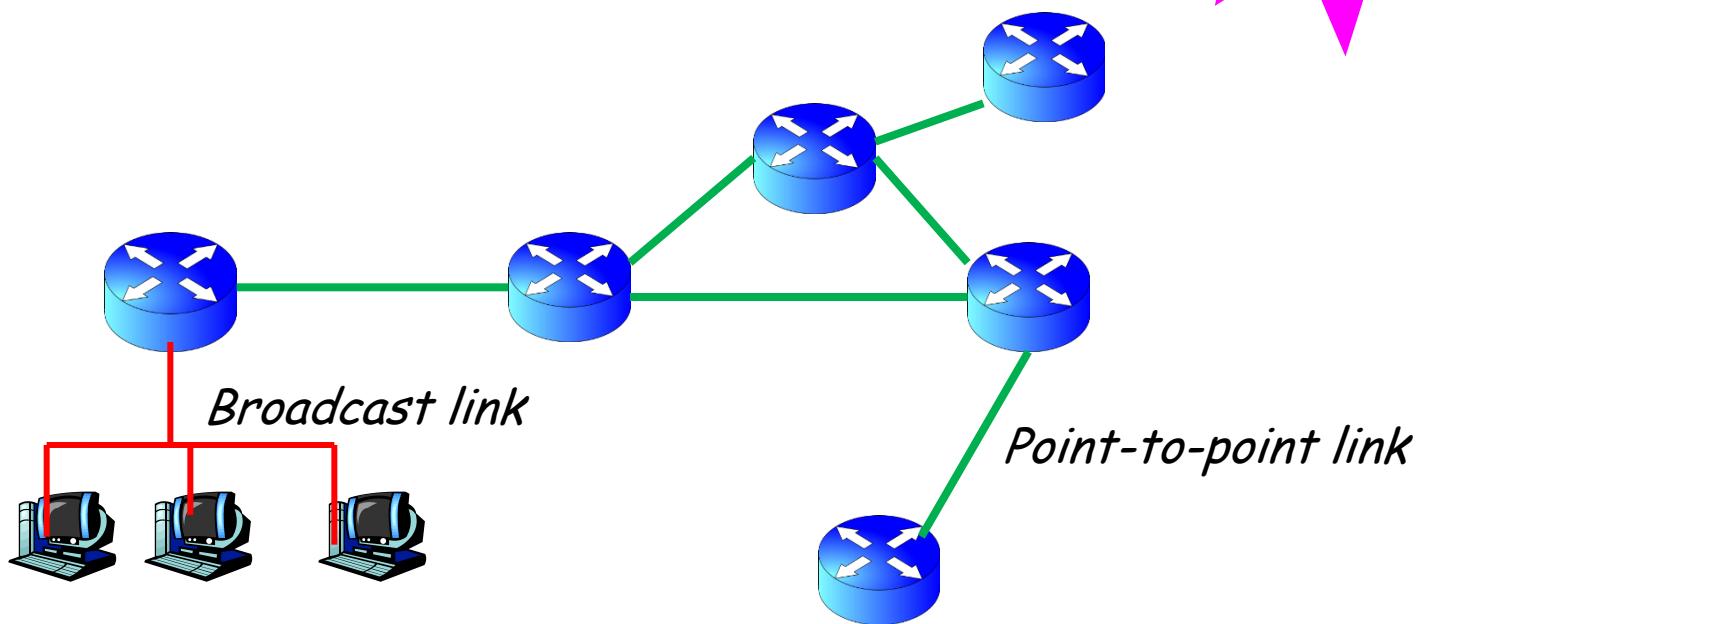
I learned about accessing broadcast links ...
Ethernet and WiFi



Are there other
kinds of links??



There are
point-to-point links



Point to Point Data Link Control

- ❖ one sender, one receiver, one link: easier than broadcast link:
 - no Media Access Control
 - no need for explicit MAC addressing
- ❖ popular point-to-point DLC protocols:
 - **PPP** (point-to-point protocol)
 - **HDLC**: High level data link control (Data link used to be considered "high layer" in protocol stack!)

PPP Design Requirements [RFC 1557]

- ❖ **packet framing:** encapsulation of network-layer datagram in data link frame
 - carry network layer data of any network layer protocol (not just IP) *at same time*
 - ability to demultiplex upwards
- ❖ **bit transparency:** must carry any bit pattern in the data field
- ❖ **error detection** (no correction)
- ❖ **connection liveness:** detect, signal link failure to network layer
- ❖ **network layer address negotiation:** endpoint can learn/configure each other's network address

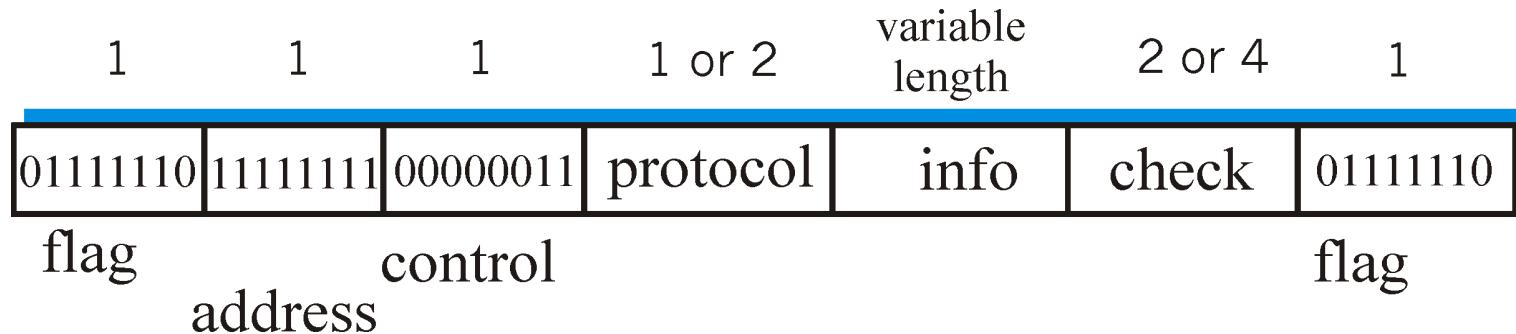
PPP non-requirements

- ❖ no error correction/recovery
- ❖ no flow control
- ❖ out of order delivery OK
- ❖ no need to support multipoint links (e.g., polling)

Error recovery, flow control, data re-ordering
all relegated to higher layers!

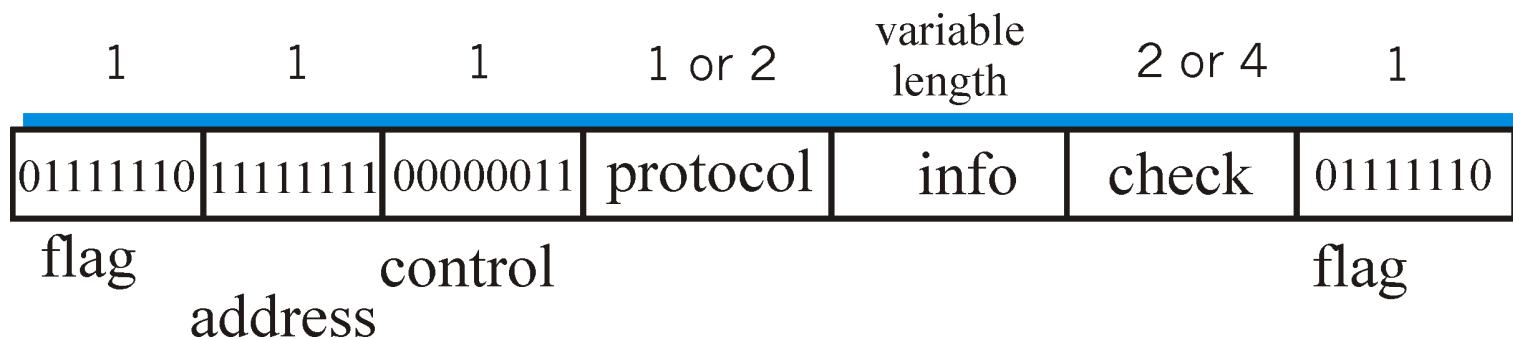
PPP Data Frame

- ❖ **Flag:** delimiter (framing)
- ❖ **Address:** does nothing (only one option)
- ❖ **Control:** does nothing; in the future possible multiple control fields
- ❖ **Protocol:** upper layer protocol to which frame delivered



PPP Data Frame

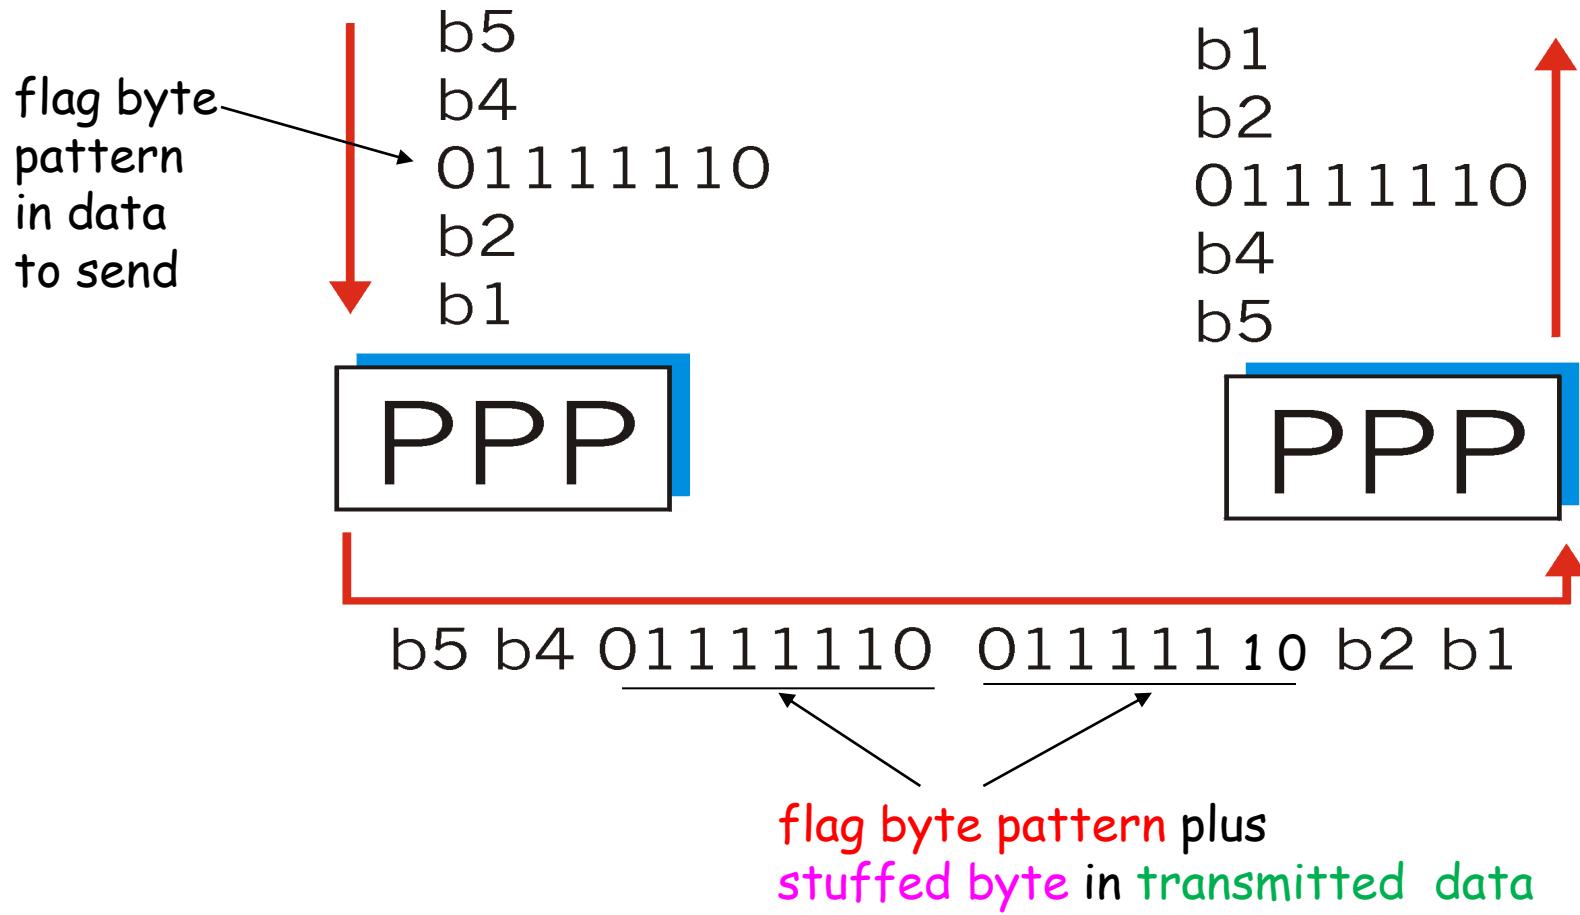
- ❖ **info:** upper layer data being carried
- ❖ **check:** cyclic redundancy check for error detection



Byte Stuffing

- ❖ “data transparency” requirement: data field must be allowed to include flag pattern <0111110>
 - **Q:** is received <0111110> data or flag?
- ❖ Sender: adds (“stuffs”) extra <0111110> byte after each <0111110> **data** byte
- ❖ Receiver:
 - two 0111110 bytes in a row: discard first byte, continue data reception
 - single 0111110: flag byte

Byte Stuffing

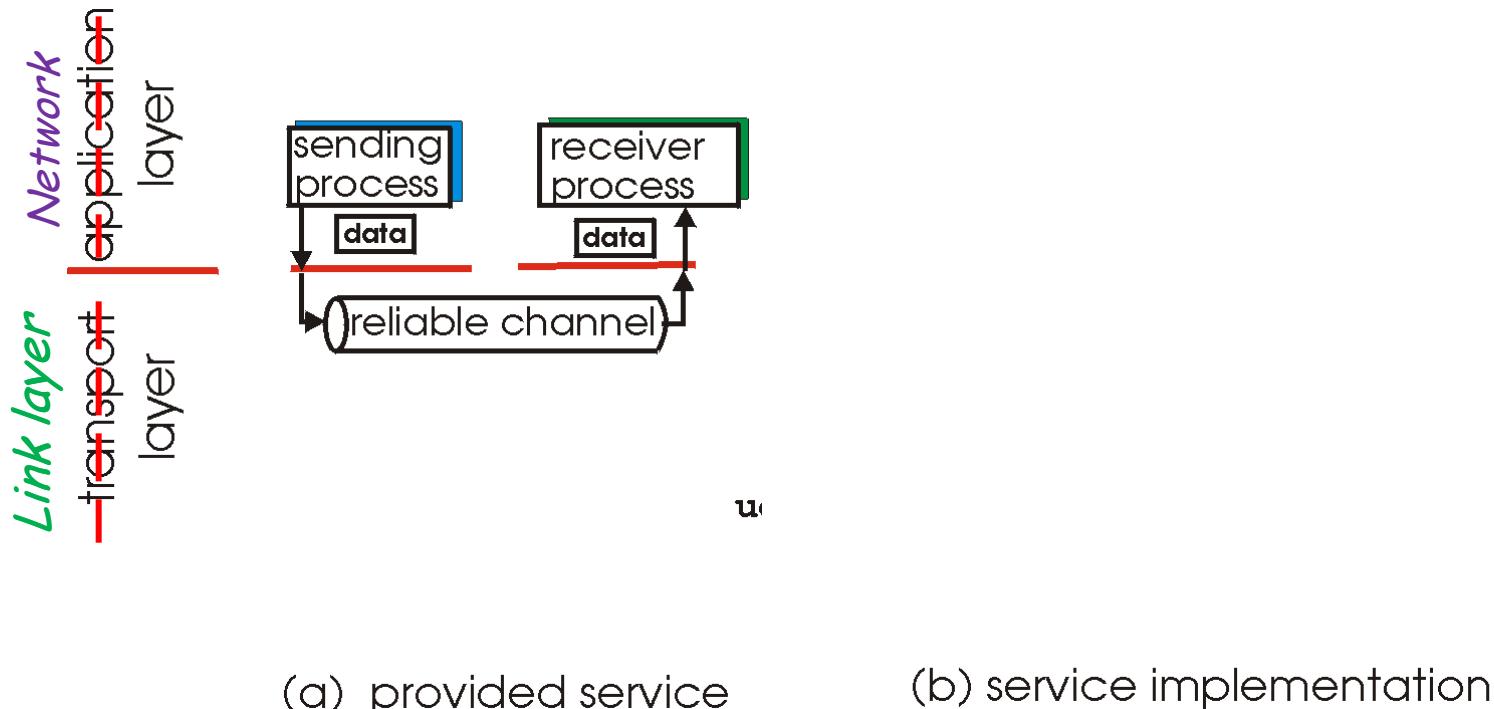


Link Layer

- ❖ 5.1 Introduction and services
- ❖ 5.2 Error detection and correction
- ❖ 5.3 Multiple access protocols
- ❖ 5.4 Link-Layer Addressing
- ❖ 5.5 Ethernet
- ❖ 5.6 Link-layer switches
- ❖ 5.7 PPP
- ❖ 3.4 Reliable data transfer

3.4 Principles of Reliable data transfer

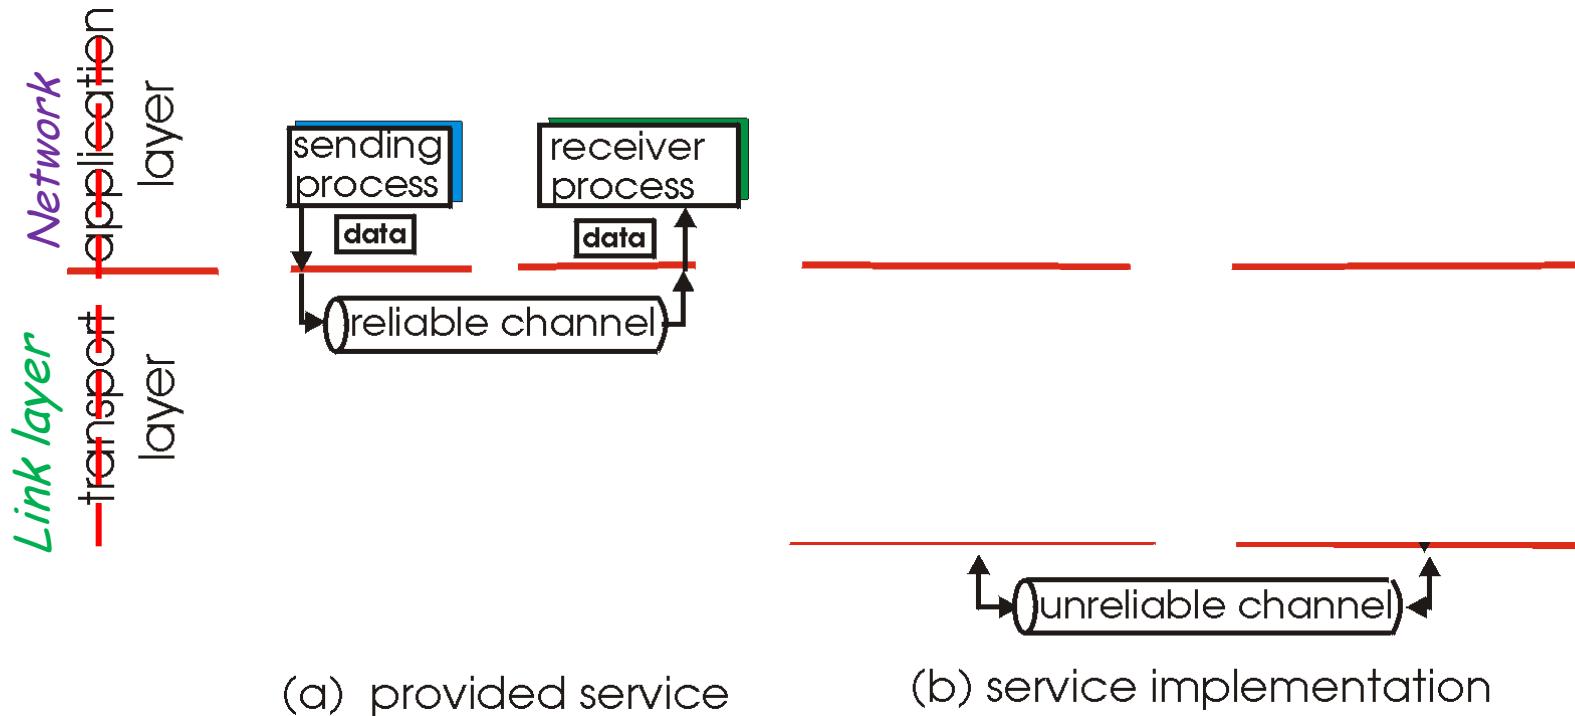
- ❖ important in app., transport, link layers



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of Reliable data transfer

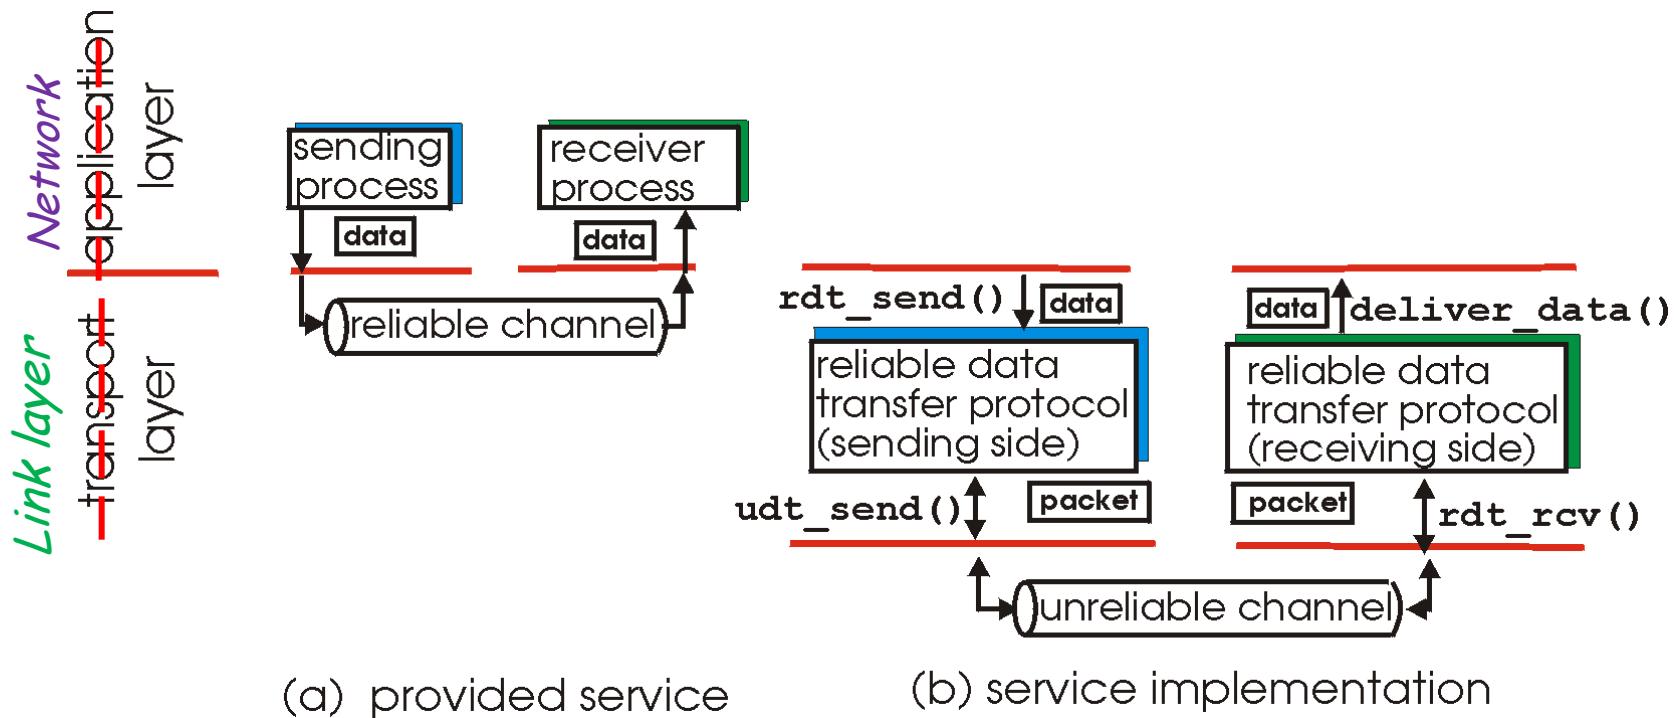
- ❖ important in app., transport, link layers



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

Principles of Reliable data transfer

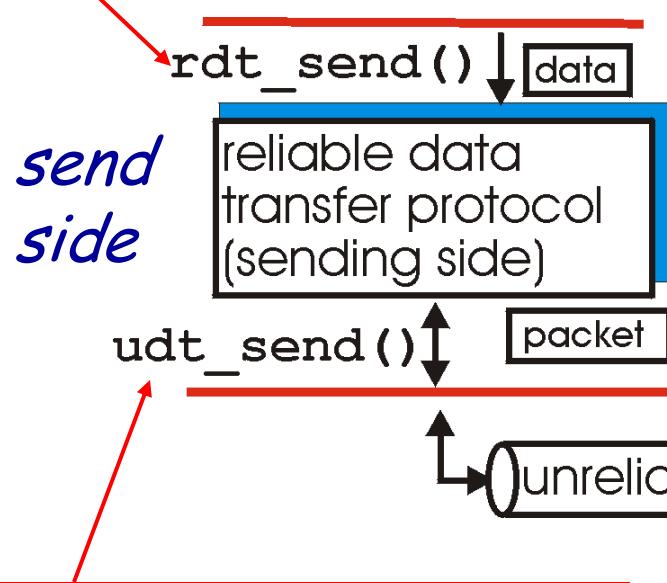
- ❖ important in app., transport, link layers



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

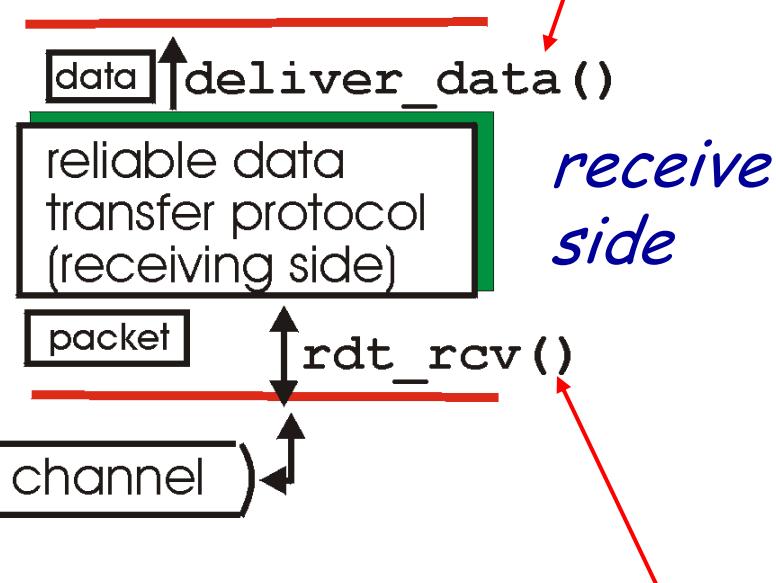
Reliable data transfer: getting started

rdt_send(): called from above, (e.g., by app.). Passed data to deliver to receiver upper layer



udt_send(): called by rdt, to transfer packet over unreliable channel to receiver

deliver_data(): called by rdt to deliver data to upper

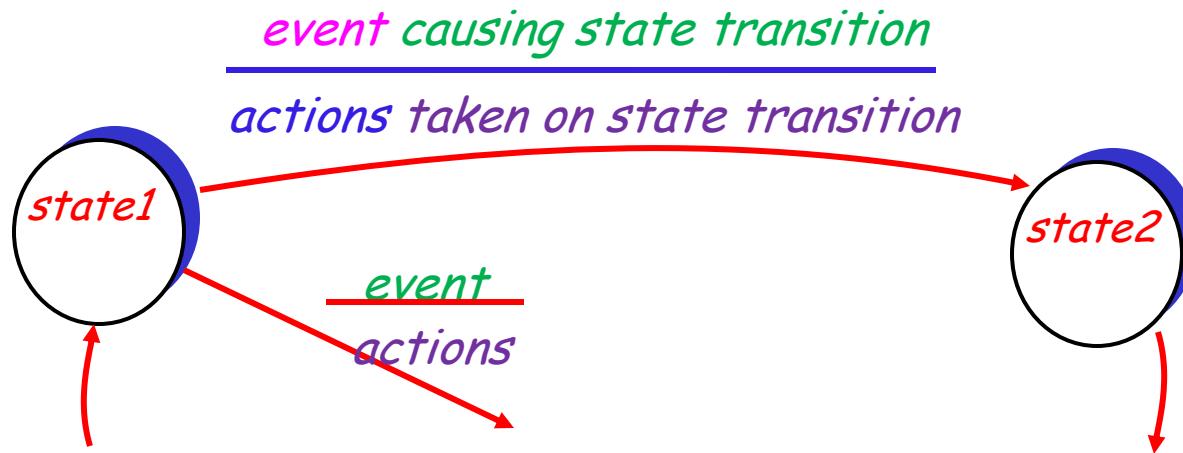


rdt_rcv(): called when packet arrives on rcv-side of channel

Reliable data transfer: getting started

We will:

- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ consider only unidirectional data transfer
 - but control info will flow in both directions!
- ❖ use finite state machines (FSM) to specify sender, receiver



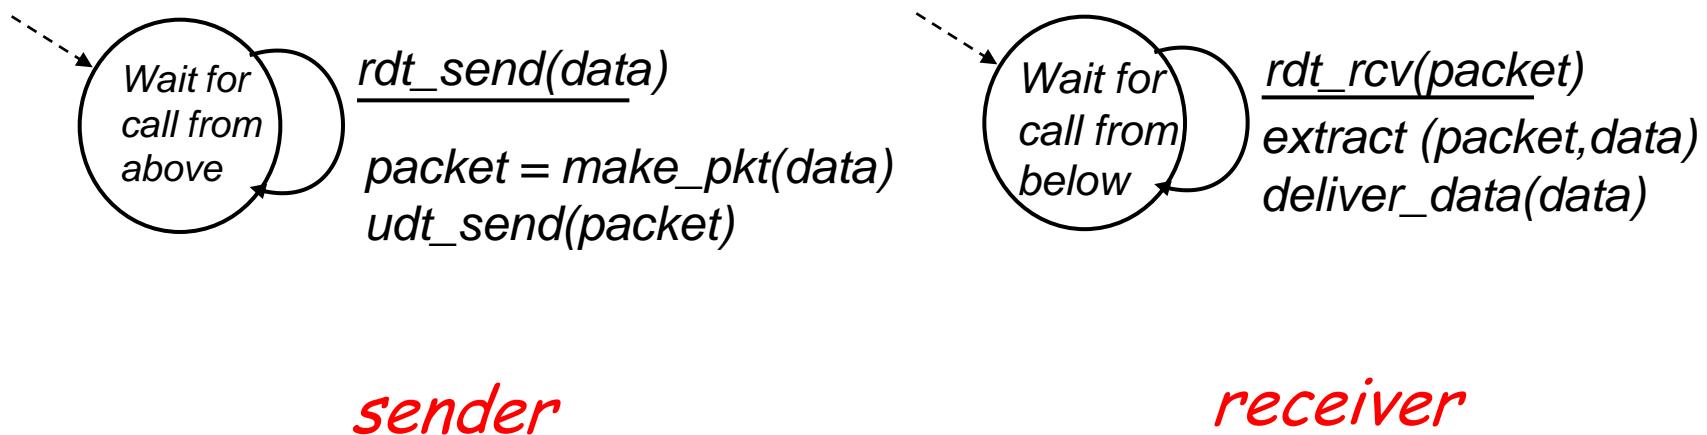
Rdt1.0: reliable transfer over a reliable channel

- ❖ underlying channel **perfectly reliable**

- no bit errors
- no loss of packets

- ❖ separate FSMs for sender, receiver:

- sender sends **data** into underlying channel
- receiver reads **data** from underlying channel



Rdt2.0: channel with bit errors

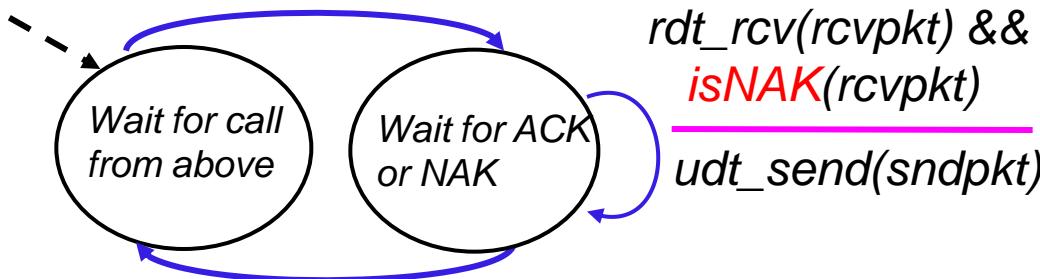
- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ the question: how to recover from errors:
(Note: Correction by receiver is very expensive)
 - acknowledgements (ACKs): receiver explicitly tells sender that pkt received is OK
 - negative acknowledgements (NAKs): receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK
 - sender transmits next pkt on receipt of ACK
- ❖ new mechanisms in rdt2.0 (beyond rdt1.0):
 - error detection
 - receiver feedback: control msgs (ACK,NAK)

rdt2.0: FSM specification

rdt_send(data)

sndpkt = make_pkt(data, checksum)

udt_send(sndpkt)



rdt_rcv(rcvpkt) && isACK(rcvpkt)

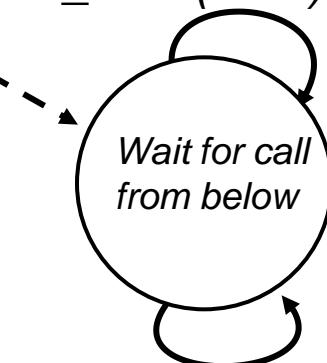
Λ

sender

receiver

rdt_rcv(rcvpkt) && corrupt(rcvpkt)

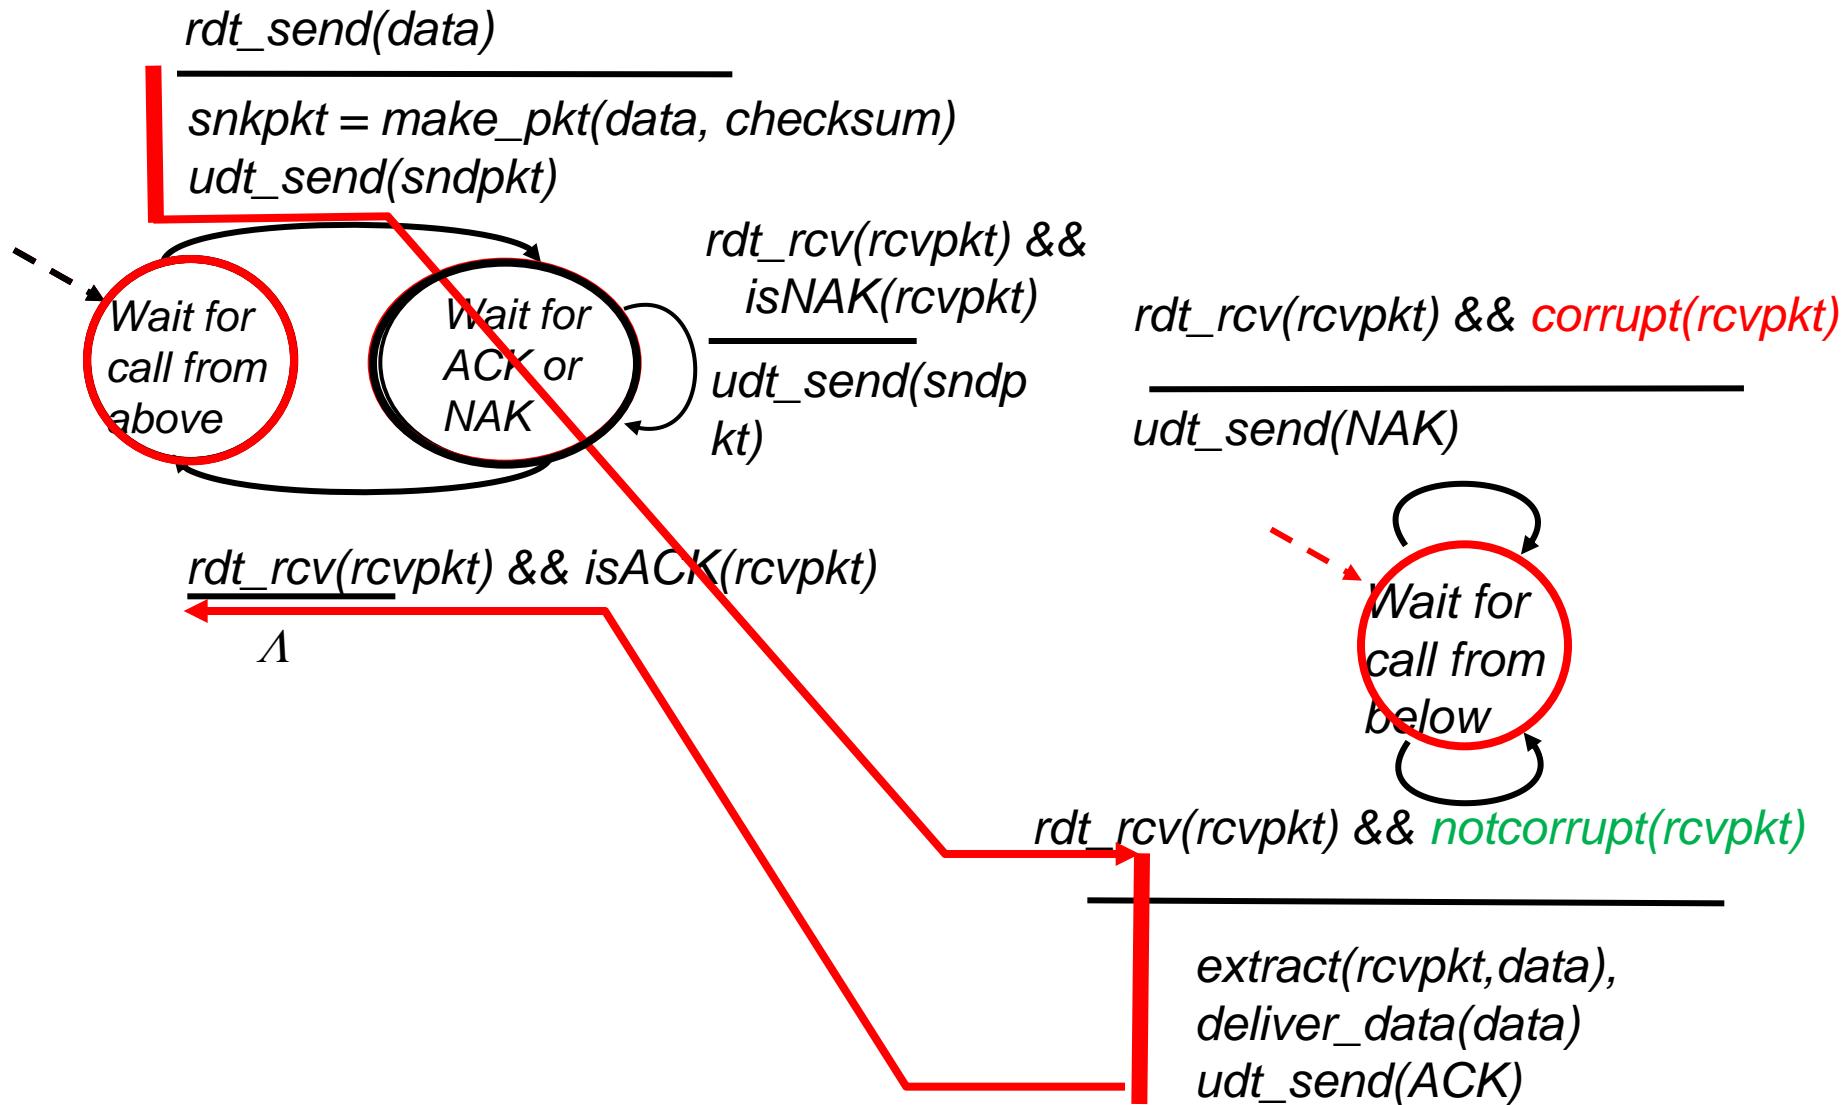
udt_send(NAK)



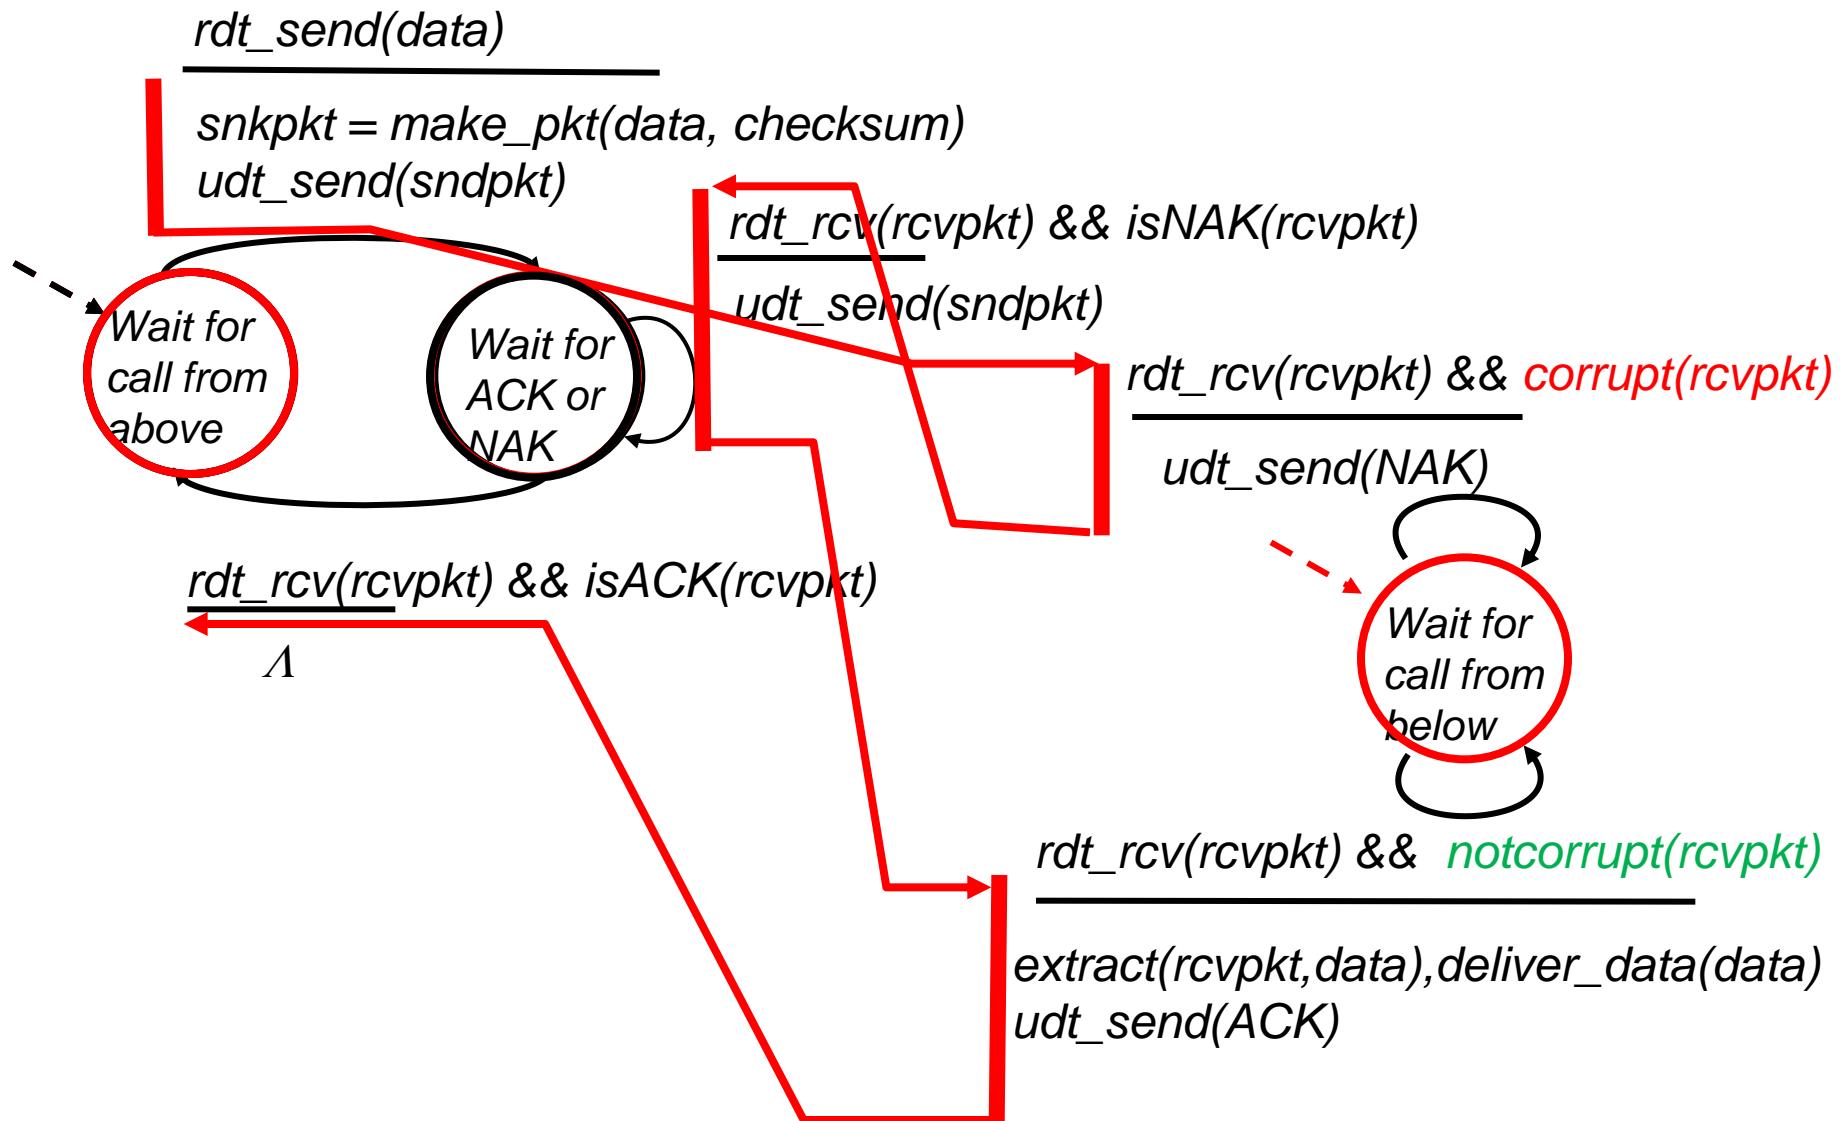
rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)

*extract(rcvpkt,data),
deliver_data(data)
udt_send(ACK)*

rdt2.0: operation with no errors



rdt2.0: error scenario



rdt2.0 has a fatal flaw!

What if ACK/NAK corrupted?

- ❖ sender doesn't know what happened at receiver!
- ❖ can't just retransmit: possible duplicate

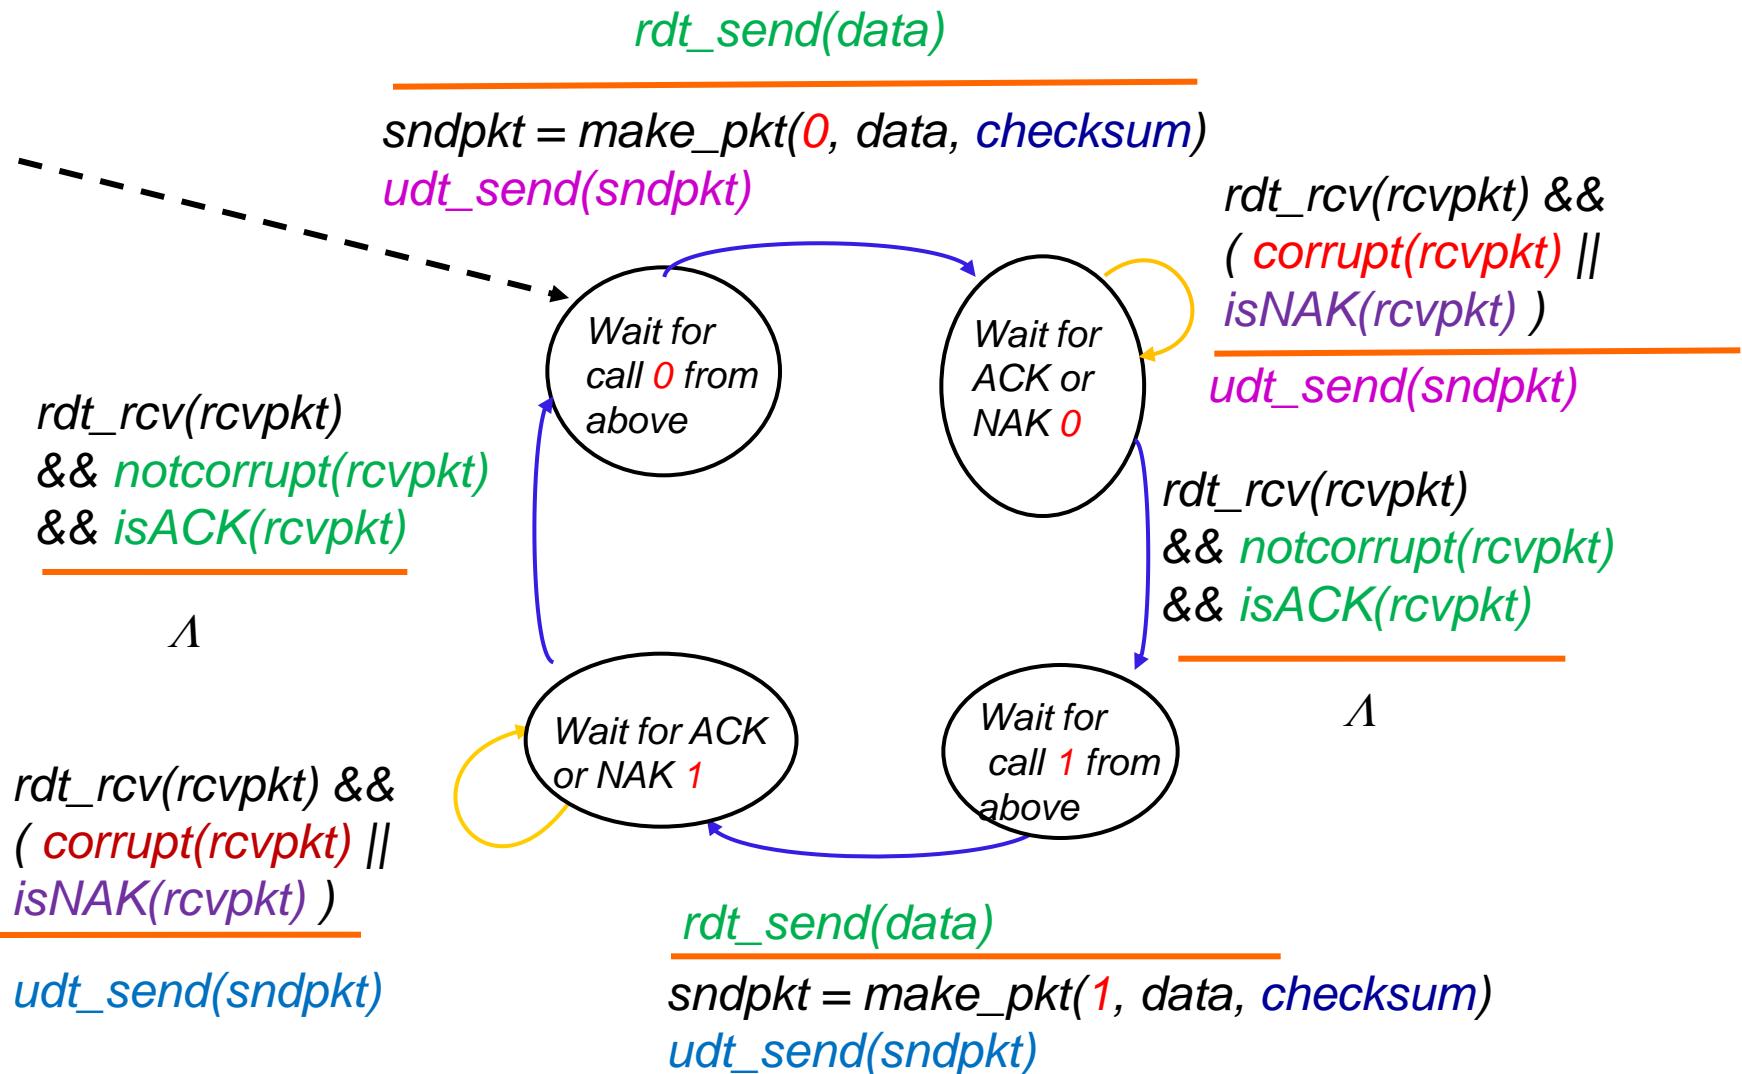
Handling duplicates:

- ❖ sender **retransmits** current pkt if ACK/NAK **garbled**
- ❖ sender adds **sequence number** to each pkt
- ❖ receiver **discards** duplicate pkt

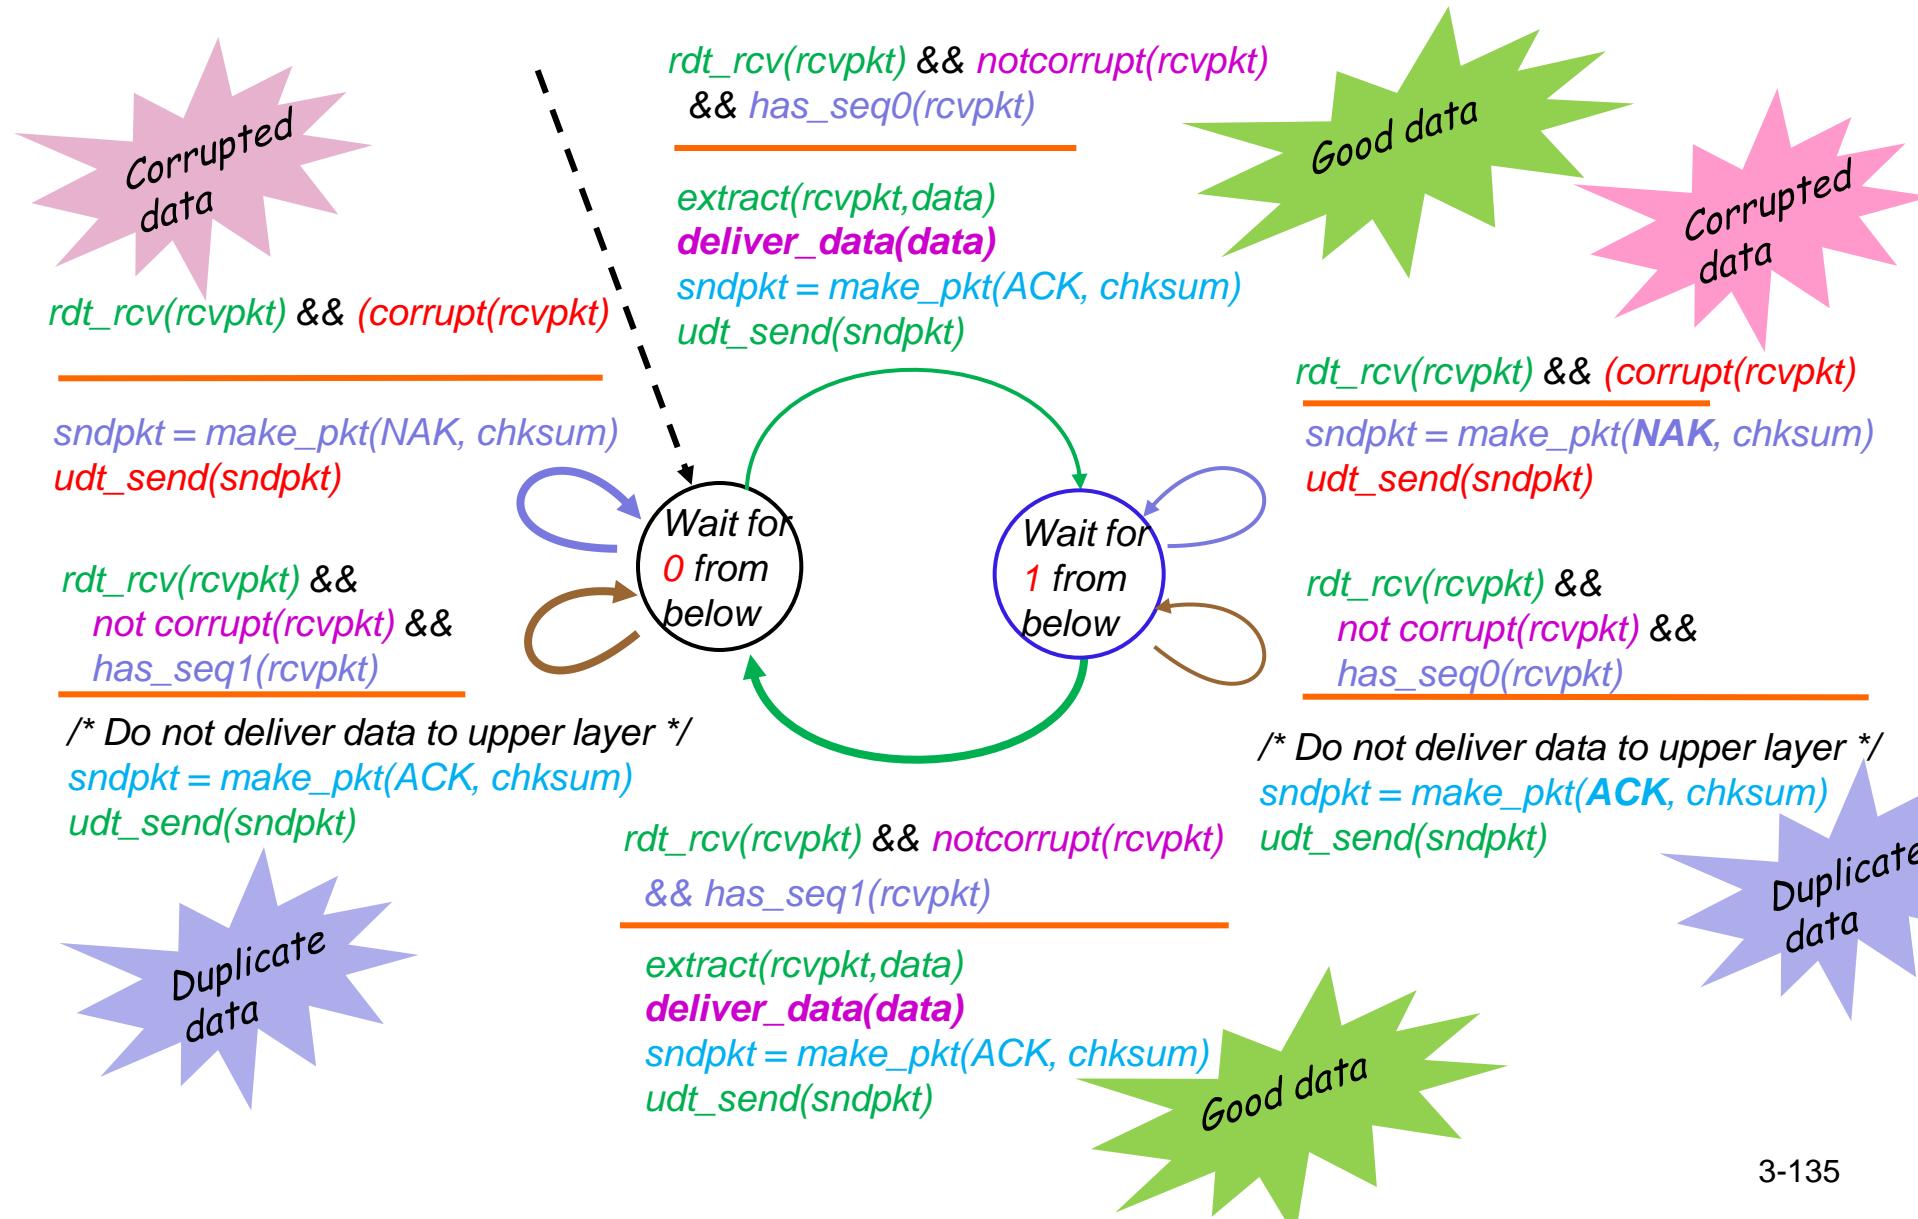
stop and wait
Sender **sends** one packet,
then **waits** for
receiver response

Note: Because of stop-and-wait, you need just two sequence numbers {0, 1}
Less than 2 is not a sequence number

rdt2.1: sender, handles garbled ACK/NAKs



rdt2.1: receiver, handles (garbled ACK/NAKs)



rdt2.1: discussion

Sender:

- ❖ seq # added to pkt
- ❖ two seq. #'s (0,1) will suffice.
- ❖ must check if received ACK/NAK corrupted
- ❖ twice as many states
 - state must "remember" whether "current" pkt has 0 or 1 seq. #

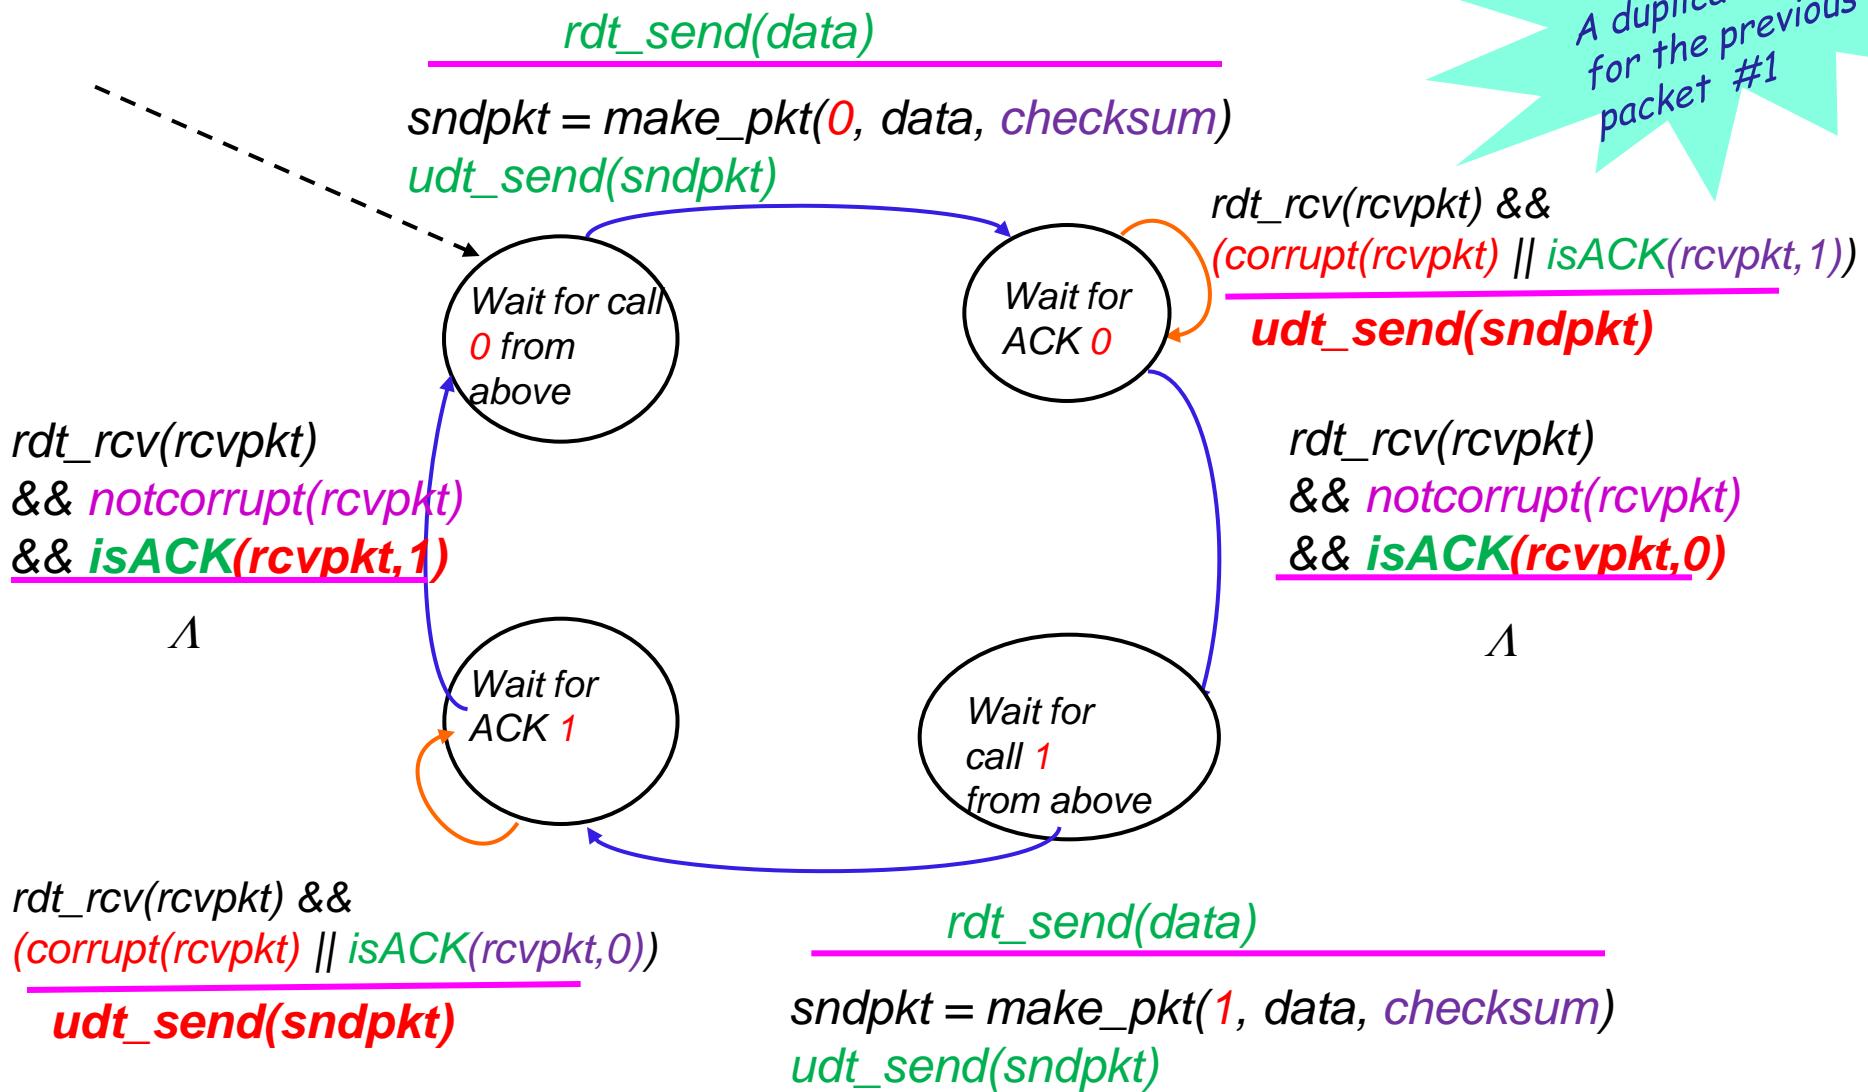
Receiver:

- ❖ must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #
- ❖ note: receiver can *not* know if its last ACK/NAK was received OK at sender

rdt2.2: a NAK-free protocol

- ❖ same functionality as rdt2.1, using *ACKs only*
- ❖ *instead of NAK, receiver sends ACK for last pkt received OK*
 - receiver must explicitly include seq # of pkt being ACKed
- ❖ *duplicate ACK at sender results in same action as NAK: retransmit current pkt*

rdt2.2: sender fragment



rdt2.2: receiver fragment

$rdt_rcv(rcvpkt) \&\& \text{notcorrupt}(rcvpkt) \&\& \text{has_seq0}(rcvpkt)$

$\text{extract}(rcvpkt, \text{data})$

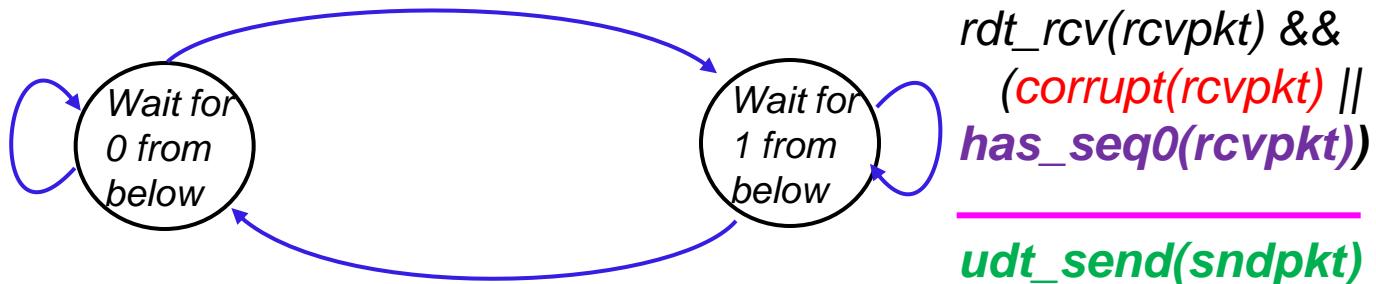
$\text{deliver_data}(\text{data})$

$\text{sndpkt} = \text{make_pkt}(\text{ACK0}, \text{chksum})$

$\text{udt_send}(\text{sndpkt})$

$rdt_rcv(rcvpkt) \&\&$
 $(\text{corrupt}(rcvpkt) \mid\mid$
 $\text{has_seq1}(rcvpkt))$

$\text{udt_send}(\text{sndpkt})$



$rdt_rcv(rcvpkt) \&\& \text{notcorrupt}(rcvpkt) \&\& \text{has_seq1}(rcvpkt)$

$\text{extract}(rcvpkt, \text{data})$

$\text{deliver_data}(\text{data})$

$\text{sndpkt} = \text{make_pkt}(\text{ACK1}, \text{chksum})$

$\text{udt_send}(\text{sndpkt})$

rdt3.0: channels with errors and loss

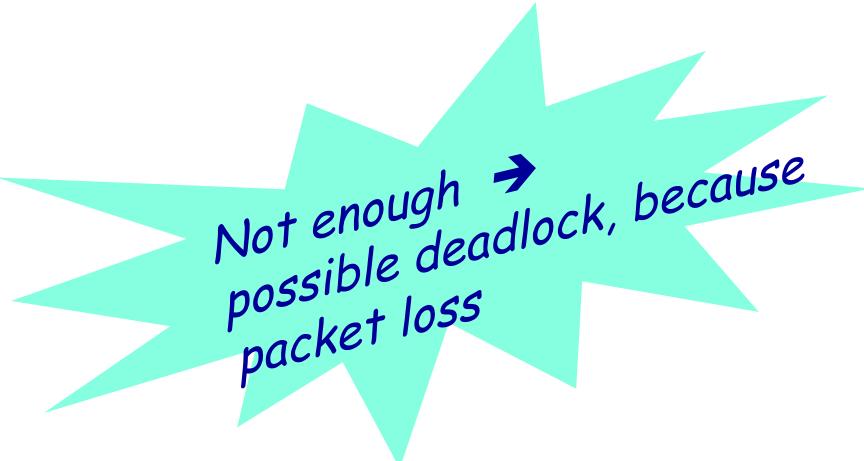
New assumption:

Underlying channel can also **lose** packets (data or ACKs)

- checksum, seq. #, ACKs, retransmissions will be of help, **but not enough**

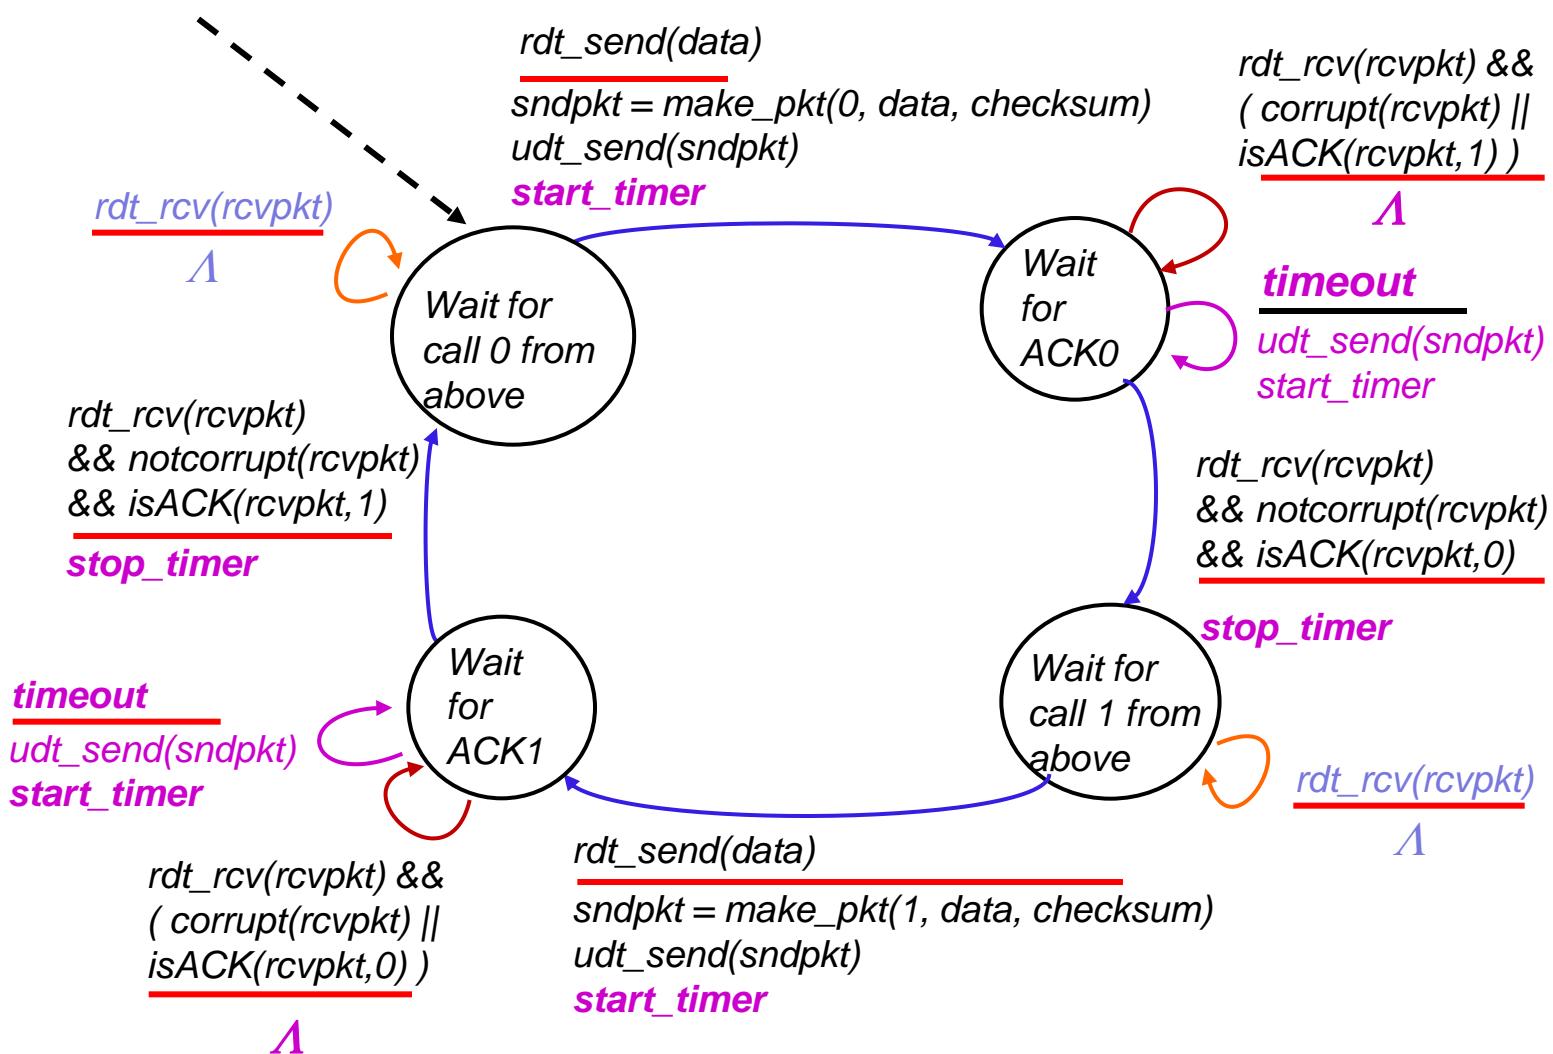
Approach: sender waits "reasonable" amount of time for ACK

- ❖ retransmits if no ACK is received in this time
- ❖ if pkt (or ACK) just delayed (not lost):
 - retransmission will be **duplicate**, but use of seq. #'s already handles this

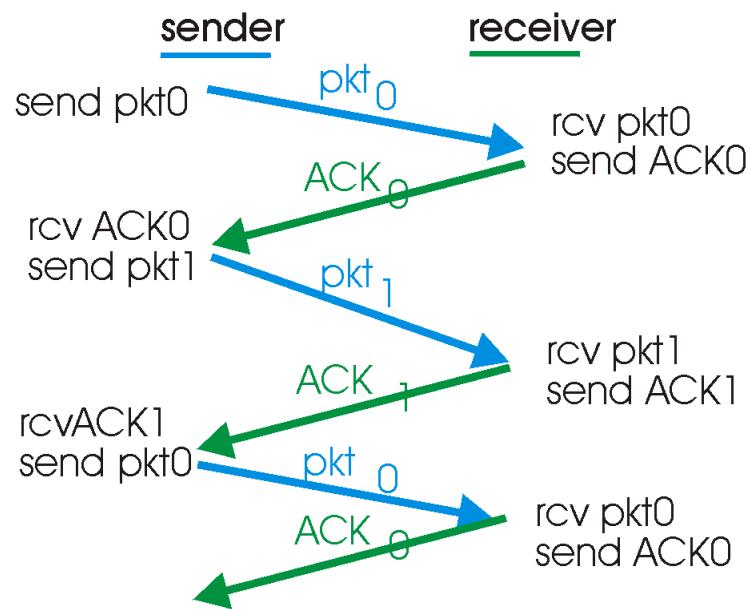


Not enough → possible deadlock, because packet loss

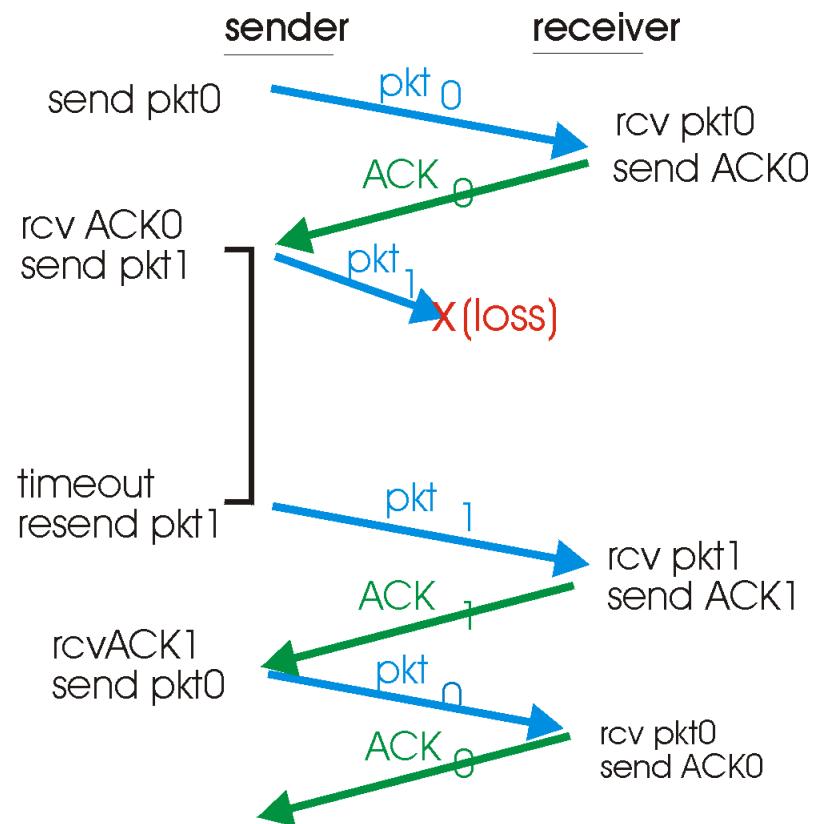
rdt3.0 sender (rdt2.2 + Timer + minor changes)



rdt3.0 in action

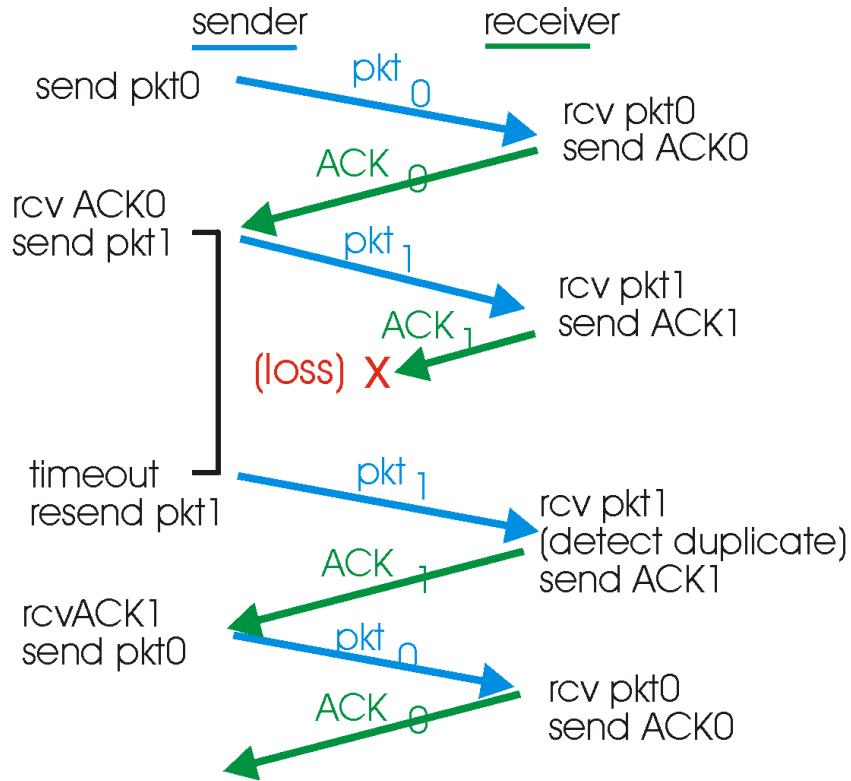


(a) operation with no loss

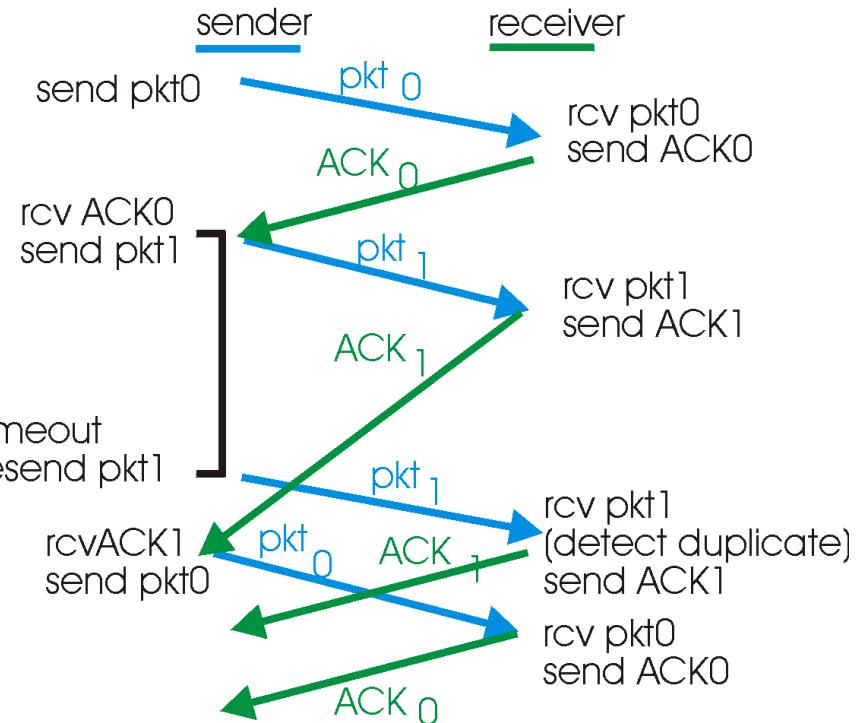


(b) lost packet

rdt3.0 in action



(c) lost ACK

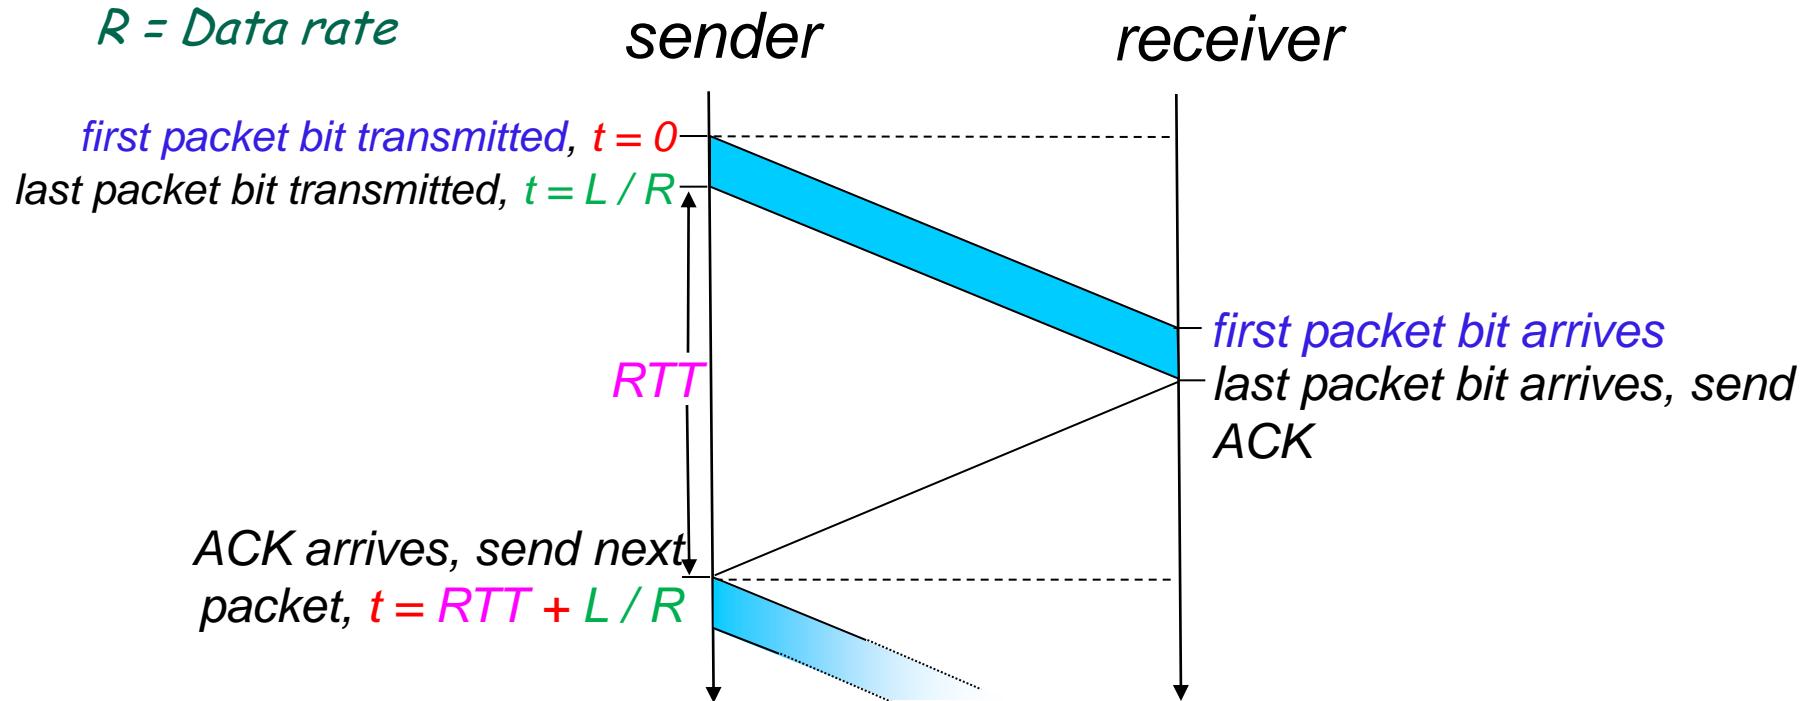


(d) premature timeout

rdt3.0: stop-and-wait operation

L = Packet length

R = Data rate



RTT: Round-Trip Time = $2 \times$ propagation delay.

Performance of rdt3.0

- ❖ ex: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

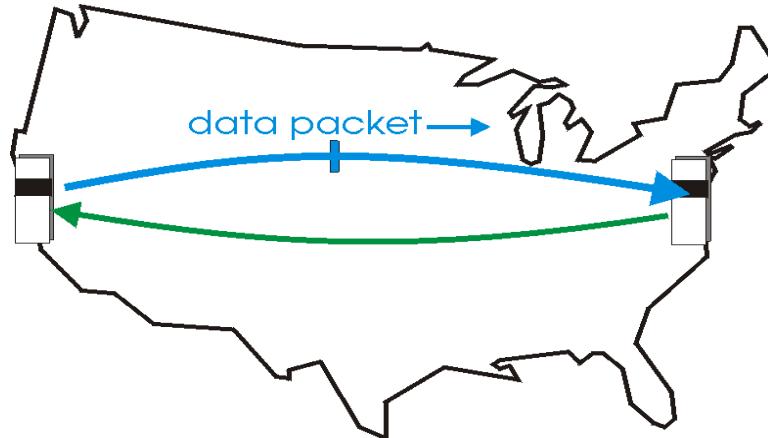
$$d_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bps}} = 8 \text{ microseconds}$$

- U_{sender} : utilization (= fraction of time sender is busy sending)

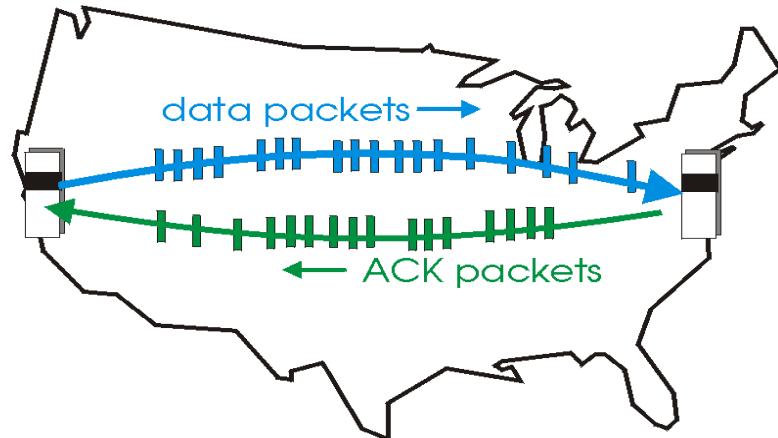
$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- if $RTT=30 \text{ msec}$, $1KB \text{ pkt every } 30 \text{ msec} \rightarrow 33KB/\text{sec}$ throughput over 1 Gbps link
- network protocol limits use of physical resources!

Pipelined protocols



(a) a stop-and-wait protocol in operation

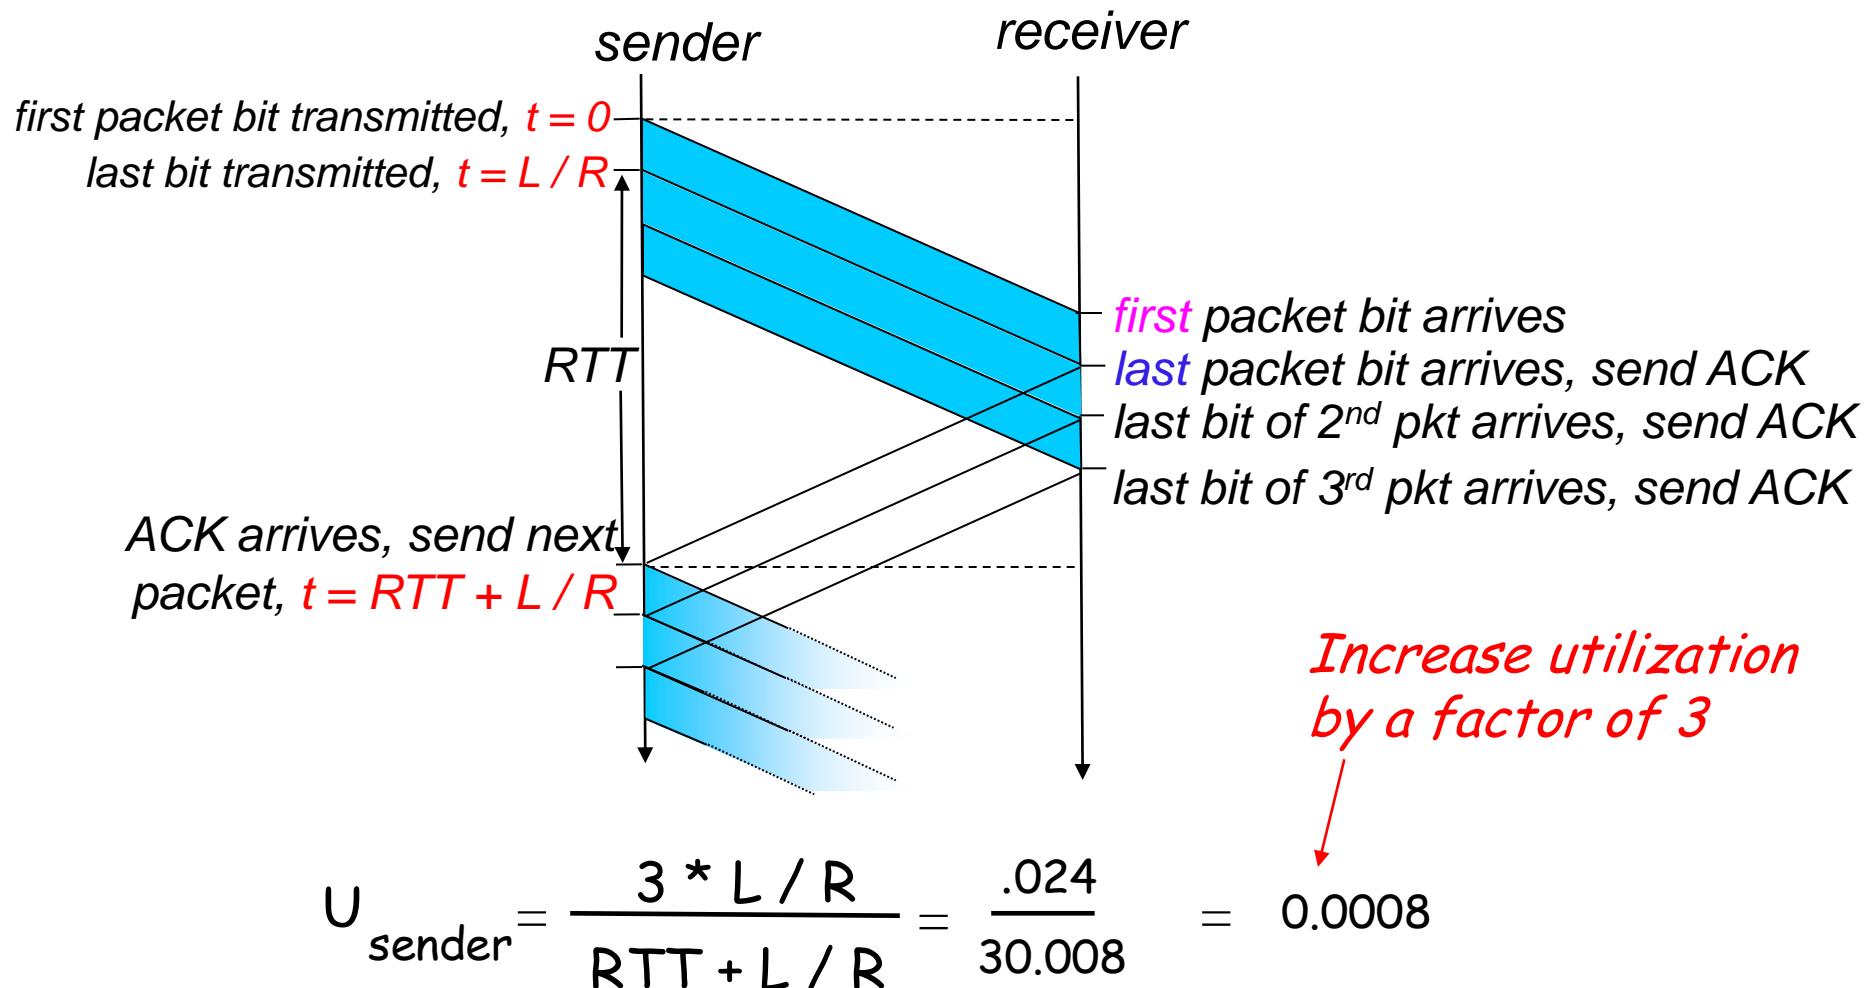


(b) a pipelined protocol in operation

pipelining: sender allows multiple, "in-flight", yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
(sequence # {0, 1} are not enough.)
- ❖ two generic forms of pipelined protocols:
go-Back-N and **selective repeat**

Pipelining: increases utilization



Pipelined Protocols (Sliding Window Protocols)

Go-back-N: big picture:

- ❖ sender can have up to N unACKed pkts in pipeline
- ❖ rcvr sends **individual ack** for each packet
- ❖ Sender interprets **acks** as **cumulative acks**

(Cumulative ACK: If ACK for pkt #5 is received, all pkts including and up to #5 have been correctly received)

- ❖ sender has a **timer** for **oldest** unACKed packet
 - if timer expires, **retransmit** all unACKed pkts

Selective Repeat: big pic

- ❖ sender can have up to N unACKed pkts in pipeline
- ❖ rcvr sends **individual ack** for each packet

- ❖ sender has **one timer** for **each** unACKed packet
 - when a timer expires, **retransmit only one pkt** corresponding to the timeout

Pipelined Protocols

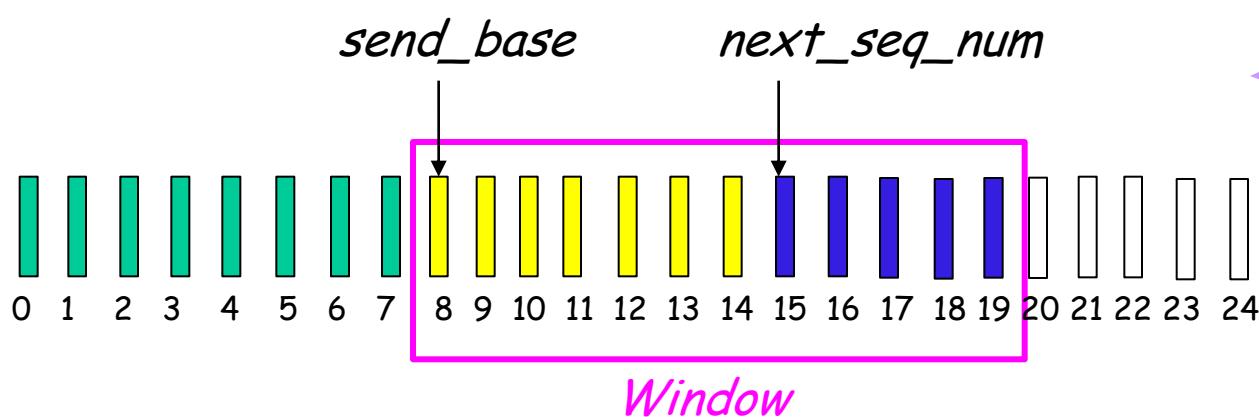
Remember:
The idea of *cumulative ACK*
is protocol specific.....

In TCP,
if the sender has received ACK
for data byte #125, it means
all bytes up to and including byte #124
have been received correctly and
the receiver is **expecting byte #125**.

Go-Back-N: Sender



- ❖ k-bit seq # in pkt header: for $k = 5$, seq.# are: 0, 1, 2, 3, 31.
- ❖ “window” of up to N, consecutive unACKed pkts allowed



- ❖ `timeout(n)`: retransmit pkt n and all higher seq # pkts in window

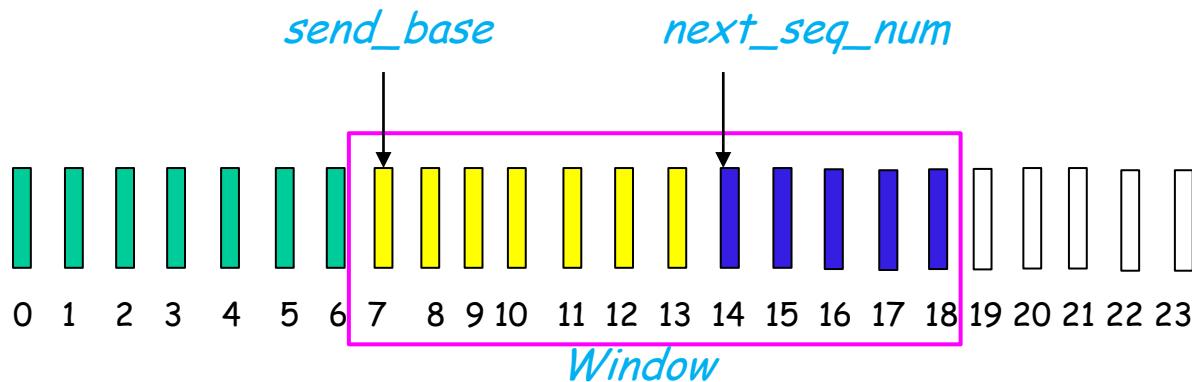
Go-Back-N: Sender

Already ACKed

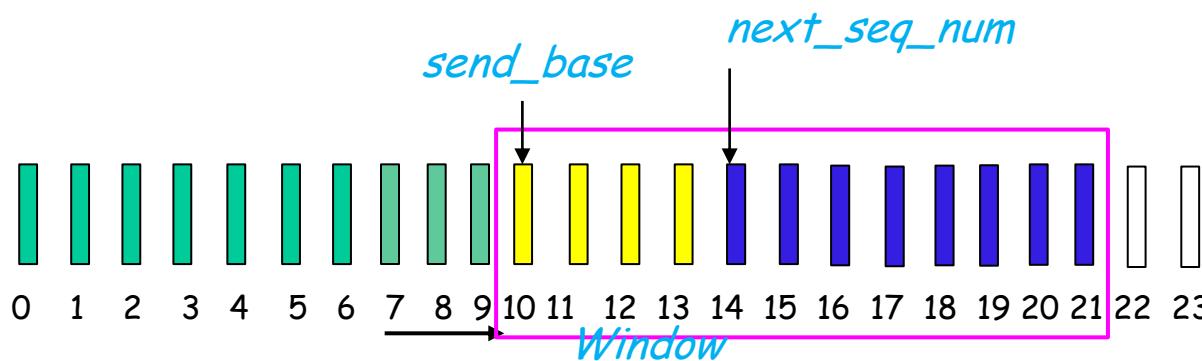
Sent, but ACK not received

Usable, not sent yet

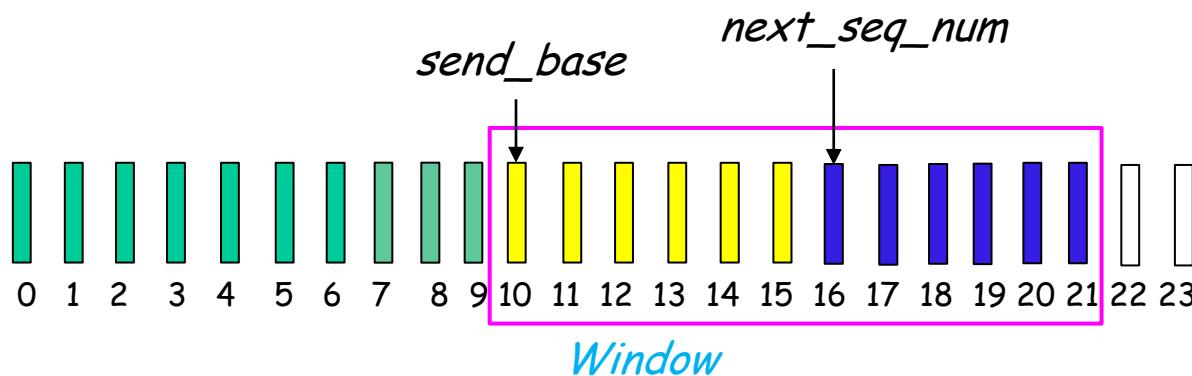
Not usable



Window snapshot



The sender receives ACKs for 3 pkts



The sender receives 2 data pkts from upper layer

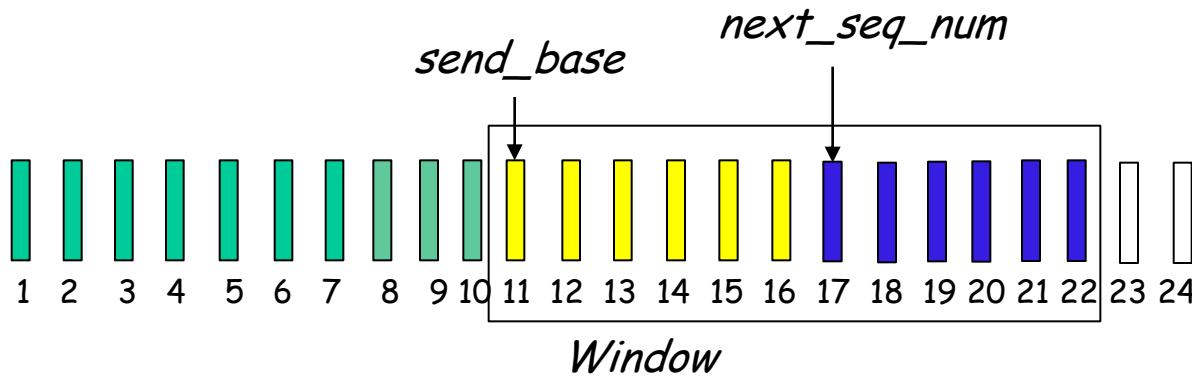
Go-Back-N: Sender

Already ACKed

Sent, but ACK
not received

Usable, not
sent yet

Not usable



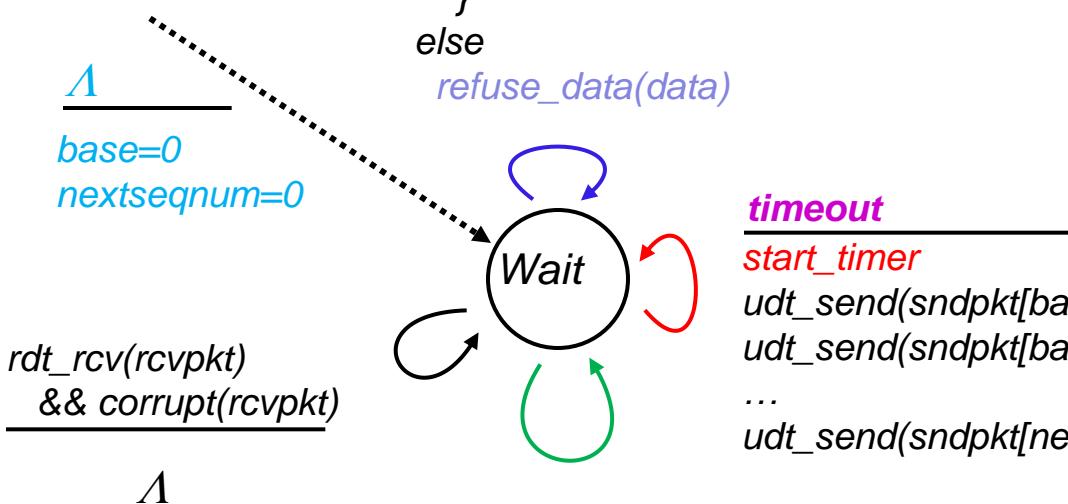
Timeout occurs

Window
snapshot

GBN: sender extended FSM

rdt_send(data)

```
if (nextseqnum < base+N) {  
    sndpkt[nextseqnum] = make_pkt(nextseqnum,data,chksum)  
    udt_send(sndpkt[nextseqnum])  
    if (base == nextseqnum) start_timer  
    nextseqnum++  
}  
else  
    refuse_data(data)
```



timeout

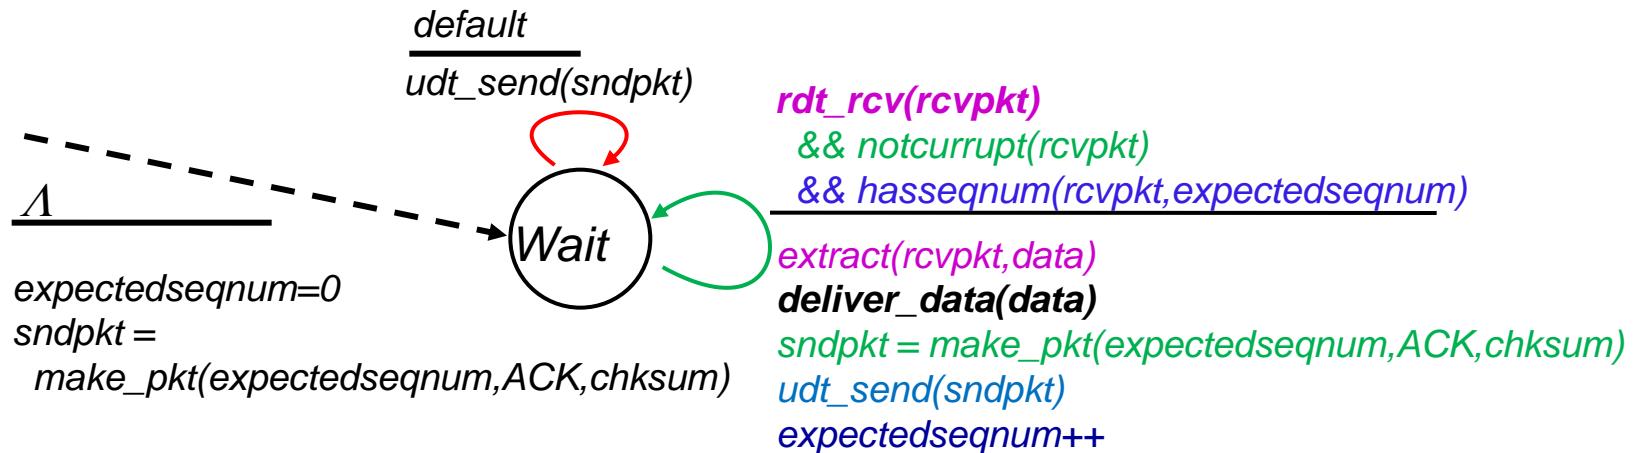
```
start_timer  
udt_send(sndpkt[base])  
udt_send(sndpkt[base+1])  
...  
udt_send(sndpkt[nextseqnum-1])
```

rdt_rcv(rcvpkt) && notcorrupt(rcvpkt)

base = *getacknum(rcvpkt)*+1

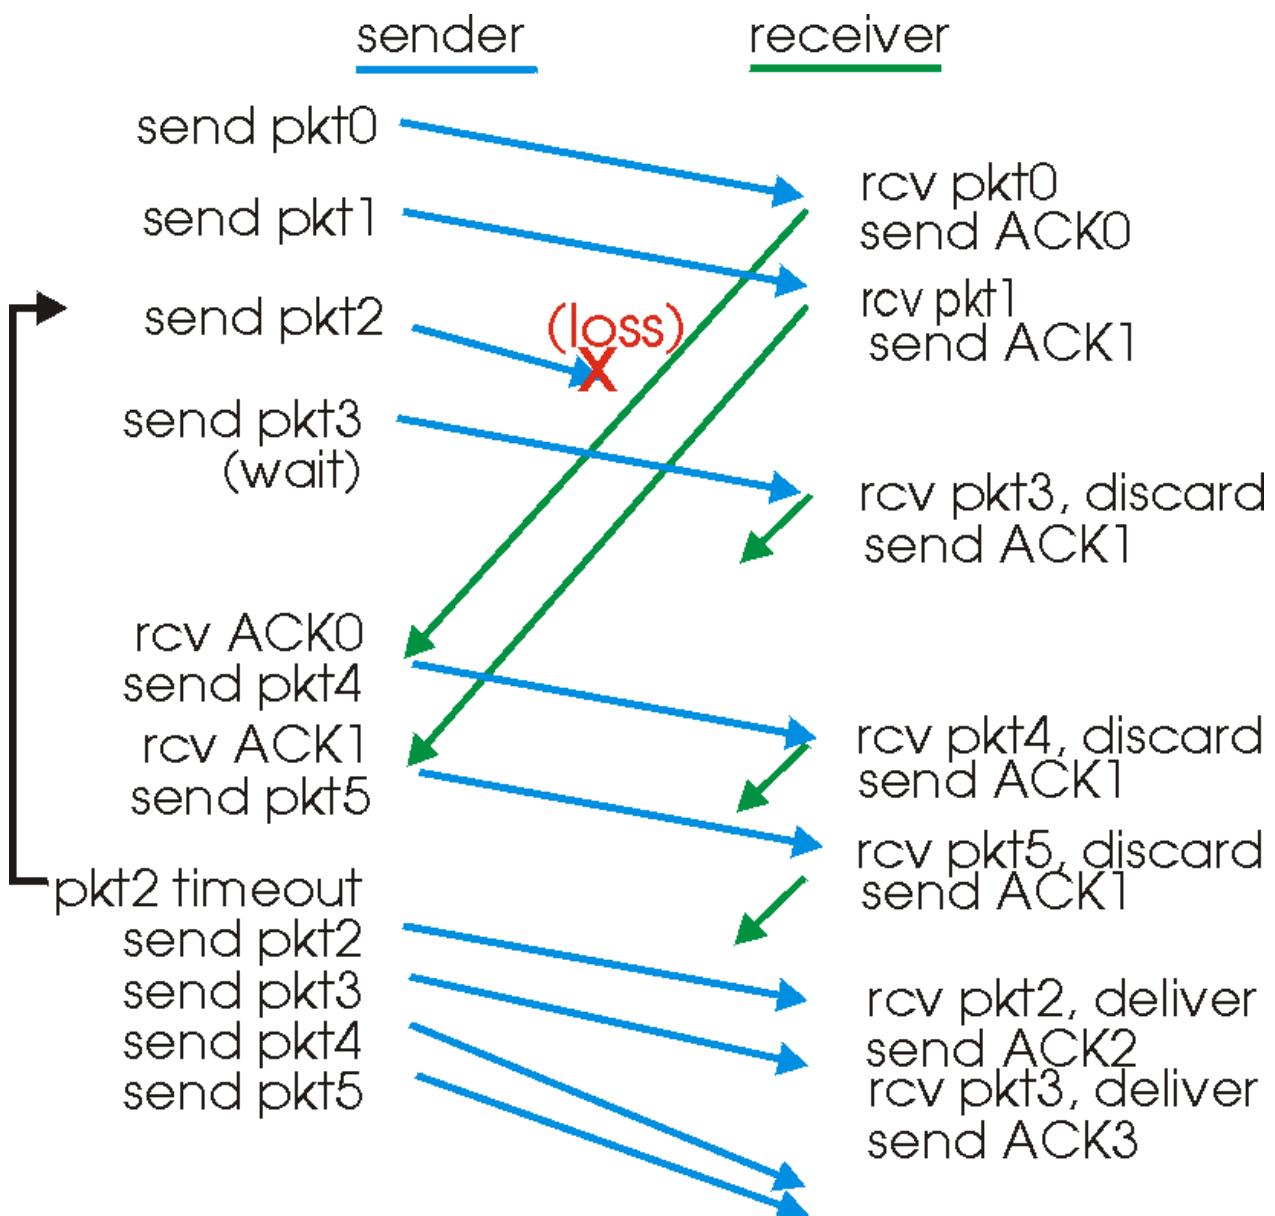
If (base == nextseqnum) stop_timer // All packets have been ACKed...
else start_timer

GBN: receiver extended FSM



- ❖ ACK-only: always send ACK for correctly-received pkt with highest in-order seq #
 - need only remember **expectedseqnum**
- ❖ out-of-order pkt:
 - discard (don't buffer) -> **no receiver buffering!**
 - Re-ACK pkt with highest in-order seq #

GBN in action



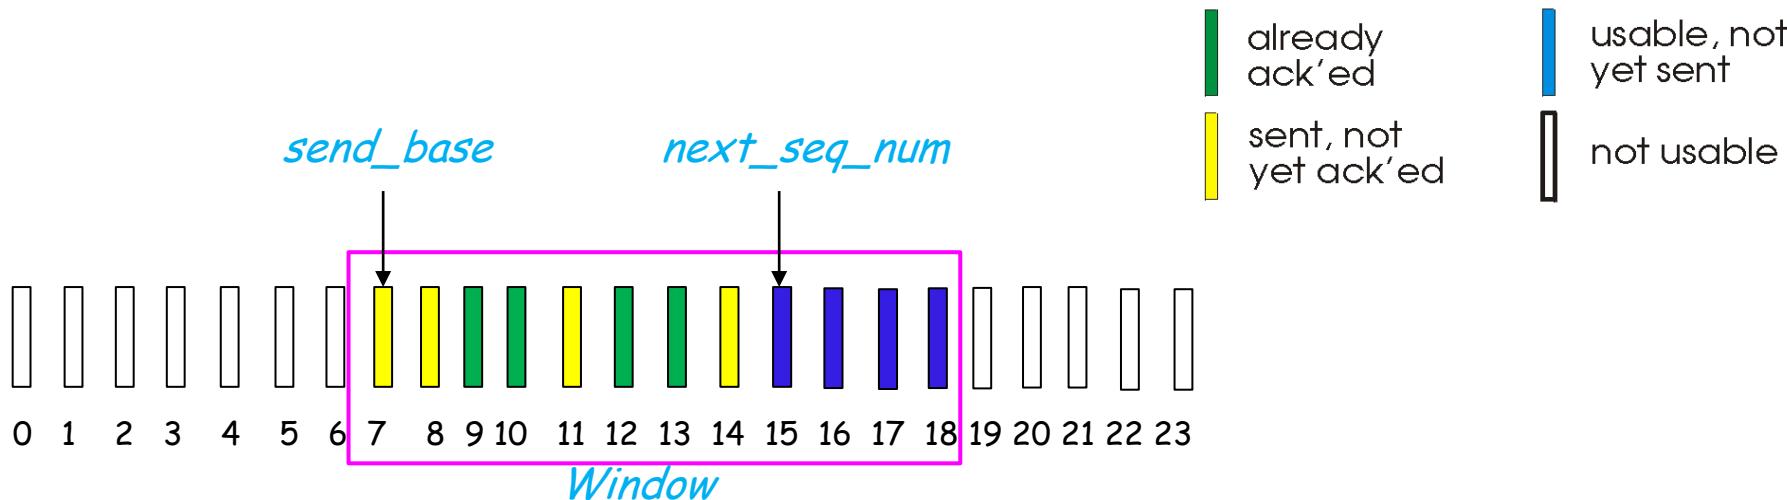
Selective Repeat



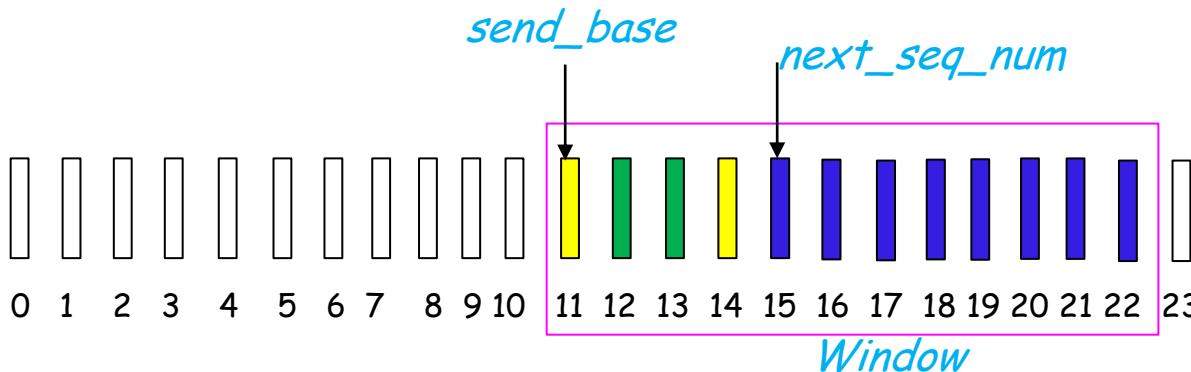
- ❖ receiver *individually ACKs all correctly received pkts*
 - buffers pkts, as needed,
for eventual in-order delivery to upper layer

- ❖ *sender only resends pkts for which ACK is not received and timeout occurs*

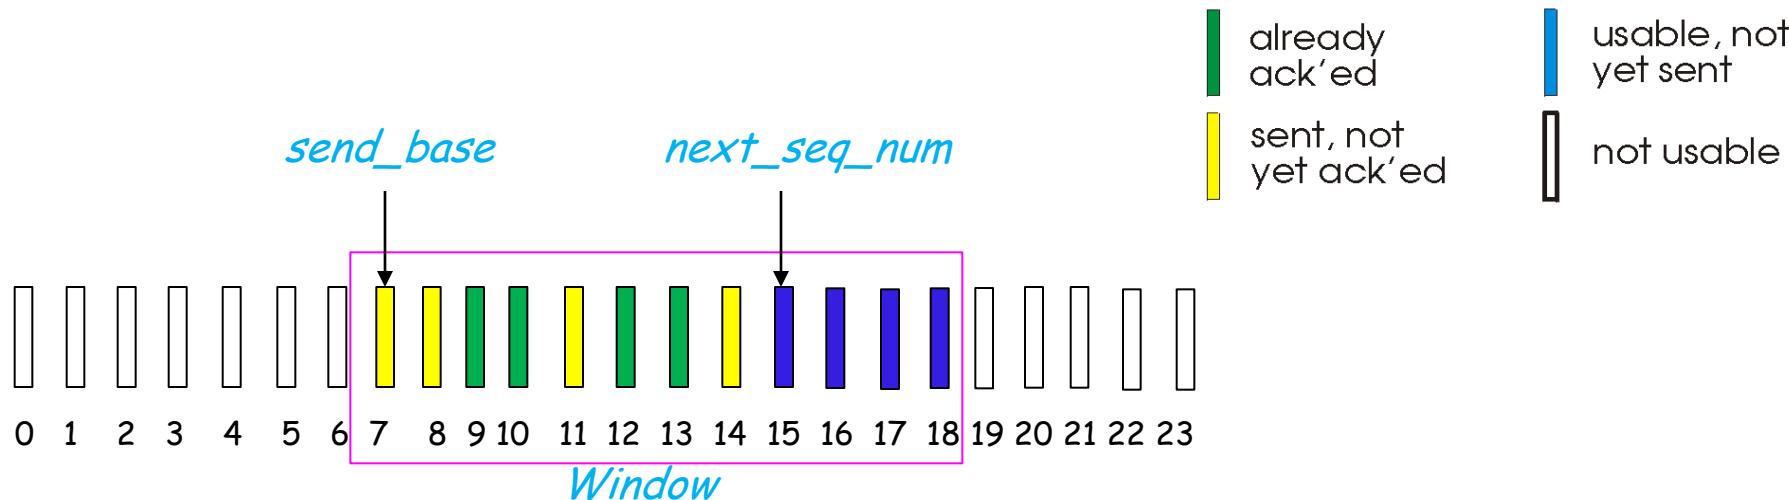
Selective repeat: sender window



ACK received for 7 and 8.



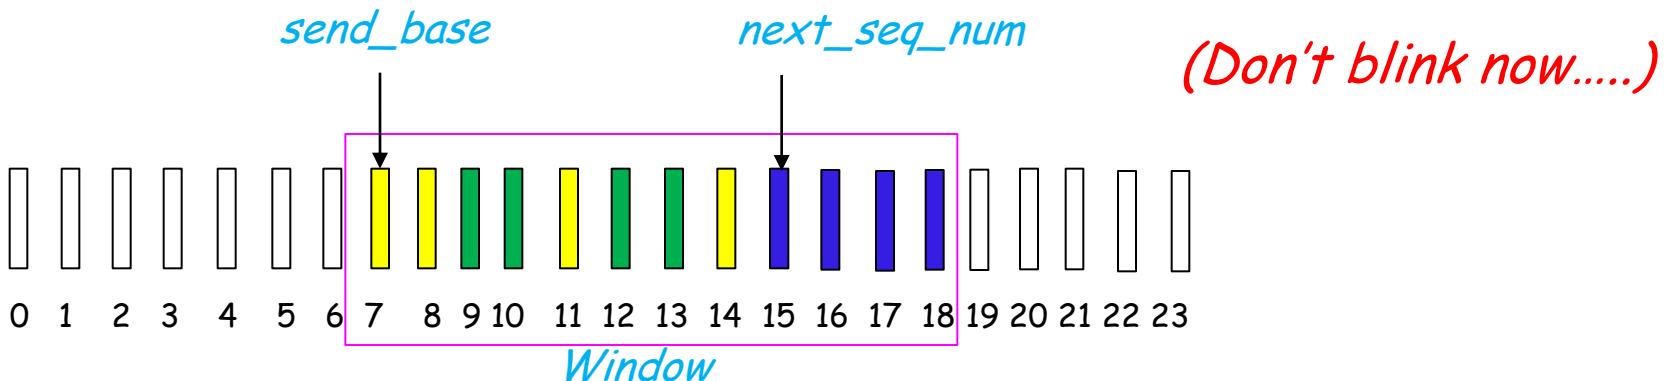
Selective repeat: sender window



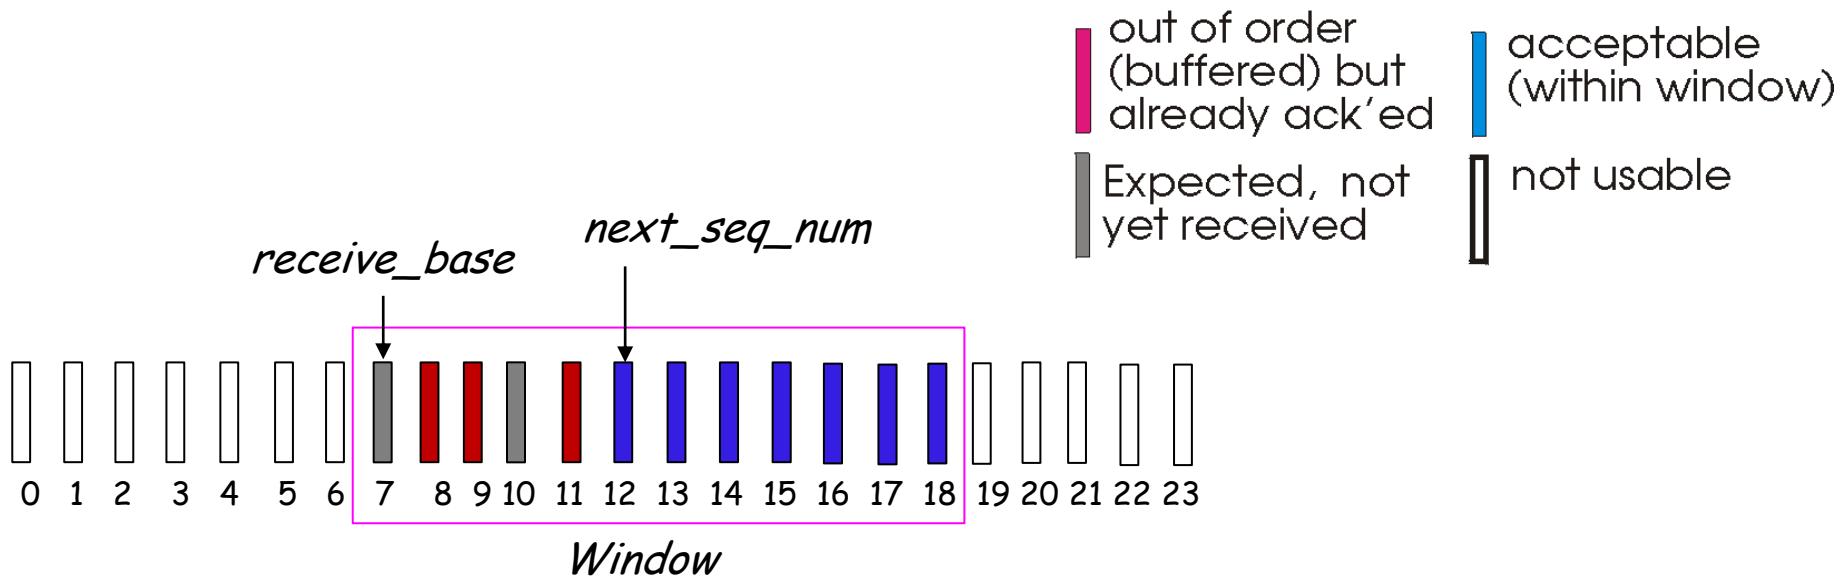
Retransmit

Timeout occurs for #7

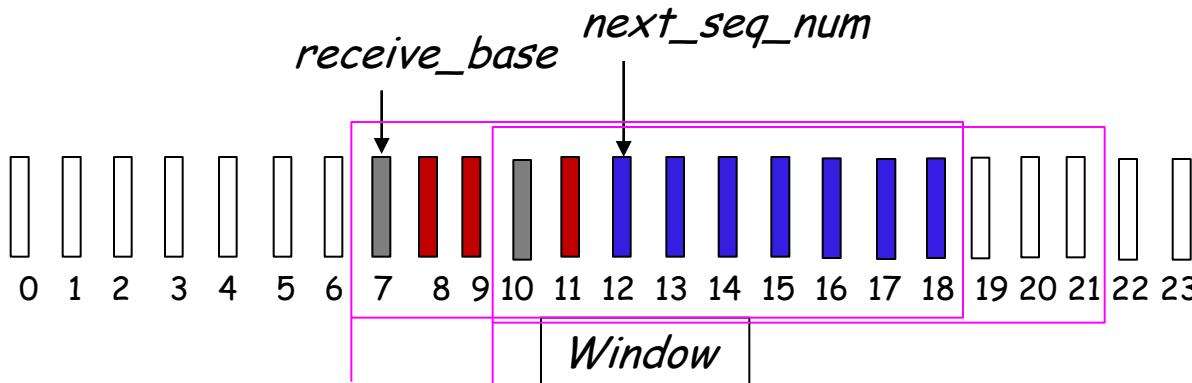
2 new packets arrive from upper layer



Selective repeat: receiver window



Pkt #7 is received



Selective repeat

sender

data from above :

- ❖ if next available seq # in window, send pkt

timeout(n):

- ❖ resend pkt **n**, restart timer

ACK(n) in [sendbase,sendbase+N]:

- ❖ mark pkt **n** as received
- ❖ if **n** smallest unACKed pkt, advance window base to next unACKed seq #

receiver

pkt n in $[rcvbase, rcvbase+N-1]$

- ❖ send ACK(n)
- ❖ out-of-order: buffer
- ❖ in-order: deliver (also deliver buffered, in-order pkts), advance window to next not-yet-received pkt

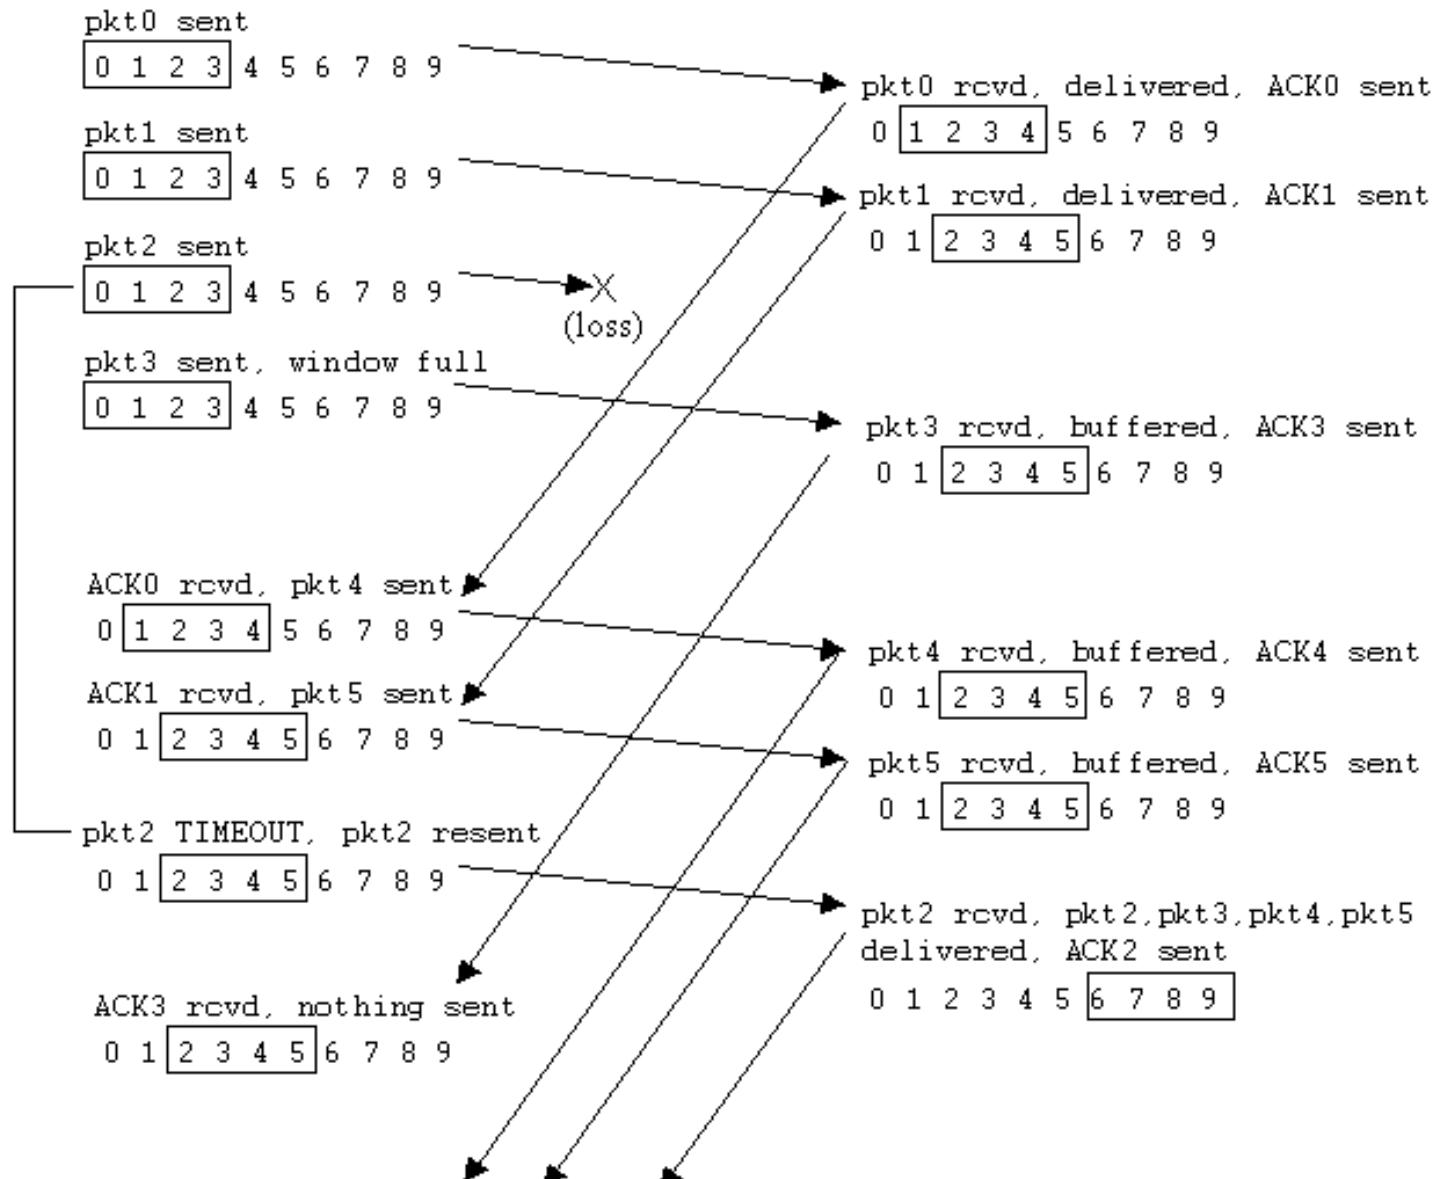
pkt n in $[rcvbase-N, rcvbase-1]$

- ❖ ACK(n)

otherwise:

- ❖ ignore

Selective repeat in action

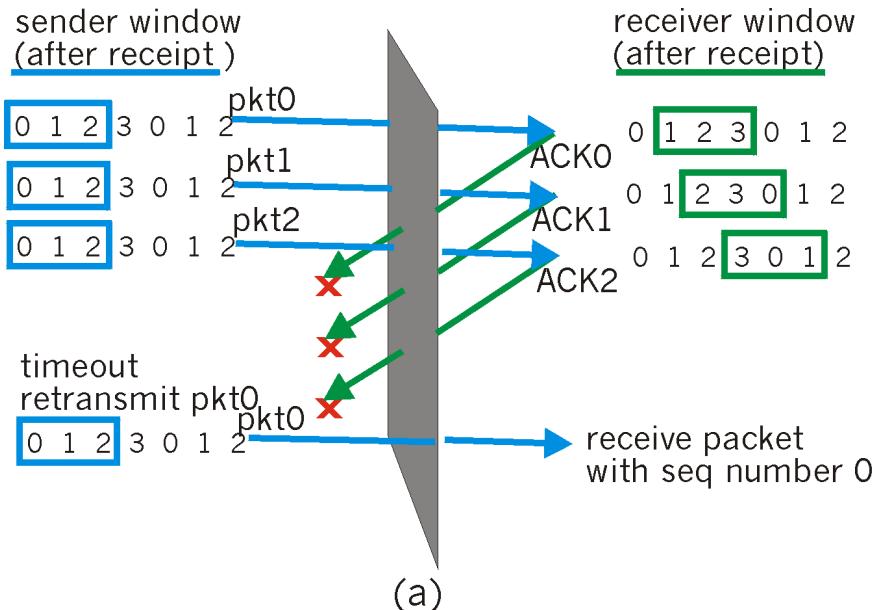


Selective repeat: dilemma

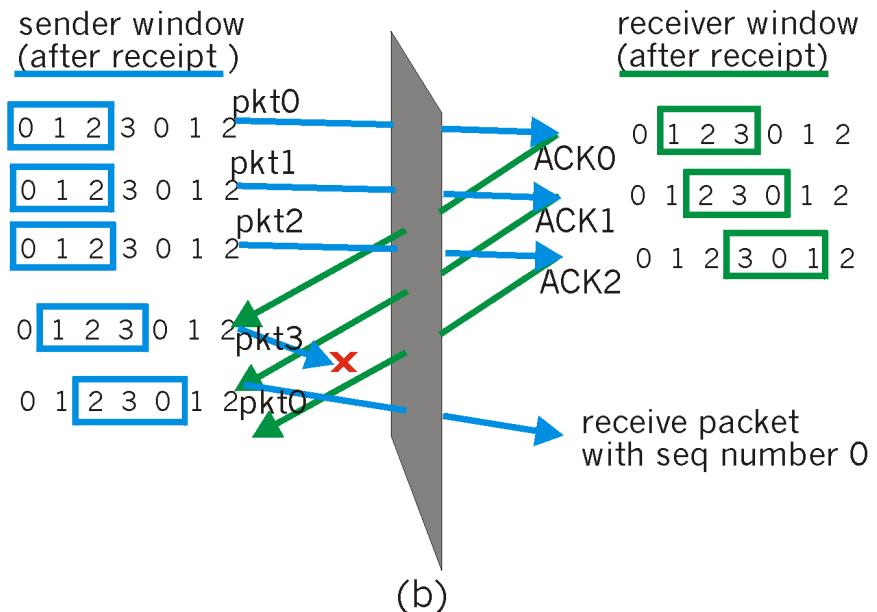
Example:

- ❖ seq #'s: 0, 1, 2, 3
- ❖ window size = 3

incorrectly passes duplicate data as new in (a)



receiver sees no difference in two scenarios!



End of Link Layer